

# AB-650

# Compact Exam Course

## Microsoft 365 and AI services administration

Version 3.0 • Updated 7 October 2026

10 focused modules • worked decisions • practical evidence tasks

Companion HTML: 20 module checks and two 30-question mixed tests with automatic marking and explanations.

Read the deciding rule, apply it to the case and prove the outcome. Use the linked Microsoft procedures and official labs for hands-on depth.

Independent preparation. No guaranteed exam score, calibrated readiness prediction or verified pass-rate claim.

# Reading map

Read one module, explain its control and complete its two checks in the companion HTML. Use the practical task to demonstrate the result. References provide current procedure details. Page entries below are clickable.

|                                                   |           |
|---------------------------------------------------|-----------|
| <b>01. Tenant, licences and recovery</b>          | <b>4</b>  |
| <b>02. Mail, Teams and SharePoint boundaries</b>  | <b>6</b>  |
| <b>03. Identity lifecycle and least privilege</b> | <b>8</b>  |
| <b>04. Authentication and Conditional Access</b>  | <b>10</b> |
| <b>05. Defender policy and incident decisions</b> | <b>12</b> |
| <b>06. Purview: classify, prevent and retain</b>  | <b>14</b> |
| <b>07. Copilot readiness and user settings</b>    | <b>16</b> |
| <b>08. Connectors and enterprise search</b>       | <b>18</b> |
| <b>09. Agent identity, access and registry</b>    | <b>20</b> |
| <b>10. AI monitoring, cost and exam decisions</b> | <b>22</b> |

# Study for the decision, not the page count

This edition replaces the long, repetitive core with ten focused modules. Each gives the mechanism, a worked case, a decision table and a practical evidence task. Extra depth means more precise reasoning, not more definitions.

**1. Read:** explain the module's main control without notes. Open its Microsoft references for an unfamiliar procedure or current prerequisite.

**2. Check:** answer the two module questions. Explain why every alternative fails the stated requirement.

**3. Apply:** complete the evidence task and the matching official lab when you have appropriate access. Simulation and live execution are recorded separately.

**4. Test:** take Mixed A without notes before opening its key. Use errors and correct guesses to choose what to restudy.

**5. Recheck:** take Mixed B on a different day. Repeat questions practise recall; they are weaker readiness evidence.

**Local study target:** at least 27/30 on each mixed test, with at least 90% in each sampled domain, plus practical evidence and correct explanations of the critical distinctions below. At least 29/30 is a stretch target. These are uncalibrated course targets; meeting them does not establish a probability of passing Microsoft's exam.

Each mixed test contains 7 tenant/workload, 12 security/governance and 11 AI questions. This follows the published domain ranges. Small domain samples are diagnostic, not precise ability estimates. The bank is original practice, not actual exam questions.

| Domain                      | Published weighting | Study emphasis                                 |
|-----------------------------|---------------------|------------------------------------------------|
| D1: tenant and workloads    | 20â–25%             | Provisioning, licences, access and recovery    |
| D2: governance and security | 40â–45%             | Identity, authentication, Defender and Purview |
| D3: AI services             | 35â–40%             | Copilot, connectors, agents and monitoring     |

Use roughly one focused session per module, allowing extra time for security and AI. Start with an untimed explanation; then practise making the decision under a limit you choose. The course does not prescribe an official exam question count or timing.

# 01. Tenant, licences and recovery

A tenant is the organisation's Microsoft 365 environment. A licence supplies product entitlement. An administrator role supplies permission to change settings. Pay-as-you-go pays for eligible consumption. None of these alone proves that a user can use a feature. A custom domain is a name you control; verification proves ownership, while DNS records determine service routing.

## Choose the exact object

For branding, distinguish Microsoft 365 organisational themes from Entra sign-in branding. Start in the relevant admin centre and check the preview as a normal user. Organisation profile, security and privacy settings are tenant settings; a user's profile is a different object. Record the original value and intended audience before a change. For domains, add the name, publish the requested ownership record and verify it. Plan service DNS separately: changing MX redirects incoming mail; a verification TXT record does not. Do not switch MX while mailboxes remain at the old provider.

## Find the licensing fault

Use the Microsoft 365 admin centre's licensing views to inspect the product, group assignment and user error. Check available seats, the service plans selected, the user's usage location and conflicting plans. Group-based licensing processes direct user members; nested group membership does not supply licences. A blank usage location in group assignment can inherit the tenant's location, so blank alone is not evidence of failure. An explicit invalid location error is stronger evidence. For Copilot, Agent 365 and Copilot Studio, check that specific product's current licence and billing prerequisites. Do not reuse one product's rules for another. During a licensing-group move, confirm that the destination assignment has processed before removing the old source.

## Diagnose health before changing configuration

Write down the affected service, people, locations and UTC time. Check Service health for a matching incident and configure notifications for the services and recipients you need. An incident's scope must match the symptom. For a branch-only problem, compare connectivity insights with an unaffected location: DNS, routing, latency and path changes are useful evidence. A green service page does not prove every network path works. Usage reports describe activity, not live availability. Escalate with correlated times and affected endpoints, rather than a screenshot of one slow page.

## Prove recovery

Microsoft 365 Backup protects selected Exchange mailboxes, SharePoint sites and OneDrive accounts. Define the recovery point objective (acceptable data loss) and recovery time objective (acceptable recovery delay). Configure the workload policy, confirm the actual protection units and wait for recovery points. Activation is not instant. Select a point within the available recovery window and a supported restore destination. An original-location restore can replace live scoped content; use a permitted alternate destination for inspection where supported. Validate files, mail, metadata and access against the requirement. Retention, version history and an enabled backup policy do not demonstrate a successful restore.

## Worked case

Requirement: 12 analysts need Copilot. Evidence: the licensing group has 10 direct users and a nested group with 2 more. All 12 already have their base services.

Decision: inspect the Copilot assignment for each user. Add the two intended users through a supported direct-membership design, then recheck processing and entitlement. The group total is not a licence-success count.

Recovery evidence: an enabled SharePoint policy protects Site A, but the missing file was on Site B. The relevant protection scope is wrong; do not promise a restore from Site A.

## Choose the right control

| Need                         | Choose                                   | Do not confuse with |
|------------------------------|------------------------------------------|---------------------|
| Prove domain ownership       | Requested verification record            | MX cutover          |
| Entitle direct group members | Group licence assignment and error check | Nested membership   |
| Explain branch performance   | Connectivity comparison                  | Usage report        |
| Recover supported content    | Backup protection and restore evidence   | Retention alone     |

## Practical evidence task

1. In a training tenant, inspect an assigned product and list direct members, processed assignments and any errors. Use a read-only check first.
2. For a recovery exercise, record workload, protected resource, recovery point, destination and validation criteria. Follow the official lab for an authorised restore.
3. Evidence: one reconciled assignment list, one scoped recovery plan and a health/network diagnosis. Mark an unperformed restore as not tested.

Use synthetic data and the official lab instructions in an appropriate training environment. This guide supplies reasoning exercises, not a live lab tenant. Reading a task does not count as completing it.

## Microsoft references and procedure details

- [license microsoft 365 ai services](#)
- [maintain microsoft 365 tenant health](#)
- [3 verify troubleshoot license assignments](#)
- [2 configure microsoft 365 backup](#)

## 02. Mail, Teams and SharePoint boundaries

A workload contains resources and permissions. A shared mailbox, a Teams channel and a SharePoint site have different access models. The exam often gives a valid setting for the wrong resource. First name the resource, the acting person and the exact action.

### Mailbox type and delegation

A user mailbox belongs to a user; a shared mailbox supports a common address such as helpdesk. Full Access lets a delegate open and manage a mailbox. Send As sends as that mailbox. Send on Behalf identifies the delegate. Opening successfully does not prove sending rights. A shared mailbox up to 50 GB generally needs no separate mailbox licence, but delegates need their own Exchange Online entitlement. A larger primary mailbox or features such as archive and hold can change licensing needs. Keep the backing account when converting a licensed user mailbox to shared. Check size and feature requirements before removing its licence; conversion alone is not evidence that removal is safe.

### Teams membership and meeting Copilot

Set team owners deliberately and verify members. Standard channels follow team membership; private and shared channels introduce separate membership boundaries. Their files can reside in separate SharePoint sites, so review the actual channel site. For meeting Copilot, distinguish the administrator's meeting policy, the organiser's meeting option and whether transcription actually starts. If a saved transcript is mandatory, use the enforced saved-transcript policy mode, not a changeable default. Permission to transcribe does not start transcription. Validate the organiser and transcription starter's policies, licensing and supported meeting experience.

### SharePoint permissions versus discovery

A group-connected team site and a communication site support different collaboration patterns. Review owners, members, visitors, direct grants and sharing links. Test a permitted person and a denied person. Restricted content discovery (RCD) reduces broad discovery; it does not remove underlying access and has documented exceptions. Restricted access control (RAC) adds a group boundary: a person needs both an allowed-group membership and underlying resource permission. Membership alone does not grant access. Microsoft Search and site exclusions govern discovery; they are not permission repairs. For Copilot data readiness, remove excessive permission and risky links before relying on search restrictions. Check the current SharePoint Advanced Management entitlement and scope for the chosen control.

### Worked case

Requirement: only Finance should access a site. Mia has a sharing link but is not in a RAC allowed group. Ravi is in that group but has no site permission.

Result: neither condition set is complete. Mia fails the group boundary; Ravi lacks underlying permission. Give Ravi only the required site permission if approved. Do not add Mia merely to make her link work.

Requirement: every Sales meeting must preserve a transcript for later Copilot use. A default that organisers can change does not enforce this requirement.

## Choose the right control

| Underlying permission | RAC allowed group | Result when RAC applies           |
|-----------------------|-------------------|-----------------------------------|
| Yes                   | Yes               | Access permitted by this boundary |
| Yes                   | No                | Denied                            |
| No                    | Yes               | Denied                            |
| No                    | No                | Denied                            |

## Practical evidence task

1. Inspect mailbox delegates and record opening rights and sending rights separately. Do not send mail in a read-only exercise.
2. In a lab site, compare permissions and discovery controls; test both allowed and denied identities after propagation.
3. Record the organiser policy mode, whether the organiser can override it and whether the transcript was actually created.

Use synthetic data and the official lab instructions in an appropriate training environment. This guide supplies reasoning exercises, not a live lab tenant. Reading a task does not count as completing it.

## Microsoft references and procedure details

- [7 restricted access control](#)
- [2 create manage mailboxes](#)
- [7 configure copilot teams meetings](#)

## 03. Identity lifecycle and least privilege

An identity is an account or principal used to act. Authentication proves who it is; authorisation decides what it can do. An owner manages an object, a role permits administration, and resource membership grants access. Keep these relationships separate.

### Users, guests, contacts and groups

Create users with the correct account details, location and service entitlement. Guest invitation creates a collaboration identity; cross-tenant settings govern trust and inbound or outbound collaboration. Invitation acceptance does not automatically grant a site. A contact is a directory entry for communication, not a sign-in account. Microsoft 365 groups support collaboration resources; security groups support access assignments. Dynamic membership is evaluated from a rule, while assigned membership is explicitly maintained. Choose the group type needed by the target service. For departures, block the relevant sign-in, address sessions and permissions, and preserve business data according to the approved lifecycle. Deleting immediately can break recovery and ownership.

### Roles, PIM and administrative units

Choose the least privileged role that can perform the task and check its scope. An administrative unit (AU) scopes supported administration of member objects; it does not create a SharePoint data boundary. PIM eligible assignment is a right to activate a privileged role when required, rather than permanent active administration. Check activation requirements, approval, duration and audit. Test with the intended administrator, because an existing broad role can hide a scoping error. Assigning a team owner is not equivalent to giving a tenant-wide admin role.

### Bulk work with Graph PowerShell

Connect to the intended tenant with approved scopes and the correct delegated or application permission model. Select only required properties and retrieve all pages where necessary. Inspect a sample before changes. For a CSV operation, validate headers, unique identifiers, blanks and duplicates; target object IDs rather than display names. Record one result per item. A successful script exit does not prove every row succeeded. Handle service throttling using the documented retry guidance; do not repeatedly replay a partially completed batch without checking results. Separate read-only inventory, proposed changes and the execution step.

### Worked case

A regional support lead must manage only users in Region West. The lead already has a tenant-wide role.

First remove the misleading test condition in a controlled lab: use an administrator whose only relevant assignment is the scoped role. Put intended users in the AU and validate a permitted West action and a denied East action.

A 40-row import reports 38 successful updates and 2 duplicate identifiers. Record 38 successes and repair the two rows. Do not report 100% completion because the process ended.

## Choose the right control

| Requirement                  | Control                        | Proof                                   |
|------------------------------|--------------------------------|-----------------------------------------|
| Temporary privileged task    | PIM activation                 | Active duration and audit               |
| Regional user administration | Supported AU-scoped role       | Allowed and denied admin action         |
| External communication entry | Contact                        | Directory record without sign-in        |
| External collaboration       | Guest plus resource permission | Guest can access only intended resource |

## Practical evidence task

1. Use Graph PowerShell to read five users with Id, DisplayName and UserPrincipalName. Confirm the tenant and granted read scope before running.
2. Review a fictional import containing duplicate IDs and missing locations. Produce a corrected proposal and an item-by-item result log.
3. Demonstrate the difference between eligible and active roles in an appropriate PIM lab; record unsupported or unavailable features.

Use synthetic data and the official lab instructions in an appropriate training environment. This guide supplies reasoning exercises, not a live lab tenant. Reading a task does not count as completing it.

## Microsoft references and procedure details

- [provision govern identities entra](#)
- [3 bulk identity management graph](#)
- [9 govern roles role groups privileged identity management](#)

## 04. Authentication and Conditional Access

Registration, authentication and policy enforcement are separate. A registered method might not be allowed for the user or suitable for the requested authentication strength. A missing MFA prompt does not establish that MFA was skipped.

### Choose the identity control

Authentication-method policies control allowed methods and targeted users. Password Protection addresses weak or banned passwords; it is not an MFA policy. SSPR lets eligible users recover their account using configured methods. Check SSPR scope, method requirements and registration. Hybrid password writeback has additional configuration and licensing requirements; cloud reset success does not prove an on-premises password changed. ID Protection adds risk evidence and risk-based decisions; distinguish user risk from the risk of a particular sign-in. A reset, stronger authentication and access block are different responses.

### Read a sign-in in the right order

Identify user, resource, client, timestamp and final result in Entra sign-in logs. Review Authentication Details and each Conditional Access policy result. Then inspect scope, exclusions, conditions and state. Report-only records a potential result without enforcing it. What If helps predict policy applicability, but a real event is better evidence of what occurred. An existing token claim can satisfy MFA without a new prompt. Allow time for logs to finish aggregating before declaring missing evidence.

### Combine policies correctly

All applicable policies must be satisfied. A blocking policy wins over a policy that would permit access after MFA. A policy excluded from the relevant client category is not applied. Legacy authentication cannot fulfil modern MFA or device requirements; an applicable requirement can therefore block the request. Pilot changes with carefully scoped test identities and preserve a tested emergency access approach. Do not resolve an unexplained denial by excluding everyone from the policy.

### Worked case

Fixture: sign-in succeeds; policy P1 requires MFA and succeeds; Authentication Details says the requirement was satisfied by a token claim; P2 is report-only and would require a compliant device.

Verdict: the event supports MFA satisfaction. P2 did not enforce device compliance. A screenshot showing no prompt does not reverse either finding.

New fixture: P1 succeeds, but an applicable P3 blocks access. The request fails. Combining policies is not a vote or a priority contest.

## Choose the right control

| Evidence                    | What it proves                      | What it does not prove     |
|-----------------------------|-------------------------------------|----------------------------|
| Report-only success/failure | Evaluated simulated policy outcome  | Enforced restriction       |
| MFA satisfied by claim      | Existing claim met this requirement | A fresh challenge occurred |
| Not applied                 | This policy did not apply           | No other policy applied    |
| Applicable block            | Request is blocked                  | MFA can override it        |

## Practical evidence task

1. Choose one lab sign-in. Record UTC time, user, resource, client, every policy result and authentication details.
2. Use What If for an allowed and a denied scenario. Compare predictions with fresh sign-ins when authorised.
3. Explain the final result in three sentences: which policies applied, how requirements were met and why the outcome followed.

Use synthetic data and the official lab instructions in an appropriate training environment. This guide supplies reasoning exercises, not a live lab tenant. Reading a task does not count as completing it.

## Microsoft references and procedure details

- 6 evaluate conditional access client apps

## 05. Defender policy and incident decisions

A policy specifies protection. A rule scopes who receives it. An alert signals something to investigate; an incident relates evidence. A delivery event, a user click and a compromise are different stages.

### Choose protection and effective policy

Anti-spam handles unwanted mail; anti-phishing addresses phishing and impersonation; anti-malware addresses malicious attachments. Safe Links examines supported links, and Safe Attachments evaluates supported attachments. For recipient-based precedence, Strict preset precedes Standard, followed by applicable evaluation policies, custom policies and built-in/default protection. Within custom policies, a lower priority number wins. The applicable policy for a feature is used; settings are not merged across all matches. Check recipient inclusions and exclusions before comparing precedence. Custom priority zero does not outrank an applicable Strict preset.

### Investigate before closing

In Defender, inspect the alert's affected user and message, relevant timestamps and related evidence. Trace receipt, protection action, delivery location, later remediation, user action and related account activity. Use message identity and recipient evidence to avoid confusing two similar subjects. Containment may include supported removal, quarantine or account protection actions according to the case. Check action status and remaining exposure. A submitted remediation request is not proof that all affected items were removed. Automated investigation helps collect and propose actions; review its evidence and approvals. Keep false-positive conclusions separate from lack of visibility.

### Run simulations as training

Attack simulation training measures behaviour in an authorised exercise. Define audience, scenario and training assignment before launch. Review delivery, interaction, completion and repeat risk with the correct denominator and reporting period. A simulated click indicates training need; it is not evidence of a real compromise. A lower reported click rate can reflect a different audience or incomplete delivery. Verify those factors before claiming an improvement.

### Worked case

A mailbox matches Strict and a custom anti-phishing policy with priority 0. The custom policy permits a sender that Strict treats differently.

Decision: first verify recipient scope. If Strict applies to the recipient and feature, its settings govern. Editing the custom priority cannot make it win.

An alert says a malicious message was delivered. A later removal action is pending. The case remains open for exposure and completion checks; pending is not remediated.

## Choose the right control

| Question                | Evidence to inspect                                                  |
|-------------------------|----------------------------------------------------------------------|
| Which settings applied? | Feature, recipient scope, exceptions, precedence                     |
| Was mail delivered?     | Message event and delivery location                                  |
| Is cleanup complete?    | Action status and remaining affected items                           |
| Was training effective? | Delivered audience, behaviour and completion over comparable periods |

## Practical evidence task

1. Use a Defender training scenario to build a five-event timeline. Label observed evidence and inferred risk separately.
2. Resolve a policy precedence fixture with one applicable preset, two customs and one exception.
3. Write a close-out record containing action status, remaining exposure, owner and next check.

Use synthetic data and the official lab instructions in an appropriate training environment. This guide supplies reasoning exercises, not a live lab tenant. Reading a task does not count as completing it.

## Microsoft references and procedure details

- [secure email collaboration defender office365](#)
- [2 preset custom policy precedence](#)

## 06. Purview: classify, prevent and retain

Classification identifies data. Protection controls permitted use. DLP applies rules to risky activity. Retention manages how long content is kept or deleted. Choosing one does not automatically achieve the others.

### Classification and publication

A sensitive information type (SIT) detects a defined data pattern with supporting evidence and confidence. A sensitivity label classifies content and may apply configured protection. A label policy publishes labels to selected users and governs available labelling behaviour. Creating a label does not make it appear for every user. If a label is missing, check publication scope, supported client and propagation before changing its protection. If a document is accessible unexpectedly, examine permissions and the actual label settings; a label name alone proves neither encryption nor restricted access.

### Specify DLP precisely

Write location, detection condition and supported action. Exchange concerns mail flow; SharePoint and OneDrive concern files and sharing; Teams concerns messages; Endpoint DLP needs onboarded devices. Do not substitute a Teams message rule for a file stored in SharePoint. For Copilot, choose a custom, Copilot-only policy. Distinguish restricting labelled grounding content, blocking supported sensitive prompt text and blocking a sensitive web search. Simulation records matches without blocking. The prompt-processing action is preview and tenant-dependent; verify support. Directly typed text and uploaded file content are not the same detection input.

### Respond to alerts and monitor AI posture

Inspect policy, rule, location, matched evidence, user or agent, activity time and action. Establish whether the policy was in simulation or enforcement at that time. A match is not a demonstrated leak, and a block is not proof that every other channel is safe. In DSPM, use supported AI activity and findings to identify sensitive-data exposure and prioritise investigation. Correlate with audit and the relevant workload. Record coverage gaps, owners and follow-up instead of treating an empty dashboard as universal safety.

### Retain content for the required lifecycle

A retention policy applies lifecycle settings across scoped locations. A retention label attaches lifecycle settings to an item, including supported record behaviour when configured. A retention label policy makes labels available or applies them through supported methods. Identify required duration, start event, location and disposition process. Preservation and deletion can interact with other policies and holds; test the effective outcome with the appropriate specialist. Sensitivity labels describe handling and protection; retention labels describe lifecycle. Backup is a recovery service, not a substitute for a legal preservation requirement.

### Worked case

Requirement: stop financial account numbers typed into a supported Copilot prompt, while retaining Exchange coverage. Use a separate Copilot policy and Exchange policy with suitable conditions and actions.

A simulation event reports the prompt match and the user still receives a response. That is consistent with simulation. Verify tenant support, enforcement state and propagation before

expecting a block.

Requirement: classify contracts as Confidential and keep them for seven years. Specify protection and retention separately; the word Confidential does not define seven-year retention.

## Choose the right control

| Requirement                          | Primary control                    |
|--------------------------------------|------------------------------------|
| Detect account-number patterns       | SIT                                |
| Publish Confidential to Finance      | Sensitivity label policy           |
| Prevent risky supported activity     | DLP location, condition and action |
| Keep an item for a defined lifecycle | Retention label/settings           |
| Assess supported AI exposure         | DSPM findings plus investigation   |

## Practical evidence task

1. Draft three DLP requirements: external email, external file sharing and Copilot prompt text. Name location, condition, action and prerequisite.
2. Review a simulated alert: identify whether the outcome matches policy state. Use synthetic sensitive data in any authorised test.
3. Produce a separate sensitivity and retention design for one document class.

Use synthetic data and the official lab instructions in an appropriate training environment. This guide supplies reasoning exercises, not a live lab tenant. Reading a task does not count as completing it.

## Microsoft references and procedure details

- [prevent data loss lifecycle purview](#)
- [2 specify data loss prevention requirements](#)

## 07. Copilot readiness and user settings

Copilot readiness has four gates: entitlement, supported application experience, connectivity and governed data. A missing button and an unsafe answer are different problems. Diagnose the gate before changing settings.

### Pilot the actual user experience

Select users with a real task and verify prerequisites for that Copilot experience. Check product licence, service availability, app version, signed-in account and applicable settings. Inspect excessive sharing in the data the task will use. Copilot can expose discoverable information that the user already has permission to access; repairing permissions matters. Pilot with both a permitted and a denied user. A successful answer proves only that task, identity and moment. Compare expected sources with grounding evidence and document failures.

### Match tenant settings to the requirement

Use the Copilot settings experience for relevant tenant options and the applicable policy service for user settings. Separate self-service purchasing, admin-centre Copilot, release preference, AI disclaimer, video/image generation, search experience and web search. A disclaimer informs users; it is not a data-access control. Microsoft 365 Copilot Search concerns enterprise discovery; web search introduces external grounding. A product setting should be checked in the requested experience and for the requested audience, not inferred from another Copilot surface. Follow current supported values and role requirements rather than memorising a menu position.

### Manage Cowork and providers

Cowork access uses current spending-policy controls selecting Cowork, with product prerequisites. The old Cowork registry entry is not an access switch. Web-search behaviour is controlled separately. Test fresh requests after propagation and distinguish user-targeted controls from device settings. For an external AI provider, evaluate the specific offering, legal processing relationship, region commitments, data flow and approved tasks. A provider brand or a tenant toggle does not prove every offering has the same privacy terms. Document approval and test the actual permitted experience.

### Worked case

One licensed pilot user can use Copilot in the web experience but cannot see it in the desktop app. Another can use the app but sees overly broad project data.

These are different gates. For the first, inspect client, account and experience prerequisites. For the second, investigate resource permissions and sharing. Buying extra licences does not fix either observed issue by itself.

The business approves internal-only answers. A disclaimer is insufficient. Use the relevant web-search control and test the supported experience; review source evidence rather than judging by how an answer sounds.

## Choose the right control

| Symptom                           | First useful check                              |
|-----------------------------------|-------------------------------------------------|
| Only one app lacks Copilot        | Client/account and experience prerequisites     |
| Unexpected internal document      | Underlying permissions and discovery            |
| External grounding is prohibited  | Applicable web-search policy                    |
| Cowork access differs by audience | Current spending-policy scope and prerequisites |
| Provider processing is uncertain  | Terms for the specific offering                 |

## Practical evidence task

1. Create a readiness checklist for two users and two applications. Keep evidence per identity and experience.
2. Review a fictional overshared site and propose the minimum permission repair, then define allowed and denied tests.
3. Record a tenant setting, its audience, propagation expectation and fresh test evidence.

Use synthetic data and the official lab instructions in an appropriate training environment. This guide supplies reasoning exercises, not a live lab tenant. Reading a task does not count as completing it.

### Microsoft references and procedure details

- [enable manage microsoft 365 copilot](#)
- [configure copilot settings extensibility](#)
- [4a manage copilot cowork](#)
- [7 manage third party ai providers](#)

## 08. Connectors and enterprise search

A connector brings another source into the search or AI experience. Content retrieval, identity mapping and access control must all work. A connected source can be available technically while still failing the user's permission requirement.

### Select the retrieval model

A synced connector indexes source content and access information for later retrieval. A federated connector retrieves supported data at query time rather than maintaining the same content index. Its runtime call can still transfer data. Choose by freshness, source capabilities, permission model and permitted data flow. Do not assume that federated means no exposure or that synced means every employee can read everything. A search exclusion changes discovery; it does not revoke direct source access.

### Validate source permissions end to end

Check the connection identity, required source permissions and mapping from source identities to Microsoft 365 identities. Verify crawl status, indexed items and expected access rules. A successful crawl proves ingestion, not permission correctness. Test one permitted person and one denied person using the same source item. For synced connector ACL changes, follow the required full-crawl process; an incremental content crawl is not proof that updated access rules were reprocessed. A failed connection can leave older indexed content available, so deletion or cleanup may require a separate action. For federated retrieval, test the runtime authorisation and returned content under the intended user context.

### Troubleshoot relevance without weakening access

When results are missing, trace source availability, retrieval model, mapping, crawl/runtime result, permissions and discovery settings. Check query and supported result type after access prerequisites. Do not make the source public as a relevance test. When results are stale, check indexing or runtime cache behaviour and last successful refresh. Record the source item, identity, query and timestamp so another administrator can reproduce the issue.

### Worked case

A synced CRM connector crawled successfully yesterday. Today a source ACL removes Lena. An incremental content crawl succeeds, but Lena still gets the old result.

Decision: verify the source ACL and mapping, perform the supported full crawl for ACL changes and test Lena again after processing. Do not conclude that removing the source permission instantly updated the index.

A federated source is queried live. The architecture avoids a persistent content index, but the response still enters the supported AI workflow. Include that transfer in the approval.

## Choose the right control

| Evidence                 | Conclusion                                     |
|--------------------------|------------------------------------------------|
| Crawl success            | Ingestion ran; access still needs testing      |
| Permitted user gets item | Positive access case works                     |
| Denied user gets item    | Boundary defect requiring investigation        |
| No content index         | Does not mean no runtime data transfer         |
| Connection failed        | Does not prove old indexed results disappeared |

## Practical evidence task

1. Design a synced-versus-federated decision for a source with hourly changes and strict ACLs.
2. Write a permission-test matrix with source ID, mapped identity, expected result and actual result.
3. Simulate a removed source permission: describe the required refresh and verification, without claiming an unperformed live test.

Use synthetic data and the official lab instructions in an appropriate training environment. This guide supplies reasoning exercises, not a live lab tenant. Reading a task does not count as completing it.

## Microsoft references and procedure details

- [configure copilot settings extensibility](#)
- [9 configure copilot connectors](#)

## 09. Agent identity, access and registry

An agent is an actor that can retrieve information or take supported actions. Approving its listing, assigning its identity and granting it resource access are separate decisions. Identify the actor at each step.

### Know the object model

An Entra agent identity blueprint defines the agent application pattern; its blueprint principal represents it in the tenant. Child agent identities represent runtime agents. Some agents also have a paired agent user for supported user-account scenarios. A human sponsor or owner is a different object. Record object IDs and relationships. Do not treat a paired user as identical to the runtime agent identity. Manage ownership and lifecycle so departures do not leave orphaned agents. Inventory, approve, operate, review and retire with resource grants and credentials accounted for.

### Govern resource access

Access packages manage approved business-resource access with supported approval, duration and review controls. They are not Agent 365 security templates. In Entra, find the agent identity, open its Access packages view, inspect the assigned package and enabled policy, review actual resources and roles, and save only the approved changes. Test the intended resource and an out-of-scope resource. Check support and entitlement for agent identity governance in the current tenant. A package assignment is evidence of configuration, not proof that every API call is correctly authorised.

### Apply Conditional Access to the actual subject

Blueprint-targeted agent policies can scope supported child agent identities; paired user accounts require their own appropriate targeting. Agent identity policies use supported block controls rather than an interactive human MFA challenge. Verify fresh token requests after propagation. An API-key call that bypasses Entra authentication is outside that token policy path. Record authentication method, subject, target resource and policy result when assessing coverage. A green user sign-in test is not evidence of agent-token enforcement.

### Review agents and their tools

In the registry, inspect publisher, purpose, owner, requested capabilities, data access and status. Review requests, publish or reject, then control installation and user access as required. Set allowed types, sharing and template defaults deliberately. An approved agent is not automatically approved for every tool. Tool registry and requests support separate review of available capabilities, publishers and block status. MCP provides an interface to tools and data; it does not establish trust or least privilege. Custom upload or BYO tool support requires checking current format, preview status and onboarding prerequisites. Reassess agents after a capability or ownership change.

### Worked case

A procurement agent has an approved listing, a paired user and a runtime identity. Its package permits the Procurement site only. A new tool requests access to all projects.

Decision: approval of the listing does not justify the new grant. Review the tool separately, identify the token subject and narrow resource permissions. Test Procurement success and HR denial.

A blueprint block stops new child-agent token requests, but a paired user can still sign in. This does not prove the blueprint block failed; the two subjects require separate evidence and scope.

## Choose the right control

| Object/control          | Job                                              |
|-------------------------|--------------------------------------------------|
| Blueprint and principal | Agent pattern and tenant representation          |
| Child agent identity    | Runtime actor                                    |
| Paired agent user       | Separate supported user-account actor            |
| Access package          | Govern business-resource access                  |
| Registry approval       | Approve availability and deployment decision     |
| Tool review             | Govern a specific capability and its permissions |

## Practical evidence task

1. Use the current official access-package exercise. Record agent ID, package, enabled policy, resource roles, owner, review and expiry requirements.
2. Map blueprint, child identity, paired user and sponsor using explicit IDs. Write separate Conditional Access tests for relevant subjects.
3. Review a fictional agent request with one excessive tool grant. Produce a justified publish, reject or return-for-correction decision.

Use synthetic data and the official lab instructions in an appropriate training environment. This guide supplies reasoning exercises, not a live lab tenant. Reading a task does not count as completing it.

## Microsoft references and procedure details

- [establish agent identities entra agent id](#)
- [curate control agents registry tools](#)
- [exercise set up permissions](#)
- [2 understand agent id object model](#)
- [7 apply conditional access agent licensing](#)
- [manage tools](#)

## 10. AI monitoring, cost and exam decisions

The right report depends on the question. Usage is not business value. An audit event is not an adoption trend. A spending alert is not necessarily a service stop. Match platform, population, time and measurement before drawing conclusions.

### Choose the evidence source

Use Microsoft 365 admin usage reports for eligible user and workload activity. Use supported Copilot dashboard insights for adoption patterns in the available population. Use Purview Audit for an individual supported activity; a Copilot interaction record does not guarantee verbatim prompt text. Use Copilot Studio analytics for its environment and agents. Use Azure monitoring and Cost Management for the relevant Foundry resource. Agent 365 overview and registry views help assess the agent estate. Copilot Control System helps bring administration, usage and service health views together; check scope and freshness rather than assuming every AI platform is in one report.

### Investigate agent governance gaps

Correlate owner, identity, tools, permissions, activity, sensitive-data findings and action status. If an agent accessed a protected resource, inspect the actual principal and token path, source permission and supported data controls. Do not declare compliance from registration alone. Record the gap, affected resources, risk evidence, accountable owner and corrective test. A missing event may reflect unsupported activity, reporting delay or insufficient visibility. State the limitation instead of inventing a safe outcome. Service health should match the failing service, geography and time; compare a product incident with individual request failures.

### Reconcile cost before optimising

Separate assigned licences, prepaid/capacity funding and metered pay-as-you-go. For Microsoft 365 spending policies, inspect per-user and total limits, targeted group and period. Where multiple policies apply, current precedence chooses the highest per-user limit, then highest overall limit, then the newest policy. Moving a user does not erase already consumed usage in the same period. Supported enforced limits can stop covered service access; Azure budgets normally notify and do not stop the resource. Reports refresh at different times. Compare the same UTC period, currency, resource and meter. Copilot credits or tokens are consumption units, not currency. Avoid counting a charge twice when it appears in both Microsoft 365 and Azure reporting.

### Measure adoption honestly

Define eligible population, active users, workload, window and refresh time. Fictional example: 60 active users out of 100 eligible users is 60% adoption; 60 out of 80 licensed users is 75% of licensed users. Both calculations are valid only with their stated denominator. A dashboard's estimated time saving is a modelled measure, not proof of causal return on investment. Compare similar groups and workload opportunities before recommending removal or expansion. Optimise unwanted consumption and poor task outcomes, rather than simply maximising activity.

### Worked case

Finance sees a £120 PAYG total in Microsoft 365 and the same tagged £120 resource charge in Azure for the same day. Treat them as two views until charge identity is reconciled, not automatically £240.

An Azure budget emails at £100. The resource continues running. That is consistent with a notification budget. Choose a supported enforcement mechanism if the requirement is to stop covered use.

An agent is registered but has no owner and an unreviewed external tool. Registration does not close those governance gaps. Assign accountability, review access and test the corrected boundary.

## Choose the right control

| Business question                    | Useful evidence                                |
|--------------------------------------|------------------------------------------------|
| Who used Copilot in Word?            | Scoped workload usage report                   |
| What did this identity do?           | Supported audit event                          |
| What did this Foundry resource cost? | Azure resource/meter cost                      |
| Is this agent governed?              | Owner, grants, tools and activity evidence     |
| Why did the service fail?            | Matching health incident plus request evidence |

## Practical evidence task

1. Reconcile two fictional reports with matching period, currency and charge identity. Mark unmatched lines for investigation.
2. Calculate adoption with eligible and licensed denominators separately; explain why the figures differ.
3. Complete the two mixed tests on separate days. For every error or correct guess, record the mistaken assumption, corrected rule and a new example.

Use synthetic data and the official lab instructions in an appropriate training environment. This guide supplies reasoning exercises, not a live lab tenant. Reading a task does not count as completing it.

## Microsoft references and procedure details

- [agent 365 govern](#)
- [monitor optimize m365 ai services](#)
- [2 pick right report source](#)
- [6 reconcile copilot agent costs](#)

## A repeatable exam method

1. **Underline the outcome:** prevent access, reduce discovery, preserve content, recover content, enable a feature or explain an event?
2. **Name the object:** tenant, person, group, message, site, client, agent identity, paired user, tool or billing resource?
3. **Read scope and state:** who is targeted, who is excluded, which policy applies, and is it simulated, enforced or still processing?

- 4. Eliminate by mechanism:** a valid product name is insufficient. Ask whether it can perform the requested action in that location.
- 5. Use the evidence:** configuration proves intent; an event proves an observed outcome. Check timestamps and action status.
- 6. Choose the smallest supported action that meets every stated constraint.** Do not add broad permissions just to make a test succeed.

## Critical distinctions to explain without notes

| Pair                                 | Deciding rule                                      |
|--------------------------------------|----------------------------------------------------|
| Licence / role                       | Use entitlement / administration permission        |
| Verification / MX change             | Domain control / mail routing                      |
| RCD / RAC                            | Discovery reduction / additional access boundary   |
| Eligible / active PIM role           | May activate / currently exercising privilege      |
| Report-only / enforced               | Observe potential outcome / apply controls         |
| MFA claim / fresh prompt             | Requirement satisfied / new challenge displayed    |
| Preset / custom priority             | Applicable preset precedes custom                  |
| Sensitivity / retention              | Handling and protection / lifecycle                |
| DLP match / blocked action           | Detected event / enforced outcome                  |
| Agent identity / paired user         | Different subjects and policy scope                |
| Registry / access package            | Availability approval / resource-access governance |
| Synced / federated                   | Indexed retrieval / query-time retrieval           |
| Usage / value                        | Activity count / demonstrated business outcome     |
| Budget notification / enforced limit | Alert / supported stopping behaviour               |

## Repair an error in one minute

Write: “I chose \_\_ because I assumed \_\_. The deciding rule is \_\_. The evidence needed is \_\_. In a new case where \_\_ changes, the answer becomes \_\_.” Restudy the smallest relevant section. Do not reread the entire guide for one scope mistake.

## Before booking or sitting the exam

Check the current official study guide and exam page for objectives and delivery details. Microsoft requires a scaled score of at least 700; this is not a claim that 70% of these questions is sufficient. This short guide covers the current objective families, but exact procedures, entitlements and preview availability require the linked current documentation and hands-on experience.

[Official objectives](#) • [Exam details](#) • [Official lab collection](#)

## Scope, sources and checks

Version 3.0, 7 October 2026. The exam study guide, exam page and relevant Microsoft training units were checked during revision. Module references provide detailed procedures. The mapping CSV retains all 58 distinct local objective pointers from the previous course and routes their 64 mappings to this edition. Broad coverage is not a claim that each objective has equal depth or a live exercise here. Use primary Microsoft material to extend unfamiliar or lightly covered topics.

Company names, examples, amounts and test cases are fictional. No tenant was accessed or changed to validate these tasks. This package supplies no verified learner pass-rate data. Product availability, licensing, portal wording and preview features can change.

The PDF is the short reading guide. The HTML adds automatic practice, answer keys and progress history. The ZIP contains those files, a question-bank CSV, objective mapping, a practical-evidence worksheet, an error log and revision notes. Prior repetitive questions and long lesson text are removed from the active course.