

AB-650

Plain-English Mastery Course

Microsoft 365 and AI services administration

Prepared for Anir Ghosal • Self-paced • Version 2.0

Updated 7 October 2026

13 modules • 60 mapped lessons • worked examples • practical exercises

Companion HTML: 360 practice activities, including 60 new mixed-check questions and automatic marking for 240 selected-response items.

95% is the course practice target. It is not a guaranteed exam score. Independent training; no Microsoft endorsement or verified market ranking is claimed.

Reading map

Start with the foundation bridge if a term is unfamiliar. Read each module introduction before its lessons. Use the HTML practice desk after each module. Page entries below are clickable.

AB-650 Beginner to enterprise practitioner	5
Foundation bridge - 3 hours	8
Curriculum and learning outcomes	9
Student guide - modules and lessons	11
M01 - Tenant foundations and licensing	12
M02 - Operational resilience and funding	15
M03 - Collaboration and safe discovery	19
M04 - Identity administration at scale	23
M05 - Authentication and access decisions	28
M06 - Threat protection and response	32
M07 - Information protection and governance	36
M08 - Copilot readiness and rollout	41
M09 - Copilot controls and integrations	45
M10 - Agent identity and access	49
M11 - Agent registry and tool governance	52
M12 - Agent security and compliance operations	56
M13 - AI economics, adoption and service health	59
Lab guide - 13 bounded exercises	62
Objective-to-lesson traceability	71
Integrated enterprise case studies	74
Instructor and facilitator guide	76
Slide deck outline and recording plan	77
Revision guide and exam priorities	83
30-day study plan - 60 hours	85
Flashcards - 60 prompts	87
Glossary - 60 working terms	96

Your route to 95% practice mastery

Use this as a learning target, not an exam prediction. First attempts matter more than scores earned after memorising the answers.

1. Understand: study one module's plain-English explanation, then its mapped lessons. Explain the control aloud without reading.

2. Practise: answer that module's reinforcement questions. Study the explanation for every mistake and every correct guess.

3. Apply: complete the corresponding fixture exercise, then the live Microsoft lab when you have an appropriate tenant. Record simulation and live work separately.

4. Check: attempt Mixed check A without notes. The local target is at least 29/30 (96.7%) overall and at least 90% in each domain.

5. Repair: return to the linked lesson for every error. Write the mistaken rule, corrected rule and one new example.

6. Recheck: attempt Mixed check B on a later day without first reading its answers. Apply the same target. Repeated attempts are practice, not unseen evidence.

7. Demonstrate: explain access versus discovery, policy scope versus enforcement, and user versus agent identity correctly; complete the lab evidence checklist. A high quiz score cannot cancel a critical misconception.

These small checks are uncalibrated and not replicas of the real exam. They are weighted 7 tenant, 12 security and 11 AI items each (23.3%, 40%, 36.7%), within the published domain ranges. They sample the course rather than exhaust every exam possibility.

A sustainable session

Allow about 15 minutes for the explanation, 30-45 minutes for lessons, 20-30 minutes for exercises and 15 minutes for retrieval practice. Take longer when you cannot explain why an alternative is wrong. The full 30-day schedule remains below.

AB-650 | Beginner to enterprise practitioner

Independent preparation • Version 2.0 • Updated 7 October 2026 • Self-paced with optional facilitated workshops

Prepared for Anir Ghosal. Concept-first explanations, Microsoft administration examples and guided evidence collection precede commands. Your Microsoft infrastructure experience is useful, but this course assumes no prior knowledge of the newer AI and agent administration features.

Read this first

This is original, independent training material, not an official Microsoft course, exam dump or endorsed product. It contains 60 mapped lessons across 13 exam modules, a foundation bridge, practical work, 300 original assessment activities, facilitator materials and a 30-day plan. The assessment bank deliberately includes related retrieval-practice sets; it is not five independent or statistically calibrated mock exams. Matching activities are printable drag-and-drop equivalents. The package does not contain recorded videos, an LMS/SCORM package or a hosted tenant.

Plan around 60 hours over 30 study days, with extra time for unfamiliar concepts and lab access. This is a study estimate. Aim for at least 95% on this course's practice checks, plus satisfactory practical evidence. This target is not a validated predictor of a Microsoft result, a cohort pass-rate claim or an exam-score guarantee.

Verified exam baseline and limitations

Item	Research finding
Title and status	AB-650, Administering Microsoft 365 and AI Services; the retrieved official listing still says beta.
Domain weights	Tenant/workload administration 20-25%; security/governance 40-45%; Microsoft 365 AI administration 35-40%.
Passing score	700 on the Microsoft scale; this is not a statement that 70% correct passes.
Prerequisite knowledge	Microsoft 365 workloads and Entra administration, Defender XDR awareness, and familiarity with Microsoft Graph PowerShell.
Official practice assessment	The exam page says it is not currently available. This original bank is not that assessment.
Language and beta results	The retrieved listing specifies English. Microsoft states beta results are not immediate.
Duration and item count	Do not assume a fixed AB-650 count or duration from generic exam ranges. Confirm the appointment details and exam introduction.
Documentation access	Microsoft describes in-exam Learn access for eligible associate/expert exams, with exclusions and no added time. Follow the actual exam interface and current rules.
Training availability	Three directly verified Learn paths contain 5, 5 and 7 modules. The exam/certification landing-page course widget appears incomplete; use the direct links below.
Version maintenance	The study guide shows 27 July 2026 as its update date. Product controls can move faster. Recheck blueprint, product behavior and licensing before the exam or a live lab.

Official exam: <https://learn.microsoft.com/en-us/credentials/certifications/exams/ab-650/>

Blueprint:

<https://learn.microsoft.com/en-us/credentials/certifications/resources/study-guides/ab-650>

Certification:

<https://learn.microsoft.com/en-us/credentials/certifications/ai-services-administrator-associate/>

Exam rules:

<https://learn.microsoft.com/en-us/credentials/support/exam-duration-exam-experience>

How to study each lesson

- Read the concept aloud and explain it using the fictional Northstar Consulting tenant.
- Predict the scenario outcome before reading the recommended action.
- Explain why each tempting alternative fails the requirement.
- Complete the practical exercise and record evidence, not just a screenshot of a saved toggle.
- Use the linked primary learning module for expanded product walkthroughs; return to any weak prerequisite.

- Attempt the matching and scenario questions without notes; inspect answers only after committing.
- Maintain an error log: lesson, mistaken assumption, corrected rule, evidence and retest date. Revisit on days 2, 7 and 14 after the initial attempt.

Lab access, costs and evidence standards

Use an isolated authorized training tenant and synthetic content. No tenant is provisioned by this package, and a free developer or trial tenant is not guaranteed. Basic administration requires the relevant Microsoft 365 workload entitlement. Advanced Purview, PIM, Defender, Copilot, Agent 365 and SharePoint controls can require additional plans, rollout eligibility or billing. Verify each feature in the current product documentation and tenant before buying anything. A read role may show a page without allowing a change.

There are two routes. Live route: perform the bounded exercise with required permissions and entitlements. Offline route: use the supplied fixtures and complete the same decision worksheet. Label each evidence item LIVE, OFFLINE or NOT TESTED. Offline reasoning is useful preparation but does not establish hands-on proficiency. Do not purchase subscriptions, enable meters, send campaigns, enforce tenant-wide policies or change production merely to complete a lab. Optional live actions need the tenant owner's ordinary authorization.

For every lab record: requirement; resource; identity and role; license/feature check; baseline; intended change; positive test; negative test; observed outcome; propagation or uncertainty; rollback and cleanup. A role or feature described as unavailable is a finding, not permission to substitute Global Administrator. Cleanup must respect existing retention and records controls.

Foundation bridge - 3 hours

F01, 45 minutes: tenant, subscription, identity, group, role and resource. Outcome: explain why a user can sign in but cannot open a document. Exercise: sketch five objects and label the relationships. Answer: authentication proves identity, while resource authorization still depends on effective permissions. Mapping: prerequisite for D1.T1, D2.I1 and D2.A1; no claim of a separate official domain.

F02, 45 minutes: cloud workloads and data paths. Outcome: distinguish Exchange mail, Teams conversations, SharePoint shared files and OneDrive personal work files. Exercise: classify a shared mailbox, project library, private-channel file and personal draft. Ask where access is actually enforced rather than choosing by the application icon. Mapping: prerequisite for D1.W1 and D2.P1.

F03, 45 minutes: AI vocabulary. A model generates outputs; grounding supplies relevant source information; an agent combines instructions, context and tools to pursue a task; a connector links a source; a tool exposes an operation. Exercise: trace a request to summarize a document versus send its summary. The second operation adds a destination and an action requiring its own authorization. Mapping: prerequisite for D3.C1, D3.I1 and D3.G1.

F04, 45 minutes: evidence and commands. Read a command as verb, resource, parameters and target. Distinguish configuration, logs and observed behavior. Exercise: explain every line of the read-only script later in this guide before running it. Mapping: prerequisite for D2.I1.7 and all verification exercises. Foundation gate: correctly explain access versus authentication, licensing versus roles and configuration versus outcome, then identify one appropriate positive and negative test.

Curriculum and learning outcomes

Module	Hours	Lessons	Outcome
M01 Tenant foundations and licensing	3	L01, L02, L03, L04	Select, justify and verify the controls in this module against a business requirement.
M02 Operational resilience and funding	3	L05, L06, L07, L08	Select, justify and verify the controls in this module against a business requirement.
M03 Collaboration and safe discovery	4	L09, L10, L11, L12, L13	Select, justify and verify the controls in this module against a business requirement.
M04 Identity administration at scale	5	L14, L15, L16, L17, L18, L19, L20	Select, justify and verify the controls in this module against a business requirement.
M05 Authentication and access decisions	5	L21, L22, L23, L24, L25, L26	Select, justify and verify the controls in this module against a business requirement.
M06 Threat protection and response	4	L27, L28, L29, L30, L31	Select, justify and verify the controls in this module against a business requirement.
M07 Information protection and governance	6	L32, L33, L34, L35, L36, L37, L38	Select, justify and verify the controls in this module against a business requirement.
M08 Copilot readiness and rollout	4	L39, L40, L41, L42	Select, justify and verify the controls in this module against a business requirement.
M09 Copilot controls and integrations	4	L43, L44, L45, L46, L47	Select, justify and verify the controls in this module against a business requirement.
M10 Agent identity and access	4	L48, L49, L50, L51	Select, justify and verify the controls in this module against a business requirement.
M11 Agent registry and tool governance	4	L52, L53, L54, L55	Select, justify and verify the controls in this module against a business requirement.
M12 Agent security and compliance operations	3	L56, L57, L58	Select, justify and verify the controls in this module against a business requirement.

Module	Hours	Lessons	Outcome
M13 AI economics, adoption and service health	3	L59, L60	Select, justify and verify the controls in this module against a business requirement.

Hours total: foundation 3 + modules 52 + final revision/capstone assessment 5 = 60. Module time includes reading, exercises, related questions and remediation. Optional hybrid workshops are additional. The 300-item distribution is 65 tenant, 125 security and 110 AI activities: 21.7%, 41.7% and 36.7%, within the published domain ranges. This distribution is a course design decision, not a prediction of the actual exam mix.

Student guide - modules and lessons

Current controls that are easy to confuse

Checked against the linked Microsoft units on 7 October 2026. Availability, licence requirements and portal labels can change. Follow the current unit for exact setup.

Decision	Rule to understand	Reference
Copilot DLP	The Copilot location uses a custom policy and cannot share that policy with other locations. Supported label-based grounding restrictions, prompt-text controls and web-search restrictions have different effects. Prompt processing restrictions are preview and tenant-dependent; they do not scan an uploaded file's contents as directly entered prompt text. Simulation does not block.	DLP requirements
Cowork	Access uses spending policies selecting Cowork, with required product prerequisites. Web-search policy is separate. The legacy Cowork registry entry does not control access.	Cowork controls
Agent Conditional Access	Target the actual subject. Blueprint scope can cover current and future child agent identities after propagation. It does not automatically cover paired agent user accounts. Agent identities use the supported block control; a human MFA challenge is not their remediation path. API-key access bypasses Entra authentication.	Agent policy scope
RCD and RAC	RCD limits broad discovery with documented exceptions; it does not revoke direct access. RAC requires allowed-group membership and underlying permission. Test direct access separately from search results.	SharePoint controls
Connectors	Synced sources use indexed content and access information. Federated sources retrieve at query time. No indexing does not mean no transfer. Verify the relevant identity mapping and permitted/denied source outcomes.	Connector models
Third-party AI providers	Assess the specific offering. Independent processors and Microsoft subprocessors have different terms. Do not infer region commitments or required roles from the provider's brand alone.	Provider review

M01 - Tenant foundations and licensing

Start with the organisation, then give people the right services

In simple English

A tenant is your organisation's Microsoft 365 environment. It contains accounts, settings and subscriptions. A domain is a name your organisation controls, such as northstar.example. A licence gives someone an entitlement to use a product. A role gives someone permission to administer something. These are separate ideas. Buying a licence does not make a person an administrator. Giving someone an administrator role does not buy the product.

How it works

First identify the tenant and the business requirement. Then identify the object you need to change: the organisation profile, a domain, a user, or a licence assignment. Choose the role needed for that task. For a custom domain, follow Microsoft's ownership verification process before using the domain. Keep verification separate from changing mail delivery. For group-based licensing, a group acts as an assignment list: eligible members receive the selected product licences. Assignment can fail, so membership is only the starting point of your check.

Worked example

Fictional example: Northstar adds 20 staff. All 20 join the licensing group, but only 18 receive the expected service. Do not record 100% success because the group has 20 members. Check each user's assignment result, available licences, usage location and service-plan conflicts. Success means the intended users have the intended service, not just a membership entry.

Use this process

1. Write the requirement in one sentence: who needs which service?
2. Record the current domain, assignment and role settings relevant to the task.
3. Make the smallest change in the training scope.
4. Check the assignment result and test as the intended user.
5. Record unresolved errors separately from completed work.

Keep this distinction: Entitlement answers 'May this user use the product?' Administrative permission answers 'May this person change the product's settings?' Neither answer proves that the user's client is ready.

Stop and answer: A domain is verified, but mail must stay with the old provider for one more week. Should you change the MX record now?

Check your reasoning

No. Keep the approved mail route until the planned cutover. Verification proves control of the domain; changing MX affects delivery.

Read with the mapped lessons below. [Microsoft reference](#)

Estimated study: 3 hours. Lab prerequisites: Tenant or Entra settings administrator; license administrator for entitlement work. Use a separate test tenant.

Think of a tenant as your organization's rented office building. Identities are the people recognized by reception, domains are the addresses people use, services are the rooms where work happens,

and licenses determine which facilities are available. This analogy has a limit: Microsoft 365 permissions are evaluated by several connected services, not by one receptionist. A successful sign-in does not mean every room is available.

Start with a requirement and follow its dependency chain. If the requirement is a new email domain, you need ownership verification, account and service configuration, and a deliberate routing decision. If the requirement is a recognizable login, investigate the branding surface. If a user lacks a capability, distinguish a missing product assignment from a disabled service plan. The wrong control can appear to save successfully while leaving the requirement unmet.

For troubleshooting, compare an affected user with a working user. Record entitlement sources, direct group membership and service-plan state. Avoid inferring success from a group-level message. A user moved between licensing groups needs the destination assignment verified before the previous source is removed. This is an example of preserving service continuity while changing configuration.

Key concepts: Branding, Domain verification, Organization settings, Group licensing.

Diagram to create: Create a two-layer tenant map: identity and administration above workload resources. Mark domain verification separately from mail-routing arrows.

Knowledge check with answer: Why can a verified domain still deliver mail elsewhere? Answer: verification proves ownership; routing follows the configured mail records.

Common mistakes: Confusing appearance with trust, a group assignment with every user entitlement, and a saved value with effective behavior.

Exam tip: Underline the requested outcome before choosing a portal. Domain, branding, entitlement and access are different requirements.

L01 - Branding

Objective pointer: D1.T1.1. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: Branding gives users recognizable organization and sign-in experiences; it does not establish trust in a message or authorize access.

Northstar Consulting scenario: A newly acquired division needs its logo, company URL, theme and sign-in background reviewed before launch.

Recommended approach: Review Microsoft 365 organization branding and Entra sign-in branding separately, then test both experiences.

- Alternative to reject: Change only the email signature. Why: Signatures do not configure either tenant branding experience.
- Alternative to reject: Use the logo as evidence that a sign-in is genuine. Why: Attackers can reproduce visual branding.
- Alternative to reject: Assign a security group to the theme. Why: Group membership is not the branding configuration.

Verification: Capture the organization experience and a fresh sign-in view using the pilot identity, checking the intended URLs and imagery.

Practical exercise: Draw two browser frames and mark which portal controls each visual element.

L02 - Domain verification

Objective pointer: D1.T1.2. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: Domain ownership verification and mail routing are separate decisions. Proving control of a namespace does not itself move email.

Northstar Consulting scenario: A company owns a new domain but must keep its existing mail provider until a scheduled migration.

Recommended approach: Add the domain and publish the supplied verification record; defer mail-routing changes until the agreed cutover.

- Alternative to reject: Change MX records immediately. Why: This can move mail before the service is ready.
- Alternative to reject: Rename every account first. Why: Verification should precede use of an unverified namespace.
- Alternative to reject: Create a SharePoint site with the domain name. Why: A site name does not prove DNS ownership.

Verification: Check verified domain status and independently confirm the production MX destination still matches the approved pre-cutover configuration.

Practical exercise: Prepare a DNS change ticket separating verification, routing, rollback and validation.

L03 - Organization settings

Objective pointer: D1.T1.3. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: A tenant contains organization-wide service choices and profile information; each setting needs a business owner and an understood scope.

Northstar Consulting scenario: HR requests a profile update while security requests a review of privacy and service defaults.

Recommended approach: Inventory the affected organization settings, identify owners, record current values and pilot supported scoped changes.

- Alternative to reject: Treat every organization setting as a per-device policy. Why: Tenant and device scopes differ.
- Alternative to reject: Use branding settings for privacy controls. Why: Appearance is not a privacy configuration.
- Alternative to reject: Give HR permanent Global Administrator. Why: The business request does not justify unrestricted administration.

Verification: Compare the saved values and effective experience against the approved requirements, including any propagation delay.

Practical exercise: Write a change record with old value, new value, scope, owner and recovery action.

L04 - Group licensing

Objective pointer: D1.T1.4. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: A license is an entitlement assembled from service plans. Group assignment must be verified for the intended users rather than inferred from membership somewhere in a hierarchy.

Northstar Consulting scenario: A licensed parent group contains 20 direct users and a nested group of 10 users. Only the direct users receive the product.

Recommended approach: Assign the entitlement to an appropriate group containing the intended direct user members and inspect assignment results.

- Alternative to reject: Wait for nested membership to inherit the license. Why: Nested groups are not processed by group-based licensing.
- Alternative to reject: Enable an unrelated service plan. Why: The missing assignment is a membership-scope issue.
- Alternative to reject: Treat group success as proof all 30 users are licensed. Why: Success covers processed users, not nested users.

Verification: Inspect direct membership, user-level entitlement sources, service-plan state and assignment errors for the ten affected accounts.

Practical exercise: Build a before-and-after table of 30 expected entitlements; simulate two insufficient-seat failures.

Primary expansion reading:

<https://learn.microsoft.com/en-us/training/modules/provision-brand-microsoft-365-tenant/>
<https://learn.microsoft.com/en-us/training/modules/license-microsoft-365-ai-services/>

M02 - Operational resilience and funding

Keep services working and know how to recover

In simple English

An outage is a service problem. A network problem is a connection problem. A recovery problem is an inability to restore the information you need. They can look similar to a user, but they require different evidence. Service health tells you about reported Microsoft service incidents.

Connectivity information helps you investigate the path between users and Microsoft 365. Backup prepares you to recover supported data.

How it works

Begin with scope. Is the problem affecting one person, one office, one application or many locations? Compare the timing with Service health. If only one office is affected, compare network results with an unaffected office. For recovery, define what data must be protected and how recent a restored copy must be. Confirm the backup configuration and practise a permitted restore. Retention and backup serve different requirements; a rule to keep records is not proof that a restore will meet a recovery objective.

Worked example

Fictional example: Teams works at home but repeatedly disconnects at one branch. There is no matching service incident. This is evidence to investigate the branch connection; it is not proof that Microsoft has no problem anywhere. Compare affected and unaffected paths, record times, and escalate with the evidence. Separately, a finance recovery exercise fails because the required resource was never protected. A successful backup of another resource cannot close that finding.

Use this process

1. Describe affected users, locations, services and first known failure time.
2. Check the relevant service incident and notifications.
3. Compare connectivity evidence instead of guessing from one symptom.
4. For backup, check protected scope and available recovery points.
5. Validate the restored content against the stated recovery requirement.

Keep this distinction: An incident notification is not a diagnosis of every user's problem. A configured backup is not a demonstrated restore. Pay-as-you-go funding is not the same as a user's product licence.

Stop and answer: A manager asks whether a backup is usable. Is a screenshot showing Enabled sufficient?

Check your reasoning

No. Check protection scope and recovery points, then use an authorised restore exercise to establish whether the required data can be recovered.

Read with the mapped lessons below. [Microsoft reference](#)

Estimated study: 3 hours. Lab prerequisites: Read roles for health and billing review; supported Backup role and approved billing for an optional live restore.

A service can fail at several layers. The cloud provider may have an incident; the local client may be outdated; a proxy may interfere; or the user may lack the necessary entitlement. Collect evidence that separates these hypotheses. A symptom that follows one location is a different clue from a symptom that follows one identity across locations.

Recovery starts with a business question: what must be recovered, from what time, and how will you know it is correct? A submitted restore is not the finish line. Validate the recovered content against an expected result. Agree who authorizes replacing current data and who confirms the business service works afterward.

AI administration also introduces consumption decisions. A role can allow someone to build something without paying for others to use it. Build a persona matrix before assigning products or enabling a meter. For calculations, keep the unit rate, consumption unit and period explicit. The numbers in this course are fictional teaching values, not Microsoft prices.

Key concepts: AI entitlements and funding, Backup and recovery, Network diagnosis, Service health.

Diagram to create: Draw a fault tree with client, identity, network and cloud-service branches. Add a separate recovery timeline with incident time and chosen restore point.

Knowledge check with answer: Why can a successful restore job still fail business acceptance?
Answer: it may restore the wrong point, wrong scope or incomplete expected content.

Common mistakes: Confusing funding with permission, advisory scope with every failure, and retention with recovery.

Exam tip: Use the evidence that isolates the fault. A broad reset is rarely the first justified response.

L05 - AI entitlements and funding

Objective pointer: D1.T1.5. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: Product access, permission to create agents and the mechanism that pays for consumption answer different questions.

Northstar Consulting scenario: Eight makers build agents for 80 occasional users. Procurement asks whether maker access alone funds all user consumption.

Recommended approach: Separate the maker roles, user product entitlements and supported consumption funding, then validate the required licensing combination.

- Alternative to reject: Assume a maker role pays for all usage. Why: An administrative or maker permission is not a billing arrangement.
- Alternative to reject: Assign every user an administrator role. Why: Administrative permission neither replaces licensing nor justifies broad access.
- Alternative to reject: Use an agent owner field as billing authorization. Why: Ownership and funding are separate controls.

Verification: Produce a matrix of persona, product, role, billing source and available capacity, verified against current subscription details.

Practical exercise: Calculate a fictional monthly bill using 5000 units at 0.02 currency units each plus 80 units of fixed cost; explain the 180 total.

L06 - Backup and recovery

Objective pointer: D1.T1.6. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: Backup prepares a recoverable copy or restore point for supported workloads; retention governs preservation and deletion obligations.

Northstar Consulting scenario: A test library was damaged yesterday. The recovery owner needs a defensible restoration rather than simply a retention setting.

Recommended approach: Confirm protected scope and a suitable restore point in Microsoft 365 Backup, choose the recovery option and validate recovered content.

- Alternative to reject: Apply a new retention label and assume yesterday is restored. Why: Retention configuration does not perform point-in-time recovery.
- Alternative to reject: Restore everything without identifying the target time. Why: An unnecessary broad restore can replace valid recent work.
- Alternative to reject: Close the incident when the restore is submitted. Why: Submission does not establish recovery success.

Verification: Compare restored test files with known expected contents, review job status and document missing or changed items.

Practical exercise: Design a restore drill with a synthetic file, recovery target, permitted destination and acceptance checks.

L07 - Network diagnosis

Objective pointer: D1.T1.7. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: Connectivity is a path from client through local network and security intermediaries to a cloud service. A healthy service can still be unreachable from one office.

Northstar Consulting scenario: Copilot works at headquarters but stalls in a branch behind a proxy; Service health shows no related advisory.

Recommended approach: Compare client connectivity at both locations and inspect the branch path, including supported Copilot traffic and WebSocket handling.

- Alternative to reject: Reassign all user licenses. Why: A location-specific failure is not established as a licensing defect.
- Alternative to reject: Declare a global Microsoft outage. Why: The observations isolate a branch difference.
- Alternative to reject: Disable all network inspection permanently. Why: This is broader than the evidence supports.

Verification: Correlate timestamps, connectivity results and proxy observations, then retest the same pilot operation after the targeted fix.

Practical exercise: Create a troubleshooting table separating client, identity, network and service evidence.

L08 - Service health

Objective pointer: D1.T1.8. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: Service health reports provider incidents and advisories; local symptom reports still need correlation before attribution.

Northstar Consulting scenario: Several users report a Teams problem during an advisory affecting only a different region.

Recommended approach: Compare the advisory's affected workload, scope and timing with the users' failures and configure relevant notifications.

- Alternative to reject: Attribute every Teams ticket to the advisory. Why: Different affected scope may explain none of these failures.
- Alternative to reject: Use the message center as definitive incident status. Why: Planned change communication and incident health serve different purposes.
- Alternative to reject: Ignore local diagnostics because an advisory exists. Why: An unrelated local cause remains possible.

Verification: Record matching or nonmatching workload, region, time window and incident details alongside pilot tests.

Practical exercise: Draft a status update distinguishing confirmed impact, suspected cause and next diagnostic step.

Primary expansion reading:

<https://learn.microsoft.com/en-us/training/modules/license-microsoft-365-ai-services/>

<https://learn.microsoft.com/en-us/training/modules/maintain-microsoft-365-tenant-health/>

M03 - Collaboration and safe discovery

Understand access before you change collaboration settings

In simple English

Exchange provides email. Teams brings people into conversations and meetings. SharePoint stores shared sites and documents. OneDrive stores a user's work files. These services work together, but their permissions are not one universal switch. Opening a shared mailbox and sending as that mailbox are different permissions. Finding a document and having permission to open it are also different.

How it works

For every collaboration request, identify the resource, the person and the required action. 'Finance needs access' is incomplete. 'Priya must read the shared mailbox and send as its address' is testable. In Teams, distinguish team membership from membership of a private channel. For meeting Copilot, check the applicable meeting policy and transcription behaviour. In SharePoint, inspect effective access, including group membership and sharing links. Use discovery controls for discovery requirements; use access controls for access requirements.

Worked example

Fictional example: a project file should remain available to its approved team but should not be broadly discovered. That calls for a discovery review. If a former team member must no longer open the file, hiding it from search does not meet the requirement. Test the saved direct URL as that person. For Restricted Access Control, the permitted group condition and the underlying SharePoint permission must both be satisfied.

Use this process

1. List the person, resource and allowed action.
2. Inspect the actual permission route, including existing links.
3. Check the collaboration policy at its own scope.
4. Test an allowed user and a deliberately denied test user.
5. Record direct access and search discovery as separate outcomes.

Keep this distinction: Full Access to a mailbox does not itself grant Send As. Restricted Content Discovery does not revoke direct access. Restricted Access Control does not itself grant the missing SharePoint permission.

Stop and answer: A user belongs to a site's RAC allowed group but has no underlying permission. Can the user open the site?

Check your reasoning

No. Both the allowed-group condition and underlying SharePoint permission are required.

Read with the mapped lessons below. [Microsoft reference](#)

Estimated study: 4 hours. Lab prerequisites: Exchange, Teams and SharePoint scoped administrators; supported Copilot and SharePoint Advanced Management entitlement for advanced live tests.

Collaboration services share identities but not every permission model. A mailbox delegate, a Teams owner and a SharePoint visitor have different capabilities. Begin with the resource and action: open, read, edit, send, share or administer. Then find the specific permission that authorizes that action. This prevents an apparently convenient broad role from becoming the default answer.

Search introduces another layer. A user can have direct access to a file without seeing it in a particular search result. Conversely, a stale search result is not proof that opening the item is still allowed. Keep access and discovery tests separate. For sensitive content, identify the grant that created unnecessary access and correct it; discovery restrictions can support an interim response but do not replace that work.

Meetings add temporal state. A policy may allow a transcript without anyone starting one. Test the organizer's option, the participant's eligibility and the actual saved artifact. In practical administration, a screenshot of a toggle is weaker evidence than the observed behavior under the right identity and meeting conditions.

Key concepts: Shared mailbox permissions, Teams membership, Meeting Copilot and transcripts, SharePoint permissions, Discovery versus access.

Diagram to create: Draw user → permission → resource, with a separate discovery branch. Make a table comparing mailbox read and send permissions.

Knowledge check with answer: Can hiding a site in broad search remove a saved link permission? Answer: no; test and repair the underlying access separately.

Common mistakes: Assuming channel membership is universal, Full Access includes Send As, or transcription allowed means a saved transcript exists.

Exam tip: Watch verbs: open, send, discover and retain imply different controls.

L09 - Shared mailbox permissions

Objective pointer: D1.W1.1. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: Opening a mailbox and sending with its identity require different delegated permissions.

Northstar Consulting scenario: A delegate can read the support mailbox but cannot send a reply that appears to come directly from support.

Recommended approach: Verify the intended sending behavior and grant Send As for that mailbox to the authorized delegate.

- Alternative to reject: Grant Full Access again. Why: Read access does not itself confer Send As.
- Alternative to reject: Choose Send on Behalf when the sender must appear solely as support. Why: That permission exposes the delegate-on-behalf relationship.

- Alternative to reject: Enable direct sign-in using a shared password. Why: Delegated access avoids a shared interactive credential.

Verification: Send a synthetic message through the delegate and inspect the recipient-visible sender, while confirming existing read access.

Practical exercise: Compare Full Access, Send As and Send on Behalf using three fictional helpdesk requirements.

L10 - Teams membership

Objective pointer: D1.W1.2. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: Teams owners manage the workspace; channel membership and associated content boundaries depend on the channel type.

Northstar Consulting scenario: A project has a team-wide channel and a private budget channel. A new team member must not see budget discussion.

Recommended approach: Verify the channel type and manage the private channel membership separately from general team membership.

- Alternative to reject: Make the new member a team owner to hide budget data. Why: Elevating administrative rights is not a restriction.
- Alternative to reject: Assume every channel has identical membership. Why: Private channels have distinct membership.
- Alternative to reject: Change the channel name to Confidential. Why: Naming does not enforce access.

Verification: Test discussion and associated file access with an included and an excluded pilot identity.

Practical exercise: Draw a team with standard, private and shared channels; mark the membership and file boundaries.

L11 - Meeting Copilot and transcripts

Objective pointer: D1.W1.3. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: Meeting policy, organizer options and the existence of a transcript jointly determine the experience. Allowing transcription is not evidence that it ran.

Northstar Consulting scenario: A meeting needs transcript-grounded Copilot help after the call, but nobody started transcription.

Recommended approach: Use a supported transcript-required configuration and ensure an authorized participant starts transcription during the test meeting.

- Alternative to reject: Assume allowing transcription creates a transcript. Why: An allowed capability is not a completed action.
- Alternative to reject: Use only live-meeting Copilot and promise a saved transcript. Why: Live assistance can operate without a saved transcript.

- Alternative to reject: Treat the default Off value as an immutable organizer ban. Why: The documented default can be overridden by the organizer.

Verification: Check effective meeting settings and the actual saved transcript and post-meeting experience with an eligible licensed participant.

Practical exercise: Prepare a four-row matrix for policy default, organizer override, transcript started and after-meeting availability.

L12 - SharePoint permissions

Objective pointer: D1.W1.4. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: Access is the result of effective permissions, group membership and sharing links, not the folder's reassuring name.

Northstar Consulting scenario: A finance site is called Private Finance but an all-staff group still has read access.

Recommended approach: Trace effective access and remove the unintended grant or link while preserving the finance team's required access.

- Alternative to reject: Exclude the site from search and declare the access fixed. Why: Discovery suppression does not remove permissions.
- Alternative to reject: Change the title to Highly Confidential. Why: A name does not alter authorization.
- Alternative to reject: Remove only a navigation link. Why: Users may still open saved URLs.

Verification: Use direct URLs with a finance identity and a nonfinance identity, recording an allow and a deny after the change.

Practical exercise: Build a permissions inventory with grant source, business owner and retain/remove decision.

L13 - Discovery versus access

Objective pointer: D1.W1.5. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: Restricted discovery reduces broad surfacing; access restrictions decide who can open content. These controls need different tests.

Northstar Consulting scenario: A site should stay accessible to its permitted users by direct link but should not be broadly surfaced through organization-wide Copilot discovery.

Recommended approach: Evaluate Restricted Content Discovery for the site and separately keep reviewing its underlying permissions.

- Alternative to reject: Remove every authorized user's permissions. Why: That violates the requirement to preserve direct access.
- Alternative to reject: Assume a missing search result proves access is denied. Why: Discovery and authorization are separate observations.
- Alternative to reject: Use a site restriction as proof OneDrive is identically configured. Why: Control support and scope differ by workload.

Verification: Record direct access and broad discovery separately; account for indexing delay and documented discovery exceptions.

Practical exercise: Compare an interim discovery restriction with permanent permission cleanup using two acceptance criteria.

Primary expansion reading:

<https://learn.microsoft.com/en-us/training/modules/configure-exchange-teams/>

<https://learn.microsoft.com/en-us/training/modules/govern-sharepoint-onedrive-copilot/>

M04 - Identity administration at scale

Give each identity only the access it needs

In simple English

An identity represents a person or another actor in a system. A user account can sign in. A guest supports external collaboration. A mail contact is an address-book entry for an external recipient; it is not an employee login. A group collects identities for a purpose. An administrator role grants management capabilities. Choose the object that fits the job before you create it.

How it works

Treat joiners, movers and leavers as a lifecycle. Create the right account, assign needed access, change it when responsibilities change, and remove access when it is no longer justified. Use Privileged Identity Management, or PIM, for supported temporary privileged access rather than assuming everyone needs a permanent role. An administrative unit can limit supported directory administration to selected objects. It does not create a separate Microsoft 365 tenant. Microsoft Graph PowerShell lets you perform supported directory tasks in bulk, but you still need the relevant permissions and a clear target list.

Worked example

Fictional example: a helpdesk team should manage only the Melbourne office's supported user tasks. An administrative unit with an appropriately scoped role fits this delegation question. It does not mean Melbourne documents are now isolated from all other staff. That is a separate data-permissions question. Before a bulk change to 100 users, inspect a small sample and the exact target list; a loop repeats mistakes just as efficiently as correct actions.

Use this process

1. Choose user, guest, contact or group from the business need.
2. Select a supported role and scope for the administrator's task.
3. For temporary privilege, check eligibility, activation and expiry.
4. For bulk work, inspect the input and preview the selected identities.
5. Verify individual outcomes, errors and the audit trail.

Keep this distinction: A group has a purpose; its name alone does not grant a role. PIM eligibility is not the same as an active assignment. Administrative units scope supported administration, not all resource access.

Stop and answer: An external accountant needs an address-book entry only. Must you create a licensed employee account?

Check your reasoning

No. A mail contact may meet the address-book requirement. If the person also needs to collaborate, evaluate guest access separately.

Read with the mapped lessons below. [Microsoft reference](#)

Estimated study: 5 hours. Lab prerequisites: Supported User, Groups or Exchange roles for the task; Privileged Role Administrator for PIM configuration where authorized; read-only Graph permissions for inventory.

Identity management is a lifecycle. A joiner needs correct attributes and scoped resources. A mover may need old access removed as well as new access added. A leaver requires both access revocation and continuity for the work they owned. These needs can conflict if deletion happens before ownership and preservation are assessed.

Delegation answers two questions: what operation can the administrator perform, and over which objects? Roles address capability; administrative-unit scope can limit supported directory administration. PIM adds another dimension: when the privileged role is active. None of these automatically grants or isolates all workload content.

PowerShell is useful when the same administrative operation must be repeated predictably. You do not need to become a software developer to read a command, inspect its target and understand its permission requirements. Start with reads. For writes, validate the target list, pilot a small sample, handle errors per record and reconcile expected versus actual changes. A command that continues after one error can leave a partial result requiring deliberate follow-up.

Key concepts: User lifecycle, Guests and external collaboration, Group purpose, Privileged roles and PIM, Administrative units, Contacts, Graph PowerShell bulk administration.

Diagram to create: Draw joiner → mover → leaver with access and ownership checks at each transition; add capability, scope and time as three role dimensions.

Knowledge check with answer: Why is a successful Graph connection not proof that 200 updates succeeded? Answer: it validates a session, not each target operation.

Common mistakes: Equating contacts with sign-in users, confusing role eligibility with activation, and treating administrative units as separate tenants.

Exam tip: For bulk questions, check input validity, permission scope, partial failure and verification.

L14 - User lifecycle

Objective pointer: D2.I1.1. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: A user identity needs controlled creation, change and retirement. Departure handling must consider sessions, data and business ownership.

Northstar Consulting scenario: An employee leaves today but their project documents must remain available to the approved successor.

Recommended approach: Apply the approved offboarding sequence for sign-in and sessions, preserve required data and transfer ownership before destructive cleanup.

- Alternative to reject: Delete the account immediately without checking dependencies. Why: That can disrupt continuity and retention obligations.

- Alternative to reject: Keep sign-in active until every file is transferred. Why: The departure access risk remains.
- Alternative to reject: Transfer the departed user's password. Why: Ownership transfer should not require sharing credentials.

Verification: Verify access revocation and successor access separately, with a record of preserved data and remaining dependencies.

Practical exercise: Write joiner, mover and leaver checklists for a fictional analyst.

L15 - Guests and external collaboration

Objective pointer: D2.I1.2. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: A guest can participate through a governed external identity; external chat capability is not equivalent to membership or file access.

Northstar Consulting scenario: A supplier needs one project site's files for six weeks but no general tenant access.

Recommended approach: Use a scoped guest collaboration arrangement with the resource owner's approval, an end-date review and the necessary resource permissions.

- Alternative to reject: Enable broad external chat and assume files are accessible. Why: Communication reachability does not grant site permissions.
- Alternative to reject: Share an employee account. Why: That removes individual accountability.
- Alternative to reject: Add the supplier to an all-company group. Why: That grants unnecessary scope.

Verification: Test the intended resource and a prohibited resource with the guest, and record the review date.

Practical exercise: Write a guest access request with sponsor, resource, duration and removal conditions.

L16 - Group purpose

Objective pointer: D2.I1.3. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: Security groups and Microsoft 365 groups support different management and collaboration needs. Membership changes can affect multiple connected resources.

Northstar Consulting scenario: A manager wants shared collaboration resources rather than only an access-control list.

Recommended approach: Choose a Microsoft 365 group where its connected collaboration resources satisfy the requirement, then assign accountable owners.

- Alternative to reject: Create only a mail contact. Why: A contact is not a collaborative membership container.
- Alternative to reject: Assume every security group provisions a shared workspace. Why: Security groups do not automatically provide the same collaboration resources.

- Alternative to reject: Let the group operate without an owner. Why: Membership and lifecycle decisions need accountability.

Verification: Inspect the group type, owners, membership and the expected connected resources.

Practical exercise: Classify four requests as collaboration, authorization, distribution or contact management.

L17 - Privileged roles and PIM

Objective pointer: D2.I1.4. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: Least privilege limits capability; time-bound activation reduces how long privileged capability remains available.

Northstar Consulting scenario: A support engineer needs elevated administration for an approved one-hour maintenance task.

Recommended approach: Use the least-privileged eligible role and the configured PIM activation workflow for the maintenance window.

- Alternative to reject: Assign permanent Global Administrator. Why: The task does not justify permanent tenant-wide power.
- Alternative to reject: Grant ownership of an unrelated team. Why: Team ownership does not provide the administrative role.
- Alternative to reject: Treat role eligibility as permanent activation. Why: Eligible and active are distinct states.

Verification: Inspect activation duration, approval or justification evidence and the role state after expiry.

Practical exercise: Write an activation request identifying task, role, time window, approver and validation.

L18 - Administrative units

Objective pointer: D2.I1.5. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: An administrative unit can scope supported delegated directory administration; it is not a separate tenant or a universal data boundary.

Northstar Consulting scenario: A regional helpdesk should administer supported user tasks for its region without gaining the same scope everywhere.

Recommended approach: Use an administrative unit and a supported role assignment scoped to that unit.

- Alternative to reject: Rename the region's users. Why: Naming alone does not scope a role.
- Alternative to reject: Use a tenant-wide role and ask staff to avoid other users. Why: A written request is not enforced scoping.
- Alternative to reject: Assume the unit isolates all SharePoint data. Why: Directory administration scope is not universal workload isolation.

Verification: Test a permitted operation on an in-scope test user and its rejection on an out-of-scope test user.

Practical exercise: Draw a tenant with two administrative units and annotate the supported delegated operations.

L19 - Contacts

Objective pointer: D2.I1.6. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: A mail contact represents an external recipient in the directory; it is not an interactive employee account.

Northstar Consulting scenario: Staff need to find a vendor's external email address in the address book, with no tenant sign-in.

Recommended approach: Create or manage a mail contact with the correct external address and display details.

- Alternative to reject: Create a privileged member account. Why: The request needs neither sign-in nor administration.
- Alternative to reject: Assign a Copilot license to the contact. Why: A directory contact is not an interactive licensed user.
- Alternative to reject: Invite a guest solely to populate this address entry. Why: That adds an identity relationship beyond the stated need.

Verification: Verify address-book resolution and the external delivery address without creating sign-in access.

Practical exercise: Compare a contact, guest and member using purpose, sign-in and resource access columns.

L20 - Graph PowerShell bulk administration

Objective pointer: D2.I1.7. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: Graph PowerShell exposes directory operations as commands. Safe bulk work uses explicit scope, validated input and per-record results.

Northstar Consulting scenario: A CSV contains duplicate user principal names and a bulk change is planned for 200 accounts.

Recommended approach: Validate and deduplicate the input, preview the exact target set, then pilot the authorized operation and log each result.

- Alternative to reject: Run the entire write immediately. Why: Invalid targets can cause partial or unintended changes.
- Alternative to reject: Treat a successful connection as proof all writes succeeded. Why: Authentication says nothing about each record.
- Alternative to reject: Request all directory write permissions for a read-only inventory. Why: The permission scope exceeds the operation.

Verification: Reconcile input targets with successful, failed and skipped records, and sample the resulting directory state.

Practical exercise: Use the supplied read-only Graph exercise and design a validation report for a proposed bulk write.

Primary expansion reading:

<https://learn.microsoft.com/en-us/training/modules/provision-govern-identities-entra/>

M05 - Authentication and access decisions

Work out why a sign-in succeeds or fails

In simple English

Authentication checks who is signing in. Authorisation decides what that identity can do.

Multifactor authentication, or MFA, asks for more than one type of proof. Self-service password reset, or SSPR, helps eligible users recover access. Conditional Access applies rules to sign-ins. A rule can require stronger proof, require a supported device condition, or block access.

How it works

Read a Conditional Access policy in three parts: who it applies to, which resource and conditions it covers, and what it requires. Then evaluate every applicable policy. An applicable block is not cancelled by a different policy that permits access after MFA. Report-only mode records what the policy would do; it does not enforce the block. A user who sees no new MFA prompt may already have satisfied MFA in an existing session. Check the sign-in event before diagnosing a bypass.

Worked example

Fictional example: one enforced policy requires MFA and another enforced policy blocks the same sign-in. The user completes MFA. Access is still blocked. The block policy applies independently. If the block policy were report-only instead, it would not enforce that block, although another policy might still deny the request. Your evidence must identify the actual policy state and the event, not just the names of two policies.

Use this process

1. Identify the user, resource, client and time of the sign-in.
2. Check authentication-method availability and registration where relevant.
3. Read each matching Conditional Access result and policy state.
4. Inspect authentication details, device information and the final result.
5. Use a controlled test after correcting the specific problem.

Keep this distinction: Allowed authentication method does not mean registered method. Registered method does not mean it was used in this event. Sign-in risk concerns a sign-in; user risk concerns possible account compromise.

Stop and answer: A report-only block policy shows Failure. Did that policy block the user?

Check your reasoning

No. It indicates that the policy would block if enforced. Check the final event and other policies to explain what actually happened.

Read with the mapped lessons below. [Microsoft reference](#)

Estimated study: 5 hours. Lab prerequisites: Authentication Administrator or appropriate policy role; Conditional Access Administrator for a dedicated sandbox; required Entra licensing.

Authentication asks who is presenting themselves. Authorization asks what that identity may do. A stronger authentication method helps establish identity, but it does not automatically give the person access to a particular site. Keep these questions separate when interpreting a support ticket.

For policy troubleshooting, establish applicability before outcome. Which identity, resource and client are involved? Is the policy enabled or report-only? Did its conditions match? An applicable policy that requires an additional condition must be satisfied along with the other applicable policies. A block cannot be canceled by a successful MFA step elsewhere.

Logs are stronger than impressions. No new prompt may mean an earlier authentication claim satisfied the requirement, not that the control failed. Record the specific event and its authentication details. When piloting policy changes, use supported report-only analysis first, preserve a tested recovery route and limit the affected population. The aim is evidence of the intended decision, including a negative test, rather than merely a successful login.

Key concepts: Authentication methods, Password Protection, Self-service password reset, Sign-in investigation, Conditional Access interaction, Risk-based access.

Diagram to create: Draw a policy decision tree: in scope? conditions match? report-only or enforced? applicable block? required controls satisfied?

Knowledge check with answer: What does report-only failure establish? Answer: the policy evaluation indicates a failure under its logic, but report-only did not enforce blocking.

Common mistakes: Treating enablement as registration, no prompt as no MFA, or the newest Conditional Access policy as the winner.

Exam tip: Build a small truth table. Evaluate every applicable policy before predicting the final outcome.

L21 - Authentication methods

Objective pointer: D2.A1.1. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: Authentication establishes identity. Method availability, registration and actual use must be distinguished.

Northstar Consulting scenario: A pilot group is allowed a stronger sign-in method, but several users have never registered it.

Recommended approach: Verify method policy scope and user registration, then run a controlled sign-in using the intended method.

- Alternative to reject: Assume policy enablement registers everyone. Why: Users may still need registration.
- Alternative to reject: Reset every password as proof of stronger authentication. Why: A reset does not demonstrate the desired method.
- Alternative to reject: Use a sensitivity label to require the method. Why: Content classification does not configure sign-in methods.

Verification: Inspect the effective method policy, registration status and authentication details for the pilot sign-in.

Practical exercise: Build a matrix of allowed, registered and actually used methods for three fictional users.

L22 - Password Protection

Objective pointer: D2.A1.2. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: Password Protection reduces weak password choices using prohibited-password evaluation; it does not replace multifactor authentication.

Northstar Consulting scenario: Users choose passwords based on the company name. Security wants that predictable pattern rejected.

Recommended approach: Evaluate the custom banned password list and Password Protection configuration, then test controlled synthetic password changes.

- Alternative to reject: Use an email anti-spam policy. Why: It does not evaluate password selection.
- Alternative to reject: Disable MFA after adding banned words. Why: Password quality and additional authentication protect different risks.
- Alternative to reject: Publish a password guideline only. Why: A guideline alone does not enforce rejection.

Verification: Record that prohibited test choices are rejected and acceptable synthetic choices work, without logging passwords.

Practical exercise: Write a password-policy test plan using placeholders rather than real credentials.

L23 - Self-service password reset

Objective pointer: D2.A1.3. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: SSPR requires an eligible, in-scope user with sufficient usable registered methods. Configuration and readiness are separate.

Northstar Consulting scenario: A pilot user is in the SSPR group but cannot complete reset because the required recovery methods are missing.

Recommended approach: Check SSPR scope, licensing and registration requirements, then resolve the missing usable method registration.

- Alternative to reject: Add the user to a Teams channel. Why: Channel membership is unrelated to SSPR readiness.
- Alternative to reject: Disable all recovery requirements. Why: This removes the intended assurance rather than fixing registration.
- Alternative to reject: Assume group membership alone guarantees reset. Why: The user must also satisfy the configured method requirements.

Verification: Complete a controlled reset for the test identity and inspect the reset or audit evidence.

Practical exercise: Create a troubleshooting decision tree for out-of-scope, unregistered and method-unavailable users.

L24 - Sign-in investigation

Objective pointer: D2.A1.4. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: A missing MFA prompt does not establish a bypass; an existing authentication claim can satisfy the requirement.

Northstar Consulting scenario: A user signs in without a new prompt and the helpdesk claims Conditional Access failed.

Recommended approach: Inspect the specific sign-in's policy evaluation, authentication details and session context before changing controls.

- Alternative to reject: Disable the policy because no prompt appeared. Why: Prompt appearance alone is insufficient evidence.
- Alternative to reject: Treat the user's successful desktop login as the cloud sign-in record. Why: Different events may have different evaluations.
- Alternative to reject: Create another identical MFA policy immediately. Why: Duplication can obscure the diagnosis.

Verification: Capture the relevant sign-in ID and explain each applicable policy's result and how authentication was satisfied.

Practical exercise: Annotate a synthetic log: MFA satisfied by prior claim, policy applied, final result success.

L25 - Conditional Access interaction

Objective pointer: D2.A1.5. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: Applicable Conditional Access requirements are cumulative; an applicable block is not canceled by another policy's grant.

Northstar Consulting scenario: One applicable policy requires MFA; a second applicable policy blocks the user's access to the same resource.

Recommended approach: Predict blocked access and use the sign-in evidence to verify both policy evaluations.

- Alternative to reject: Assume passing MFA overrides the block. Why: A grant does not negate an applicable block.
- Alternative to reject: Apply only the most recently edited policy. Why: Policies are not evaluated as a last-edit-wins list.
- Alternative to reject: Ignore resource targeting. Why: Applicability depends on assignments and conditions.

Verification: Inspect the final blocked result and the contribution of each applicable policy for the same token request.

Practical exercise: Solve a four-row policy matrix with user, resource, conditions and grant or block outcomes.

L26 - Risk-based access

Objective pointer: D2.A1.5. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: Sign-in risk concerns a particular authentication event; user risk concerns potential compromise of the identity. Remediation must match the signal and supported controls.

Northstar Consulting scenario: Security sees a suspicious sign-in and separately an account marked at risk; the administrator needs a scoped response.

Recommended approach: Distinguish the risk signals, check supported licensing and policy conditions, and validate the intended response in a pilot.

- Alternative to reject: Treat every risk signal as the same event. Why: User and sign-in risk describe different assessments.
- Alternative to reject: Use an unrelated DLP alert as proof of sign-in risk. Why: The signal must be verified in the identity evidence.
- Alternative to reject: Require an agent identity to complete a human MFA prompt. Why: Nonhuman identities cannot satisfy interactive remediation.

Verification: Record the risk type, matched policy, intended action and observed test outcome without assuming every risk means confirmed compromise.

Practical exercise: Write separate handling plans for a risky sign-in and a potentially compromised user.

Primary expansion reading:

<https://learn.microsoft.com/en-us/training/modules/implement-authentication-access-entra/>

M06 - Threat protection and response

Investigate threats with a timeline, not a guess

In simple English

Microsoft Defender for Office 365 helps protect email and collaboration. A threat policy defines protection behaviour. An alert says something needs investigation. An incident can bring related evidence together. An attack simulation is a controlled learning exercise, not proof that a real attack happened. These are connected tools, but each answers a different question.

How it works

Start by understanding which protection applies to the recipient and workload. Policy scope and precedence matter: a policy that exists but does not apply cannot explain protection for that person. When an alert appears, inspect the message, recipients, delivery outcome and related events. Separate 'message delivered', 'user clicked' and 'account compromised'. Each claim needs evidence. Choose response actions that match the finding and your response authority. Check the result before closing the incident.

Worked example

Fictional example: a harmful email reaches six mailboxes. One recipient clicks its link. This supports investigating that recipient and the message's impact; it does not establish that all six passwords were stolen. Remove or contain confirmed harmful content through the appropriate response process, investigate the clicked user's sign-ins, and record unknowns. A training campaign later tests the behaviour you want to improve, such as reporting suspicious messages.

Use this process

1. Determine the affected recipients and effective protection settings.
2. Build a timeline from delivery, click and identity evidence.
3. State what is confirmed, suspected and still unknown.
4. Select and verify proportionate remediation.
5. Design follow-up training around the observed behaviour.

Keep this distinction: An alert is a lead, not the whole investigation. Deleting a message does not establish that identity risk has been resolved. A simulation completion count does not prove safer behaviour by itself.

Stop and answer: A link click is confirmed. Is account takeover confirmed too?

Check your reasoning

No. Investigate additional evidence. The click is important, but it does not prove successful credential theft or use.

Read with the mapped lessons below. [Microsoft reference](#)

Estimated study: 4 hours. Lab prerequisites: Appropriate Defender security read or response role; Defender for Office 365 Plan 2 or equivalent entitlement for advanced investigation and simulation features.

Security operations move from signal to scope to action. An alert is a starting point, not a complete narrative. Ask what happened, what evidence supports it, which identities or resources are involved and what remains unknown. Record the action and the outcome separately so another analyst can reproduce the decision.

Threat policies have their own evaluation behavior. A priority number is meaningful within the relevant policy type, not as a universal ordering of all security controls. A setting effective for spam is not necessarily the control required for impersonation. Examine feature, policy source, recipients and settings together.

Human learning is part of prevention. A simulation should have a defined population, approval, proportionate content and useful follow-up. Avoid turning a learning exercise into a public ranking of staff. Improvement is better assessed through reporting and repeated behavior than through a single delivery or click count alone. The fictional exercises here use no real malicious links or credentials.

Key concepts: Alert triage, Threat policy scope, Threat policy precedence, Threat investigation and response, Attack simulation training.

Diagram to create: Draw an incident evidence chain from alert to message to affected identity to remediation result. Add a separate per-feature policy ladder.

Knowledge check with answer: Why does closing an alert not prove containment? Answer: status may change without removing the message or addressing compromised access.

Common mistakes: Assuming custom priority zero beats a preset, every recipient was compromised, or simulation delivery proves learning.

Exam tip: Choose the response proportional to the verified scope; preserve the distinction between suspicion and evidence.

L27 - Alert triage

Objective pointer: D2.S1.1. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: An alert is a signal requiring assessment; its status is not the same as proof that the underlying threat was removed.

Northstar Consulting scenario: An email alert was marked resolved, but the analyst has not checked whether the delivered messages remain in mailboxes.

Recommended approach: Verify the evidence, affected recipients and remediation outcome before treating the threat as contained.

- Alternative to reject: Use the resolved status as proof of message removal. Why: Workflow status and remediation are different facts.
- Alternative to reject: Delete all recent email. Why: That is disproportionate and destroys legitimate content.
- Alternative to reject: Reassign a license to close the alert. Why: Licensing does not resolve the threat evidence.

Verification: Link the alert to message or incident evidence and the recorded remediation result.

Practical exercise: Prepare an incident note with signal, scope, action, outcome and unresolved questions.

L28 - Threat policy scope

Objective pointer: D2.S1.2. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: Threat-protection settings must match the feature and recipients. A policy's name does not prove effective coverage.

Northstar Consulting scenario: Finance needs protection against impersonation, but the administrator changed only a spam threshold.

Recommended approach: Review the anti-phishing controls relevant to impersonation and their effective recipient scope.

- Alternative to reject: Increase backup frequency. Why: Backup does not stop the impersonation message.
- Alternative to reject: Assume a spam threshold configures every phishing control. Why: Different protection features have different settings.
- Alternative to reject: Change the display name of the finance group. Why: A name does not enable threat protection.

Verification: Confirm the applicable protection feature, protected identities or domains where relevant and effective recipients.

Practical exercise: Build a threat-to-control table covering malicious links, attachments, impersonation and spam.

L29 - Threat policy precedence

Objective pointer: D2.S1.2. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: Custom policy priority orders comparable custom policies; applicable preset protection can take precedence for a feature.

Northstar Consulting scenario: A priority-zero custom anti-spam policy seems ignored for a user included in Standard preset protection.

Recommended approach: Inspect feature-specific precedence and recipient inclusion in the preset before changing custom priority.

- Alternative to reject: Set custom priority to zero again. Why: It already has that priority and does not override the preset.
- Alternative to reject: Merge every policy's settings mentally. Why: Lower settings are not simply blended into the winning policy.
- Alternative to reject: Assume one global winner applies identically to every feature. Why: Precedence is evaluated by protection feature.

Verification: Identify the actual policy source for the feature and explain why it wins for that recipient.

Practical exercise: Write a policy-evaluation worksheet for two recipients with different preset membership.

L30 - Threat investigation and response

Objective pointer: D2.S1.3. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: Investigation connects evidence to affected assets; response must record what action was taken and whether it worked.

Northstar Consulting scenario: A phishing message reached multiple users and one user reports clicking its link.

Recommended approach: Investigate message and related identity evidence, scope affected users, and carry out the authorized targeted response.

- Alternative to reject: Remove only the original alert. Why: The underlying exposure remains.
- Alternative to reject: Assume every recipient's device is compromised. Why: Delivery is not proof of compromise.
- Alternative to reject: Send the suspicious message to all staff for inspection. Why: That spreads potentially harmful content.

Verification: Correlate message identifiers, recipient evidence, click or sign-in signals where available and the result of each remediation action.

Practical exercise: Build an incident timeline using fictional delivery, click, sign-in and remediation times.

L31 - Attack simulation training

Objective pointer: D2.S1.4. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: A simulation is a controlled learning activity with approved scope and follow-up; it is not permission to send uncontrolled phishing.

Northstar Consulting scenario: The organization wants to improve reporting behavior after a phishing incident.

Recommended approach: Plan an approved simulation for a bounded population with appropriate training and measurable follow-up.

- Alternative to reject: Use real stolen credentials in the exercise. Why: Training does not require real credential compromise.
- Alternative to reject: Measure success only by messages delivered. Why: Delivery does not show learning or reporting improvement.
- Alternative to reject: Send an unapproved simulation to all customers. Why: That exceeds the training population and authorization.

Verification: Compare reporting and repeat-behavior measures across the planned exercise and follow-up, noting cohort differences.

Practical exercise: Draft a simulation plan with audience, owner, scenario, communications and learning measures.

Primary expansion reading:

<https://learn.microsoft.com/en-us/training/modules/secure-email-collaboration-defender-office365/>

M07 - Information protection and governance

Give each data-protection control one clear job

In simple English

Classification identifies sensitive content. A sensitivity label describes a classification and can apply configured protection. A label policy makes labels available to selected people. Data loss prevention, or DLP, controls supported handling of sensitive content. Retention governs how long content is kept or when it is deleted. These controls complement one another; they do not all do the same job.

How it works

Write a data-protection requirement as location, condition and action. Where is the risk: email, a file, a Teams message, a device or a supported Copilot interaction? What should match: a sensitive information type, a label or another supported condition? What must happen: audit, notify or restrict a supported action? Confirm that the selected workload supports the combination. Simulation records matches without enforcing the action. Publishing a sensitivity label does not automatically label every existing document.

Worked example

Fictional example: finance must keep a record for an approved period and prevent inappropriate sharing. Retention addresses the keep/delete requirement. Permissions, configured label protection and applicable DLP address access and handling requirements. One control is not evidence that the other requirement is met. Use synthetic matching content and a near-match to check detection; then inspect the exact event and policy mode.

Use this process

1. State the data, location, intended recipients and business requirement.
2. Choose classification, sensitivity protection, DLP or retention by purpose.
3. Check supported conditions, actions and feature entitlement.
4. Test a true match and a near-match in the training scope.
5. Investigate alerts and document confirmed enforcement separately from simulation.

Keep this distinction: A sensitivity label is not necessarily encryption: its configuration matters. Retention is not DLP. A Copilot policy is not universal protection for every agent and every upload.

Stop and answer: A DLP simulation reports a match. Can you claim the data transfer was blocked?

Check your reasoning

No. A match is detection evidence. Simulation does not enforce the blocking action.

Read with the mapped lessons below. [Microsoft reference](#)

Estimated study: 6 hours. Lab prerequisites: Purview read roles for inspection; appropriate DLP, information-protection and lifecycle roles for isolated policy work; feature-specific licensing.

Start with what the business needs to protect and what must happen. Classification identifies meaningful content. A label can communicate its sensitivity and apply configured protection. DLP governs supported actions involving sensitive content. Retention governs how long content must remain and how disposal is handled. Backup supports recovery. One product family may expose several of these, but their purposes are not interchangeable.

Detection quality matters. A broad pattern may find both real identifiers and harmless lookalikes. A known roster suggests a different approach from free-form narrative documents. Test positives and negatives. If a policy catches everything, it may block necessary work and train users to ignore warnings; if it catches too little, the apparent simplicity hides risk.

For each workload, specify location, condition and action. Do not project an endpoint action onto a Teams message or assume a Copilot rule covers every agent channel. Treat preview capabilities and licensing as explicit prerequisites. When reviewing alerts, establish whether an action was attempted, allowed or blocked. When reviewing posture findings, assign owners and validate remediation rather than assuming the dashboard repaired the issue.

Key concepts: Classification techniques, Sensitivity labels and publishing, DLP requirements, Copilot DLP boundaries, Retention design, DLP alert response, DSPM for AI.

Diagram to create: Draw five labeled boxes: classify, protect, restrict action, preserve, recover. Connect each to a different business requirement.

Knowledge check with answer: Why is a DLP alert not automatically proof of a leak? Answer: it may describe a prevented attempt; inspect the action outcome.

Common mistakes: Confusing sensitivity with retention, label creation with publication, a citation with processing, and posture visibility with remediation.

Exam tip: Translate the scenario into data shape, location, action and time. Those four clues usually narrow the control.

L32 - Classification techniques

Objective pointer: D2.P1.2. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: Pattern detection, matching known records and recognizing document meaning address different classification problems.

Northstar Consulting scenario: A control must match actual customer identifiers from a maintained roster rather than every string shaped like an identifier.

Recommended approach: Evaluate exact data match against the approved reference dataset and test both known and unknown synthetic values.

- Alternative to reject: Use only a broad regular expression. Why: It recognizes shape rather than membership in the roster.
- Alternative to reject: Use a retention period as the detector. Why: Preservation duration does not identify sensitive values.
- Alternative to reject: Assume any classifier has zero false positives. Why: Detection quality needs measurement.

Verification: Test known records, lookalike nonrecords and realistic benign content, recording false positives and false negatives.

Practical exercise: Compare pattern-based SIT, exact data match and trainable classification for three sample data shapes.

L33 - Sensitivity labels and publishing

Objective pointer: D2.P1.2. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: A sensitivity label describes and can protect content; a publishing policy makes labels available to the intended users. Encryption depends on configuration.

Northstar Consulting scenario: A new label exists, but the pilot users cannot select it in their supported applications.

Recommended approach: Check its publishing policy, scope and client availability, then validate both label selection and any configured protection.

- Alternative to reject: Recreate the label repeatedly. Why: The issue may be distribution rather than definition.
- Alternative to reject: Assume every label encrypts every file. Why: Protection depends on label configuration and workload.
- Alternative to reject: Apply a retention label instead. Why: Retention solves a different requirement.

Verification: Use an in-scope account to label test content and an excluded account to verify the intended publishing boundary.

Practical exercise: Draft a label specification with classification meaning, protection settings and publishing audience.

L34 - DLP requirements

Objective pointer: D2.P1.1. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: A DLP requirement connects a location, detection condition and supported action. One workload's action should not be assumed for all others.

Northstar Consulting scenario: Finance must prevent synthetic account details being emailed externally while allowing the approved internal workflow.

Recommended approach: Specify Exchange as the location, the relevant detection and external-recipient conditions, and a supported restrictive action with test cases.

- Alternative to reject: Configure only endpoint USB control. Why: That does not address this outbound email requirement.
- Alternative to reject: Write a rule without its workload location. Why: Available conditions and actions depend on the location.
- Alternative to reject: Assume a sensitivity label alone covers every outgoing path. Why: Labeling and DLP enforcement are separate decisions.

Verification: Test matching external mail, permitted internal mail and benign external mail; inspect policy events and user experience.

Practical exercise: Write a six-location DLP requirements matrix for Exchange, SharePoint, OneDrive, Teams, endpoint and Copilot.

L35 - Copilot DLP boundaries

Objective pointer: D2.P1.1. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: Copilot DLP must use supported scenario-specific conditions and actions; it does not automatically govern every agent interaction.

Northstar Consulting scenario: A team wants labelled documents excluded from Copilot processing while preserving separately authorized document access.

Recommended approach: Specify the supported Copilot DLP label condition and processing restriction, then test the exact experience and limitations.

- Alternative to reject: Remove every user's document permission regardless of the requirement. Why: That changes access beyond the stated processing restriction.
- Alternative to reject: Assume the rule blocks all agents in every platform. Why: Agent controls have separate workload boundaries.
- Alternative to reject: Treat a citation as proof the document text was processed. Why: A citation can remain even when content processing is restricted.

Verification: Record the relevant policy scope, labelled test item and processing behavior; distinguish a citation from use of item content.

Practical exercise: Prepare separate requirements for labelled grounding, sensitive prompt handling and agent communication.

L36 - Retention design

Objective pointer: D2.P1.3. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: Retention rules govern preservation and disposal over time. Labels can express item-specific needs within a broader location policy.

Northstar Consulting scenario: A site retains ordinary documents for three years; a particular record needs seven years from the same start date.

Recommended approach: Apply the appropriate item retention label and deployment policy while accounting for the broader retention requirements.

- Alternative to reject: Use a sensitivity label to set seven-year preservation. Why: Sensitivity and retention labels serve different functions.
- Alternative to reject: Assume a shorter delete action defeats active retention. Why: An active retention requirement prevents ordinary permanent deletion.
- Alternative to reject: Treat backup as the approved records schedule. Why: Recovery and lifecycle requirements are separate.

Verification: Document the record's applicable retain actions, start event and expected earliest ordinary disposal time.

Practical exercise: Draw a timeline showing a three-year site retain action and seven-year item retain action.

L37 - DLP alert response

Objective pointer: D2.P1.4. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: A DLP alert needs context about policy, item, actor and action before a responder can judge exposure and tune controls.

Northstar Consulting scenario: A DLP alert identifies an attempted external sharing action, but its status does not show whether the action was blocked.

Recommended approach: Inspect the matched rule, workload, affected content and recorded enforcement outcome before classifying the event.

- Alternative to reject: Assume every alert proves a successful data leak. Why: An attempted action may have been prevented.
- Alternative to reject: Disable the policy after one complaint. Why: One event does not establish that the rule is wrong.
- Alternative to reject: Treat the alert name as complete evidence. Why: The underlying event provides the decision context.

Verification: Record attempted versus completed activity, blocked versus allowed outcome and any justified follow-up.

Practical exercise: Write a triage note that separates confirmed exposure from a blocked attempt.

L38 - DSPM for AI

Objective pointer: D2.P1.5. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: Data security posture management identifies and prioritizes data risks; findings need owners, remediation and later verification.

Northstar Consulting scenario: A posture assessment finds many sensitive documents exposed to broad internal groups before an AI rollout.

Recommended approach: Prioritize the high-impact exposures, assign remediation owners and verify corrected sharing and relevant controls.

- Alternative to reject: Treat the dashboard as automatic remediation of every permission. Why: Visibility is not the same as a completed fix.
- Alternative to reject: Ignore internal oversharing because no guest is involved. Why: Unnecessary internal access is still a data risk.
- Alternative to reject: Measure success only by viewing the assessment. Why: An activity count does not prove reduced exposure.

Verification: Compare targeted before-and-after access findings and direct tests, noting assessment refresh timing.

Practical exercise: Create a five-item risk register with severity, owner, control, due date and verification.

Primary expansion reading:

<https://learn.microsoft.com/en-us/training/modules/protect-information-purview/>

<https://learn.microsoft.com/en-us/training/modules/prevent-data-loss-lifecycle-purview/>

M08 - Copilot readiness and rollout

Roll out Copilot only when the people and data are ready

In simple English

Copilot can help users work with information they are permitted to use. Grounding means supplying relevant information to help produce a response. A citation points to source material; you still need to check whether the answer represents that material correctly. A licence is only part of readiness. You also need supported access, applications, data and support arrangements.

How it works

Assess user readiness and data readiness separately. For users, check identity, entitlement, supported experience and operational support. For data, check whether important material is accurate, usable and shared with the right people. Broad existing access can make sensitive material easier to discover. Start with a small, defined group and known sources. Record the task, the expected source and the observed answer. Expand based on evidence instead of the number of licences purchased.

Worked example

Fictional example: 30 pilot users can open Copilot, but a finance folder is shared too widely. The technical access test passes; the data-readiness gate does not. Correct the sharing problem before widening exposure. If a permitted file is not cited, investigate source access and discovery/indexing before granting more access. An empty answer is not a reason to make the whole site public.

Use this process

1. Choose a small pilot and a useful, measurable task.
2. Check the users' identity, licence and client prerequisites.
3. Review permissions and source quality for pilot content.
4. Test permitted retrieval and a known denied boundary.
5. Record quality, unresolved risks and support needs before expansion.

Keep this distinction: Purchased seats are not active use. Active use is not useful output. Useful output is not proof that every data boundary is safe.

Stop and answer: All pilot users are licensed, but a sensitive site is overshared. Is the rollout ready?

Check your reasoning

No. Correct or contain the data exposure and verify the intended access boundary before expanding the rollout.

Read with the mapped lessons below. [Microsoft reference](#)

Estimated study: 4 hours. Lab prerequisites: Readiness and usage-report access plus scoped Copilot administration; licensed pilot users for live experiences.

A useful rollout has a business hypothesis: which task should improve, for whom, by how much and with what quality threshold? Define the baseline before asking users whether the new tool is impressive. A faster draft with more checking work may not save time overall.

Readiness has a technical side and a data side. Supported clients and entitlements make the capability available; appropriate source access makes it suitable for use. Existing oversharing can become much more visible when retrieval is easier. Correct unnecessary access at the source and verify it under the affected identity.

A phased rollout lets you learn with a small population. Set exit criteria before the pilot: required controls, support readiness, acceptable output quality and evidence of useful use. The decision to expand should follow these results. A severe unresolved exposure is a blocker even if satisfaction is high. Review the result with both the service owner and the data owner.

Key concepts: Copilot readiness, Data readiness, User experiences, Phased rollout decisions.

Diagram to create: Draw a launch gate with technical readiness, data readiness, support readiness and task-quality evidence all feeding the decision.

Knowledge check with answer: Why is license assignment insufficient for launch? Answer: it proves neither supported experiences nor safe data access or useful outcomes.

Common mistakes: Calling interest adoption, assuming AI creates new permission grants, and averaging mandatory control failures into a good score.

Exam tip: Look for the missing prerequisite or evidence, not the most ambitious rollout action.

L39 - Copilot readiness

Objective pointer: D3.C1.1. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: Readiness combines identity, entitlement, supported clients, service configuration and operational support. A purchased seat is only one component.

Northstar Consulting scenario: A pilot has licenses but mixed client versions and no support process for access failures.

Recommended approach: Assess the tenant and in-app prerequisites, correct readiness gaps and define pilot support before expansion.

- Alternative to reject: Declare readiness from seat count alone. Why: Entitlement does not prove supported client or service configuration.
- Alternative to reject: Roll out to everyone before testing. Why: That scales unresolved defects.
- Alternative to reject: Use user enthusiasm as technical acceptance. Why: Interest does not validate readiness.

Verification: Maintain a pass, fail or untested checklist for prerequisites and complete a representative pilot task.

Practical exercise: Create a readiness scorecard with evidence owners and explicit launch blockers.

L40 - Data readiness

Objective pointer: D3.C1.2. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: AI can make already-accessible information easier to find; existing oversharing therefore matters before rollout.

Northstar Consulting scenario: A pilot answer surfaces old HR content that the user can already open through a broad group grant.

Recommended approach: Correct the unintended source access and review similar sensitive content before expanding the rollout.

- Alternative to reject: Assume Copilot invented new permissions. Why: The observed grant already explains access.
- Alternative to reject: Hide only one answer. Why: The underlying source exposure remains.
- Alternative to reject: Give everyone access to make results consistent. Why: That widens the problem.

Verification: Prove the affected user's direct access changes as intended and then retest representative retrieval after propagation.

Practical exercise: Produce an oversharing remediation ticket with root grant, affected population and validation.

L41 - User experiences

Objective pointer: D3.C1.6. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: Adoption support should target actual job tasks, supported experiences and user feedback while preserving controls.

Northstar Consulting scenario: Users have Copilot access but do not know how to use it for their approved meeting and document workflows.

Recommended approach: Provide role-specific task guidance and supported experience checks, then measure repeated useful use and feedback.

- Alternative to reject: Grant extra tenant administration to all users. Why: Administrative rights are not training.
- Alternative to reject: Report licenses assigned as successful adoption. Why: Availability does not establish use.
- Alternative to reject: Remove data controls to make demonstrations impressive. Why: That sacrifices the rollout's requirements.

Verification: Review workload-level activity together with task completion and user-reported quality, not merely seat counts.

Practical exercise: Design a 15-minute onboarding exercise where users verify a generated summary against its source.

L42 - Phased rollout decisions

Objective pointer: D3.C1.1. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: A pilot is an experiment with acceptance criteria, not a ceremony before a predetermined launch.

Northstar Consulting scenario: A small rollout succeeds on speed but has unresolved restricted-data exposure and no tested support escalation.

Recommended approach: Hold expansion until the exposure and operational blockers meet the predefined acceptance criteria.

- Alternative to reject: Expand because response speed improved. Why: A benefit does not cancel a mandatory control failure.
- Alternative to reject: Average security failures into a passing satisfaction score. Why: Mandatory controls are not offset by unrelated successes.
- Alternative to reject: Delete the negative feedback. Why: That removes evidence rather than resolving the issue.

Verification: Re-run the failed access and escalation scenarios and record approval against the same launch criteria.

Practical exercise: Write a pilot exit decision with benefit evidence, blockers, remediation owners and a retest date.

Primary expansion reading:

<https://learn.microsoft.com/en-us/training/modules/enable-manage-microsoft-365-copilot/>

M09 - Copilot controls and integrations

Keep Copilot settings, providers and connectors separate

In simple English

Web search brings public-web information into supported Copilot responses. Internal search helps find work information. A connector brings an external source into supported search or retrieval. An AI provider supplies a model for eligible experiences. Cowork has its own access and spending controls. The right choice depends on the requirement; these are not interchangeable features.

How it works

First identify whether the request changes information sources, the model provider, user access or funding. For connectors, distinguish an indexed copy from live retrieval. Review identity mapping and source permissions. For providers, check the specific offering and data terms before enabling it. For Cowork, review spending-policy scope; the old agent entry is not the current access control. A web-search setting changes grounding behaviour, not Cowork entitlement.

Worked example

Fictional example: Northstar wants an approved external policy source available to staff. A connector may meet that need, but it must preserve the intended item boundary. Test a known permitted record and the same record as a denied identity. If the business instead wants a different model, review provider settings and terms. Do not enable a provider just to solve an indexing problem.

Use this process

1. Write the requested change as source, provider, access, experience or funding.
2. Identify the owning control and the specific product behaviour.
3. Review source permissions or provider terms as appropriate.
4. Pilot the supported change and allow documented propagation.
5. Verify the boundary with the intended identity and retain unresolved findings.

Keep this distinction: No indexing does not mean no runtime data transfer. A small Cowork spending limit does not mean no access. Microsoft-subprocessor status alone does not settle every geographic processing question.

Stop and answer: A user must have no Cowork access. Is setting a small positive credit limit sufficient?

Check your reasoning

No. Review every applicable spending policy that selects Cowork and remove the user from access-granting scope as required.

Read with the mapped lessons below. [Microsoft reference](#)

Estimated study: 4 hours. Lab prerequisites: AI administration and relevant Microsoft 365 Apps Cloud Policy role; external-source permission for connector work; approved billing only if exercising metered services.

AI settings are distributed because they control different boundaries. The capability a user sees, the model a service can use, the source it can retrieve and the meter that pays for execution are separate. Build a control register rather than expecting one master switch.

Connectors extend the information available to an experience. An indexed connection brings content and access metadata into an index; a runtime retrieval model has different storage and permission checkpoints. Both require identity and data-flow analysis. A working connection is not enough: a permitted user must retrieve the expected record, while a user without source access must fail for the correct reason.

New features change quickly. Cowork is a practical example of why an old registry entry cannot be assumed to remain the access control. Check current product documentation, tenant availability and effective scope. For a training exercise, inspecting configuration can be sufficient; enabling a billing policy is a separate financial action and is not necessary to learn the distinction.

Key concepts: Copilot web search, Copilot Search, Tenant Copilot settings, Cowork administration, Providers and connectors.

Diagram to create: Draw separate lanes for experience visibility, access grant, model selection, data retrieval and billing. Add indexed versus runtime connector flows.

Knowledge check with answer: Why does no indexing not mean no data transfer? Answer: a runtime connector still retrieves information for the requested operation.

Common mistakes: Treating disclaimers as enforcement, hidden features as revoked access, tiny budgets as no access, and Ready status as verified isolation.

Exam tip: Read which boundary the requirement names. Choose the control for that boundary and verify its current implementation.

L43 - Copilot web search

Objective pointer: D3.C1.3. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: Web grounding and work-data grounding have different information boundaries. Policy scope must be checked in the effective experience.

Northstar Consulting scenario: A pilot should retain permitted work-data use but restrict web grounding according to the organization's requirement.

Recommended approach: Configure the supported web-search policy for the intended users and modes, then verify effective scope and retrieval timing.

- Alternative to reject: Disable all SharePoint access. Why: That disrupts internal work content rather than targeting web search.
- Alternative to reject: Assume no citation proves web search is disabled. Why: A response might not have needed web search.
- Alternative to reject: Use a billing limit as a web-search policy. Why: Funding and search behavior are separate controls.

Verification: Inspect the effective policy assignment and test relevant modes after policy retrieval, recording uncertainty where runtime output is inconclusive.

Practical exercise: Create a mode-by-mode policy worksheet for work, web and chat experiences.

L44 - Copilot Search

Objective pointer: D3.C1.4. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: Search configuration influences finding permitted information; permission and source quality still govern useful results.

Northstar Consulting scenario: Staff cannot find approved material through Copilot Search even though they can open the source directly.

Recommended approach: Check supported search configuration, source availability and indexing or retrieval status while preserving permissions.

- Alternative to reject: Make the entire source public. Why: That is not a justified search repair.
- Alternative to reject: Assume direct access guarantees immediate search appearance. Why: Search discovery has additional dependencies.
- Alternative to reject: Treat missing search results as proof of permission denial. Why: Access must be tested separately.

Verification: Compare a known permitted item and a known denied item, documenting source access and discovery independently.

Practical exercise: Write a search test matrix containing query, identity, expected source and observed result.

L45 - Tenant Copilot settings

Objective pointer: D3.C1.5. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: AI tenant settings need distinct decisions for purchasing, admin experiences, releases, disclaimers and generation features.

Northstar Consulting scenario: The company wants controlled self-service purchases and a disclaimer while piloting new image features with approved users.

Recommended approach: Document and configure each supported setting at its actual scope, with an owner and validation for the pilot.

- Alternative to reject: Assume a disclaimer prevents data loss. Why: A notice is not an enforcement control.
- Alternative to reject: Use one purchase toggle as every AI control. Why: Purchasing, feature behavior and information protection differ.
- Alternative to reject: Set every release preference to broad rollout without evaluation. Why: That contradicts a controlled pilot.

Verification: Capture the effective settings and test the intended user experiences; note any setting that lacks the desired granularity.

Practical exercise: Build a settings register for self-service purchase, admin Copilot, releases, disclaimer, video and image generation.

L46 - Cowork administration

Objective pointer: D3.C1.7. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: Cowork access, discovery, model availability and web-search behavior are separate controls. Current access administration must be checked against current documentation.

Northstar Consulting scenario: An administrator blocks the old Cowork agent entry but leaves a user in a spending policy selecting Cowork.

Recommended approach: Review current spending-policy access scope and remove the user from every applicable Cowork access-granting scope if access must be prevented.

- Alternative to reject: Rely solely on the legacy registry entry. Why: The current documented access control has moved.
- Alternative to reject: Set an extremely small positive credit limit and call it no access. Why: A small limit still permits some use.
- Alternative to reject: Hide discovery and assume access is revoked. Why: Visibility does not equal authorization.

Verification: Review all applicable spending policies plus product prerequisites, and distinguish effective access from discoverability.

Practical exercise: Prepare a no-charge configuration review; do not enable usage merely to prove that the control exists.

L47 - Providers and connectors

Objective pointer: D3.C1.8;D3.C1.9. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: External AI and source integrations need explicit data-flow and identity analysis; connector models differ in indexing and query-time retrieval.

Northstar Consulting scenario: An external repository is connected for a pilot, but security has not checked identity mapping, data handling or provider terms.

Recommended approach: Review the provider's processing arrangement and connector model, then verify source permissions with allowed and denied users.

- Alternative to reject: Assume no indexing means no runtime data transfer. Why: Federated retrieval still transfers requested information.
- Alternative to reject: Map every external identity to an all-staff account. Why: That defeats isolation.
- Alternative to reject: Accept connection status as sufficient data-security evidence. Why: A healthy connector can still have incorrect access mapping.

Verification: Trace a synthetic record through the documented flow and test both authorized retrieval and the source or indexed access boundary.

Practical exercise: Draw indexed and federated data flows; add a separate provider-review checklist for processing terms and permitted data.

M10 - Agent identity and access

Treat an agent as an actor with its own identity

In simple English

An agent can perform work through tools and services. To govern it, you need to know which identity acts, what access it has and which person is accountable. Microsoft Entra Agent ID distinguishes a blueprint, its tenant principal, a runtime agent identity and an optional paired user account. These objects have different jobs; a familiar display name does not make them the same object.

How it works

The blueprint provides the authentication foundation. The child agent identity represents the running agent. A sponsor provides business accountability; a Microsoft Entra owner provides technical administration. Registry ownership is another layer. Keep these responsibilities current when staff move or leave. Apply supported access controls to the actual token subject. A policy aimed at agent identities does not automatically cover a paired agent user account or a human user involved in a separate flow.

Worked example

Fictional example: three finance agents come from one blueprint. The organisation wants supported agent Conditional Access to cover current and future child identities. Blueprint-level targeting addresses that scope. It does not mean an API-key connection is governed by Entra Conditional Access, and it does not silently include paired user accounts. Name each identity and authentication path in your design.

Use this process

1. Identify the agent and each identity or credential path it uses.
2. Record sponsor, technical owner and registry responsibility.
3. Check permissions and access-review requirements.
4. Select the policy subject and supported control deliberately.
5. Verify fresh relevant sign-in evidence after policy propagation.

Keep this distinction: An agent identity cannot complete a human MFA challenge. Agent identity, paired user account and human user are separate policy subjects. Moving a registry owner does not automatically update every identity owner.

Stop and answer: An agent uses an API key to reach a resource. Will an Entra token policy necessarily control that call?

Check your reasoning

No. The API-key path bypasses Entra token authentication. That path needs controls at the relevant service and credential boundary.

Read with the mapped lessons below. [Microsoft reference](#)

Estimated study: 4 hours. Lab prerequisites: Appropriate Entra Agent ID governance and Conditional Access roles; feature-specific Agent 365 and Entra licensing.

An agent is not simply another employee account. Its identity object, blueprint relationship, optional user account and delegated user context can produce different policy subjects. Before applying a control, identify who requests the token and which resource is involved.

Ownership also has layers. A business sponsor answers whether the agent should do the work; a technical owner maintains its identity; registry ownership supports curation. Put these responsibilities in a table and verify the actual objects instead of assuming one owner field resolves every obligation.

Lifecycle controls connect people changes to agent changes. When a sponsor leaves, transfer accountability or retire the agent deliberately. For access policy, supported blueprint scope can address present and future child identities, but it does not automatically cover separate user accounts or a legacy identity retained during migration. Inspect fresh token-request evidence after propagation. A successful API call alone does not tell you which policy was evaluated.

Key concepts: Agent identity lifecycle, Agent access governance, Agent ownership layers, Agent Conditional Access.

Diagram to create: Draw sponsor and technical owner connected to a blueprint, two child identities and one distinct agent user account. Mark the policy subject for each request.

Knowledge check with answer: Why cannot an agent identity complete human MFA? Answer: the noninteractive subject requires controls supported for that identity type.

Common mistakes: Confusing a child with its blueprint, registry owner with sponsor, and identity-targeted policy with user-account coverage.

Exam tip: Identify the token subject before choosing a policy or predicting inheritance.

L48 - Agent identity lifecycle

Objective pointer: D3.I1.1. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: An agent identity needs accountable ownership from creation through retirement. Its sponsor's departure is an operational event.

Northstar Consulting scenario: The business sponsor leaves while an agent continues processing project information.

Recommended approach: Use the supported agent lifecycle workflow to transfer sponsorship or retire the agent according to the approved decision.

- Alternative to reject: Leave the agent permanently unowned. Why: No accountable person remains for business decisions.
- Alternative to reject: Copy the departed employee's credentials into the agent. Why: That creates an inappropriate identity dependency.
- Alternative to reject: Assume deleting a registry description revokes all resource access. Why: Presentation and access lifecycle are different.

Verification: Inspect the resulting sponsor or retirement state and verify required access changes and audit evidence.

Practical exercise: Design a sponsor-departure workflow with transfer approval, failure handling and review date.

L49 - Agent access governance

Objective pointer: D3.I1.2. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: An agent's access should match its task and resource boundaries; supported access packages and policy controls need deliberate targeting.

Northstar Consulting scenario: A procurement agent needs temporary read access to a defined resource set for an approved pilot.

Recommended approach: Use the supported governed access arrangement for the agent and resource, with limited scope, ownership and an end-date review.

- Alternative to reject: Give the agent every permission used by Global Administrator. Why: The scope exceeds the task.
- Alternative to reject: Grant permanent access because the pilot is short. Why: Short duration strengthens the case for explicit expiry or removal.
- Alternative to reject: Assume publishing the agent grants all required resource permissions. Why: Availability and resource authorization differ.

Verification: Validate a required operation and a prohibited operation under the actual agent identity, then confirm expiry or removal behavior.

Practical exercise: Write an access request specifying identity, resources, actions, business purpose and end condition.

L50 - Agent ownership layers

Objective pointer: D3.I1.3. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: Business sponsorship, technical identity ownership and registry ownership need not be the same person or object.

Northstar Consulting scenario: A registry lists an owner, but the identity team finds no current business sponsor for the agent identity.

Recommended approach: Reconcile registry ownership, Entra technical ownership and business sponsorship as separate responsibilities.

- Alternative to reject: Assume the registry field populates every ownership layer. Why: Ownership surfaces are not interchangeable.
- Alternative to reject: Remove all owners to avoid duplication. Why: That removes accountability.
- Alternative to reject: Treat a billing contact as automatic resource approver. Why: Billing and authorization roles are distinct.

Verification: Produce an ownership matrix tied to the relevant objects and verify that each accountable person accepts the responsibility.

Practical exercise: Draft a RACI for business approval, identity administration, registry curation and incident response.

L51 - Agent Conditional Access

Objective pointer: D3.I1.2. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: Agent-token policy evaluation depends on the subject, resource and supported controls; agents cannot satisfy a human authentication prompt.

Northstar Consulting scenario: Current and future child identities from one agent blueprint require a supported blocking policy. An agent user account also exists.

Recommended approach: Use supported blueprint-level scope for child identities and assess the agent user account separately with its supported controls.

- Alternative to reject: Target one child and assume future children are covered. Why: Individual identity scope does not cover new siblings.
- Alternative to reject: Require the agent identity to perform interactive MFA. Why: The nonhuman subject cannot satisfy that human interaction.
- Alternative to reject: Assume blueprint targeting automatically includes the agent user account. Why: The user account is a separate policy subject.

Verification: Use fresh token-request evidence for included and excluded identities after propagation, and assess the separate user-account policy.

Practical exercise: Draw blueprint, child identity and agent user account as distinct nodes with separate policy scopes.

Primary expansion reading:

<https://learn.microsoft.com/en-us/training/modules/establish-agent-identities-entra-agent-id/>

M11 - Agent registry and tool governance

Review an agent before people can use its capabilities

In simple English

An agent registry helps administrators discover and manage agents. Publication, installation, availability and access are related but distinct decisions. A tool is a capability an agent can call. Approving an agent's name or description is not enough; you must understand its tools, permissions, data sources and intended audience.

How it works

For a new request, identify the agent's purpose, publisher, owner and users. Review required tools and data access. Decide whether it should be published, rejected or returned for changes. Use the registry's supported actions for the intended outcome. Then verify what the target users can actually do. A block or removal needs outcome evidence; do not assume a catalogue change removes every underlying credential or source permission.

Worked example

Fictional example: a draft agent should search an approved policy library. Its proposed tool can also modify records in a business system. The extra capability does not follow from the search requirement. Ask for a narrower design or reject the request pending correction. If the agent is later approved, test the allowed task and a disallowed operation. Keep tool permission and agent availability evidence separate.

Use this process

1. Inventory purpose, publisher, owner, audience, tools and sources.
2. Compare each permission with a required task.
3. Record the publication or rejection decision and conditions.
4. Scope installation and access using supported registry controls.
5. Test effective availability and tool behaviour for the intended audience.

Keep this distinction: Discovery is inventory, not automatic approval. Publication is not blanket data permission. Removing a visible agent does not prove its separate credentials are revoked.

Stop and answer: A read-only assistant requests a write-capable tool. Should you approve it because the description says read-only?

Check your reasoning

No. Review and constrain actual capabilities. A description does not enforce the boundary.

Read with the mapped lessons below. [Microsoft reference](#)

Estimated study: 4 hours. Lab prerequisites: Appropriate agent registry curation and tool-management permissions; do not assume view access includes install or consent rights.

An agent catalog is a governed inventory, not a universal list of everything anyone could have built. Reconcile it against platforms and team reports. Record draft visibility and integration gaps as findings. An incomplete count should prompt investigation, not an assertion that missing agents cannot exist.

Publishing decisions need specific evidence: purpose, owners, permissions, data sources, tools and negative tests. Approving a useful purpose is not approving every requested capability. Return a request for changes when the privilege is broader than the business task.

Tools deserve their own review because they determine what the agent can do. Read-only retrieval, sending a message and changing a shared record have different consequences. A tool block and an agent block can contain different scopes. During response, choose the smallest effective containment supported by the evidence, but verify it actually stops the problematic operation.

Key concepts: Agent settings, Agent inventory, Agent publication review, Registry actions and tools.

Diagram to create: Draw request → review → approve, return or reject, then show agent access and tool access as separate gates.

Knowledge check with answer: Why is a registry owner not proof that the agent has safe tools?
Answer: accountability and permitted capabilities are different properties.

Common mistakes: Equating installed with approved for everyone, registry inventory with all possible agents, or view permission with management permission.

Exam tip: Check whether the question targets tenant defaults, one agent, one user population or a tool.

L52 - Agent settings

Objective pointer: D3.I1.4. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: Tenant agent settings shape permitted creation, sharing, templates and user availability; these are governance choices with different scopes.

Northstar Consulting scenario: A pilot allows an approved agent type to a limited group and prohibits broad creator sharing.

Recommended approach: Review allowed types, sharing, template and user-access settings, then test the actual pilot and excluded population.

- Alternative to reject: Assume approval of one agent changes every tenant setting. Why: Instance decisions and tenant settings differ.
- Alternative to reject: Publish all templates to everyone. Why: That exceeds the bounded pilot.
- Alternative to reject: Use a sensitivity-label name as the sharing configuration. Why: Content labels do not substitute for agent settings.

Verification: Record effective settings and both permitted and prohibited sharing or access attempts.

Practical exercise: Prepare an agent-settings register with default, pilot exception, owner and expiry.

L53 - Agent inventory

Objective pointer: D3.I1.5. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: A registry enables discovery and curation, but inventory visibility depends on integration and discovery coverage.

Northstar Consulting scenario: The dashboard count is lower than the number of agents reported by teams using multiple platforms.

Recommended approach: Reconcile inventory coverage, connected platforms and draft visibility before declaring that unlisted agents do not exist.

- Alternative to reject: Assume the dashboard is a complete inventory of every possible platform. Why: Coverage and draft visibility can vary.
- Alternative to reject: Delete reported agents from the audit worksheet. Why: That conceals the gap.
- Alternative to reject: Treat usage count as the same as catalog count. Why: Available agents and active agents are different populations.

Verification: Maintain a reconciled inventory noting source platform, registry state, owner and uncovered agents.

Practical exercise: Build a five-agent inventory with one draft, one inactive agent and one disconnected platform.

L54 - Agent publication review

Objective pointer: D3.I1.6. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: Publication is a governance decision that should evaluate purpose, ownership, permissions, data flows and evidence.

Northstar Consulting scenario: A requested agent has a useful purpose but asks for broad mailbox write access without justification.

Recommended approach: Return or reject the request pending scoped permissions and adequate test evidence.

- Alternative to reject: Publish first and review permissions later. Why: That exposes the requested privilege before review.
- Alternative to reject: Approve because the description sounds useful. Why: A business benefit does not justify every permission.
- Alternative to reject: Treat a familiar publisher as proof of fit. Why: The actual requested access still needs assessment.

Verification: Keep the review decision, required changes and verification evidence attached to the specific version or request.

Practical exercise: Complete a publication review rubric with approve, return for changes or reject outcomes.

L55 - Registry actions and tools

Objective pointer: D3.I1.7;D3.I1.8. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: Agent availability and the tools it can invoke are related but distinct control surfaces.

Northstar Consulting scenario: An approved custom agent invokes an unwanted tool. The incident requires disabling that capability while retaining unrelated approved tools.

Recommended approach: Use the supported tool control for the unwanted tool and verify the agent's resulting behavior; assess agent blocking if needed for containment.

- Alternative to reject: Unpublish an unrelated agent. Why: That does not address the tool in question.
- Alternative to reject: Assume installing the agent makes every tool acceptable. Why: Tools require their own governance.
- Alternative to reject: Delete only the tool's description. Why: Metadata removal does not prove blocked invocation.

Verification: Observe the blocked invocation and continued permitted behavior, with corresponding governance records.

Practical exercise: Design a custom-agent intake checklist and a separate tool review with publisher, permissions and block test.

Primary expansion reading:

<https://learn.microsoft.com/en-us/training/modules/curate-control-agents-registry-tools/>

M12 - Agent security and compliance operations

Monitor what agents actually do

In simple English

Agent governance continues after publication. You need evidence of activity, permissions, data handling and ownership. A compliance gap is a requirement that is not yet met or not yet evidenced. A quiet dashboard can mean low activity, incomplete coverage or missing data. It is not automatic proof that the agent is safe.

How it works

Start with the expected behaviour: what task, which sources, which tools and which audience? Compare observed activity with that baseline. Investigate unusual volume, unexpected destinations or sensitive-data events in context. Check the scope of your monitoring before interpreting an absence of alerts. Assign a named owner to each gap, define corrective work and verify the outcome. Agent 365 supports the monitoring and governance work, but the administrator must still interpret the evidence.

Worked example

Fictional example: an approved agent usually reads a policy source. Activity suddenly includes another destination. Investigate the identity, tool, request and source permissions. Do not call it a breach solely because the destination is new, and do not dismiss it solely because the agent was approved last month. Compare the activity with the authorised purpose and apply the response process when evidence justifies it.

Use this process

1. Define expected agent behaviour and required monitoring coverage.
2. Review activity with identity, tool, resource and time context.
3. Separate confirmed behaviour from missing or delayed evidence.
4. Assign remediation, owner and due date for each gap.
5. Retest the boundary and document the closure evidence.

Keep this distinction: Approval is a decision at a point in time. Monitoring tests continued behaviour. Compliance claims require evidence against a stated requirement.

Stop and answer: An agent has no alerts, but you have not checked monitoring coverage. Can you declare it compliant?

Check your reasoning

No. Establish coverage and assess the applicable requirements before drawing that conclusion.

Read with the mapped lessons below. [Microsoft reference](#)

Estimated study: 3 hours. Lab prerequisites: Agent 365 observation and appropriate security or compliance roles; connected data sources and audit availability for live evidence.

An approved agent can behave unexpectedly later. Workload changes, new data, tool modifications or loops can change its risk. Monitoring should compare actual behavior with expected business activity and associate each finding with an accountable owner.

Trace sensitive information through the complete action path. Being allowed to read a file does not automatically authorize transmitting it to another recipient. Identify the source, identity, tool, destination and enforcement point. Test the actual path with synthetic data. An instruction in the agent prompt can guide behavior but is weaker than enforced access and data controls.

Compliance work is evidence management. List each requirement, its implemented control, the test result and limitations. A stable agent with no errors may still lack a required owner review or retention arrangement. Mark untested controls honestly. Closure requires evidence of remediation and an agreed residual-risk decision, not merely a green dashboard.

Key concepts: Agent activity monitoring, Agent data protection, Agent compliance gaps.

Diagram to create: Draw source → agent → tool → destination, with identity and data-control checkpoints; add an evidence-to-requirement compliance table.

Knowledge check with answer: Why is a tenfold activity spike not proof of compromise? Answer: legitimate demand, retries or a loop can also explain it.

Common mistakes: Treating approval as permanent safety, prompt wording as enforcement and product availability as organizational compliance.

Exam tip: Follow the data and the identity; distinguish attempted, permitted, blocked and unknown outcomes.

L56 - Agent activity monitoring

Objective pointer: D3.G1.1. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: Monitoring links an agent's identity and actions to expected behavior, owners and anomalies. Counts alone do not establish safety.

Northstar Consulting scenario: An agent's calls increase tenfold overnight, but the cause might be legitimate demand or a repeated loop.

Recommended approach: Correlate activity, affected tools, errors, timing and business context before choosing a targeted response.

- Alternative to reject: Declare compromise solely from the count. Why: Volume is a signal requiring context.
- Alternative to reject: Ignore the increase because the agent was previously approved. Why: Approval does not make later behavior safe.
- Alternative to reject: Increase the budget immediately. Why: That may fund an unintended loop.

Verification: Document the activity pattern and root-cause evidence, plus the effect of any containment or correction.

Practical exercise: Analyze a synthetic hourly usage series with a repeated error loop and a legitimate batch period.

L57 - Agent data protection

Objective pointer: D3.G1.2. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: Agent data protection must match the actual workload and action path; a broad AI label is not a universal enforcement guarantee.

Northstar Consulting scenario: An agent reads a permitted file but attempts to send its sensitive contents through a different communication path.

Recommended approach: Map the read and send paths and apply the supported identity, permission and data-protection controls to the actual operation.

- Alternative to reject: Assume the file's permitted read automatically permits external sending. Why: Read and dissemination are separate decisions.

- Alternative to reject: Assume Copilot DLP universally covers every agent channel. Why: Coverage is workload-bound.

- Alternative to reject: Rely only on a prompt asking the agent to be careful. Why: An instruction is not a substitute for enforced controls.

Verification: Test a synthetic sensitive outbound operation and a benign permitted operation under the actual agent execution path.

Practical exercise: Draw data source, agent, tool and recipient boundaries and identify where enforcement is verified.

L58 - Agent compliance gaps

Objective pointer: D3.G1.3. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: A compliance finding is a gap between required controls and observed evidence, not merely an unfavorable dashboard color.

Northstar Consulting scenario: An agent handles regulated records but has no documented owner review or tested retention requirements.

Recommended approach: Record the specific gaps, assign accountable remediation and require evidence before closing them.

- Alternative to reject: Close the finding because the agent has no recent errors. Why: Reliability does not prove compliance.

- Alternative to reject: Mark every control compliant because the product is Microsoft. Why: Organizational configuration and use still matter.

- Alternative to reject: Replace evidence with a general policy statement. Why: The implemented control must be demonstrated.

Verification: Link each requirement to a configured control, tested outcome, owner and unresolved limitation.

Practical exercise: Create a compliance matrix containing met, unmet and not-yet-tested states.

Primary expansion reading:

<https://learn.microsoft.com/en-us/training/modules/agent-365-monitor-manage/>

<https://learn.microsoft.com/en-us/training/modules/agent-365-govern/>

M13 - AI economics, adoption and service health

Separate adoption, business value, cost and reliability

In simple English

Adoption tells you whether people use a service. Business value tells you whether useful work improves. Cost tells you what resources or subscriptions are consumed. Service health tells you whether the service is working. A single chart rarely answers all four questions. Match the report to the decision you need to make.

How it works

Define the reporting period, the users or agents in scope and the measure before comparing reports. A licensed-user count is not the same as an active-user count. A usage event is not necessarily a billable event. Reconcile charges with their funding route and avoid counting the same cost twice. For value, use an agreed task measure such as time to produce an accepted draft, including review effort. For reliability, combine the incident evidence with the users' actual symptoms.

Worked example

Fictional teaching data: 80 people are licensed, 48 are active in the chosen period, and 36 complete the target workflow. Active adoption among licensed users is $48/80 = 60\%$. Workflow completion among licensed users is $36/80 = 45\%$. These are different measures. If the workflow's review burden rises, higher activity alone does not establish better productivity. State the denominator instead of calling every percentage adoption.

Use this process

1. Write the business question before opening a dashboard.
2. Match the report's period, denominator and scope.
3. Reconcile costs by service and funding route.
4. Compare usage with task outcomes and support burden.
5. Use service health and workload evidence to investigate reliability.

Keep this distinction: Seats are capacity. Active users show use. Accepted outcomes show one form of value. Charges require billing evidence. None of these should silently stand in for another.

Stop and answer: An activity report and a billing report show different totals. Is one necessarily wrong?

Check your reasoning

No. Check metric definitions, periods, processing delay and charge coverage. They may count different things.

Read with the mapped lessons below. [Microsoft reference](#)

Estimated study: 3 hours. Lab prerequisites: Billing or cost read access, usage-report access and Service health read access appropriate to the task.

Separate availability, activity and outcomes. Licenses show that people could use a capability. Active-user measures show defined behavior during a period. Neither alone proves a valuable business result. Pair activity with an accepted-task measure such as a source-checked summary

meeting a quality rubric.

Cost needs careful boundaries. If two dashboards display the same metered charge, adding them inflates the total. Reconcile identifiers, periods and units. Keep distinct fixed and variable charges visible, and record what is excluded. Cost per accepted outcome is total relevant cost divided by outcomes meeting the agreed standard; it differs from cost per prompt.

A service-health problem can reduce usage, but low usage can also result from unsuitable tasks, training gaps or access defects. Investigate before attributing cause. An operations report should state what is known, what is inferred and what action will test the inference. This makes the report useful to both finance and the service owner.

Key concepts: AI cost reconciliation, Adoption and AI service health.

Diagram to create: Draw three columns: entitlement, activity and accepted outcomes. Add a cost-reconciliation table with overlap removal.

Knowledge check with answer: What is adoption for 40 active users among 100 licensed users?
Answer: 40 percent for that defined population and period; it is not a value claim.

Common mistakes: Double-counting cost surfaces, comparing unequal report windows and equating prompts with productivity.

Exam tip: Write the denominator and period before calculating. Check whether the requested metric is usage, cost, value or health.

L59 - AI cost reconciliation

Objective pointer: D3.M1.1. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: Cost reporting needs consistent periods, billing boundaries and units; overlapping views must not be added as separate charges.

Northstar Consulting scenario: A Microsoft 365 AI report shows 1200 currency units, including 300 also shown in an Azure view for the same metered use.

Recommended approach: Reconcile the common charge identifiers and period, then count the overlapping 300 only once.

- Alternative to reject: Add both totals as independent charges. Why: That double-counts the same metered use.
- Alternative to reject: Ignore all Azure costs. Why: Some separate charges may genuinely exist.
- Alternative to reject: Report prompt count as currency. Why: Usage and financial measures require a defined conversion.

Verification: Produce a reconciled ledger with unique charges, overlaps, excluded taxes or credits and the reporting period.

Practical exercise: Reconcile 900 fixed cost, 300 shared metered cost and 100 additional distinct cost to a total of 1300.

L60 - Adoption and AI service health

Objective pointer: D3.M1.2;D3.M1.3;D3.M1.4. Outcome: given a requirement, select the appropriate control, reject the alternatives and specify evidence of the result.

Theory: Usage, useful outcomes and service reliability are different measures. Healthy adoption reporting defines the population and time window.

Northstar Consulting scenario: There are 100 licensed users, 40 active users and a suspected service incident. Leadership asks whether the rollout is successful.

Recommended approach: Report the defined 40 percent active-license rate, investigate service health separately and collect task-quality evidence before claiming value.

- Alternative to reject: Call 100 purchased seats 100 percent adoption. Why: Licensing is not activity.
- Alternative to reject: Treat a healthy service dashboard as proof users gain value. Why: Availability is not business benefit.
- Alternative to reject: Infer a service outage solely from low usage. Why: Other causes such as training and workload fit remain possible.

Verification: Reconcile the report population and period, compare workload-level use and inspect matching service-health evidence.

Practical exercise: Write a dashboard commentary with adoption, cost per accepted outcome, incident status and next action.

Primary expansion reading:

<https://learn.microsoft.com/en-us/training/modules/monitor-optimize-m365-ai-services/>

Lab guide - 13 bounded exercises

Each lab uses the shared evidence standard above. Complete lessons before the lab; use the supplied CSV fixtures when a required tenant feature is missing. Numbered steps are a workflow, not a promise that portal labels remain unchanged. Current linked product guidance governs UI-specific implementation.

LAB-01 - Tenant foundations and licensing

Prerequisites and access: Tenant or Entra settings administrator; license administrator for entitlement work. Use a separate test tenant.

1. Inventory a test tenant without changing production.
2. Record its tenant identifier and approved admin role.
3. Open organization settings and record one branding and one profile setting.
4. Open domain settings and document verification status; use an offline DNS ticket if you do not own a spare domain.
5. Inspect a pilot licensing group and distinguish direct from nested membership.
6. Compare two test users' assigned products and service plans.
7. Identify one missing entitlement in the supplied fixture.
8. Draft a targeted correction rather than assigning a broader product automatically.
9. Verify expected user coverage.
10. Restore any temporary branding after capturing the result..

Expected deliverable and acceptance: Tenant baseline, DNS change ticket and entitlement reconciliation. Pass when the trainee separates verification from routing and detects the nested-membership gap.

Troubleshooting: if the feature is absent, check role, subscription and rollout status using the linked source. If a saved setting appears ineffective, verify target identity and scope, then documented propagation. If access changes unexpectedly, restore the recorded baseline for the disposable lab scope and stop expanding the change.

Cleanup evidence: record objects removed, settings restored and any retained artifact or unresolved dependency. Never delete a protected record just to obtain a clean lab.

LAB-02 - Operational resilience and funding

Prerequisites and access: Read roles for health and billing review; supported Backup role and approved billing for an optional live restore.

1. Read Service health and record a current or historical advisory without assuming it affects the pilot.
2. Write a notification plan rather than subscribing unrelated people.
3. Compare the two synthetic connectivity records.

4. Identify the branch-only failing path.
5. Inspect Backup availability and licensing without enabling billable protection.
6. In a funded sandbox only, protect one synthetic object and follow the supported restore workflow.
7. Record target time, destination, job state and restored-content validation.
8. For the offline route, choose the correct restore point from the fixture and explain uncertainty.
9. Reconcile the AI-cost exercise.
10. Remove only disposable lab objects and any lab-only billing configuration after review..

Expected deliverable and acceptance: Incident triage worksheet, restore acceptance record and funding matrix. Offline completion is a design exercise, not evidence of a live restore.

Troubleshooting: if the feature is absent, check role, subscription and rollout status using the linked source. If a saved setting appears ineffective, verify target identity and scope, then documented propagation. If access changes unexpectedly, restore the recorded baseline for the disposable lab scope and stop expanding the change.

Cleanup evidence: record objects removed, settings restored and any retained artifact or unresolved dependency. Never delete a protected record just to obtain a clean lab.

LAB-03 - Collaboration and safe discovery

Prerequisites and access: Exchange, Teams and SharePoint scoped administrators; supported Copilot and SharePoint Advanced Management entitlement for advanced live tests.

1. Create or use a disposable collaboration workspace.
2. Assign two pilot identities: permitted and excluded.
3. On a test shared mailbox, document read and send delegation separately.
4. Send only to a lab recipient if testing delivery.
5. Create a team and inspect a standard and private channel membership boundary.
6. Upload a synthetic finance document to a test site.
7. Record its effective permission grants.
8. Test its direct URL as both users.
9. Where licensed, evaluate the supported discovery restriction and allow documented propagation; do not conclude failure from an immediate query.
10. In a licensed lab meeting, review policy and organizer options, start a transcript with permission and validate the saved result.
11. Restore pilot settings and delete disposable content subject to any retention controls..

Expected deliverable and acceptance: Permission matrix, two-identity access tests and meeting-state worksheet. Pass requires distinguishing discovery, authorization and transcript state.

Troubleshooting: if the feature is absent, check role, subscription and rollout status using the linked source. If a saved setting appears ineffective, verify target identity and scope, then

documented propagation. If access changes unexpectedly, restore the recorded baseline for the disposable lab scope and stop expanding the change.

Cleanup evidence: record objects removed, settings restored and any retained artifact or unresolved dependency. Never delete a protected record just to obtain a clean lab.

LAB-04 - Identity administration at scale

Prerequisites and access: Supported User, Groups or Exchange roles for the task; Privileged Role Administrator for PIM configuration where authorized; read-only Graph permissions for inventory.

1. Create a target inventory from the supplied identities fixture.
2. Mark members, guests and contacts.
3. Identify the duplicate proposed user principal name.
4. Assign business owners to the guest review and leaver transfer.
5. In a lab tenant, inspect a supported administrative-unit role assignment and compare in-scope and out-of-scope test objects.
6. If PIM is available, inspect an eligible test role and its activation requirements; do not elevate a production account.
7. Run the read-only Graph script after approved consent and tenant confirmation.
8. Reconcile the export with the expected inventory.
9. Draft a five-user write pilot plan without executing bulk writes.
10. Disconnect the Graph session and retain sanitized evidence..

Expected deliverable and acceptance: Lifecycle checklist, delegated-scope test and read-only inventory. Pass requires explaining all exported fields and identifying the duplicate before any write.

Troubleshooting: if the feature is absent, check role, subscription and rollout status using the linked source. If a saved setting appears ineffective, verify target identity and scope, then documented propagation. If access changes unexpectedly, restore the recorded baseline for the disposable lab scope and stop expanding the change.

Cleanup evidence: record objects removed, settings restored and any retained artifact or unresolved dependency. Never delete a protected record just to obtain a clean lab.

LAB-05 - Authentication and access decisions

Prerequisites and access: Authentication Administrator or appropriate policy role; Conditional Access Administrator for a dedicated sandbox; required Entra licensing.

1. Use the synthetic sign-in fixture before touching any tenant policy.
2. Predict each user's outcome.
3. Open sign-in logs in a permitted lab and locate one controlled test event.
4. Record identity, resource, client, authentication and policy results.
5. Review method availability and registration for a pilot user.

6. Inspect SSPR scope and required methods.
7. Draft a Conditional Access policy scoped only to test users, with a documented recovery route.
8. Where authorized and licensed, create it report-only and inspect an included and excluded event.
9. Do not convert it to enforcement merely to finish this lab.
10. Export sanitized findings and remove the disposable report-only policy after review..

Expected deliverable and acceptance: Policy truth table and event-based explanation. Fail if the learner claims that report-only blocked a user or that one grant overrides a block.

Troubleshooting: if the feature is absent, check role, subscription and rollout status using the linked source. If a saved setting appears ineffective, verify target identity and scope, then documented propagation. If access changes unexpectedly, restore the recorded baseline for the disposable lab scope and stop expanding the change.

Cleanup evidence: record objects removed, settings restored and any retained artifact or unresolved dependency. Never delete a protected record just to obtain a clean lab.

LAB-06 - Threat protection and response

Prerequisites and access: Appropriate Defender security read or response role; Defender for Office 365 Plan 2 or equivalent entitlement for advanced investigation and simulation features.

1. Read the synthetic incident timeline.
2. Identify the delivered message and affected recipients.
3. Separate delivery, click and suspected identity compromise.
4. In a permitted Defender environment, inspect an example alert and its available linked evidence without executing remediation on real users.
5. Review a test recipient's policy source for the relevant feature.
6. Draft the targeted response and success criteria.
7. Prepare a simulation plan for fictional lab recipients only.
8. Do not launch a campaign from this workbook.
9. Write an incident closure note naming completed and outstanding actions.
10. Have a peer challenge one unsupported inference..

Expected deliverable and acceptance: Investigation timeline, policy-source analysis and simulation design. Live remediation or campaign delivery is not required for completion.

Troubleshooting: if the feature is absent, check role, subscription and rollout status using the linked source. If a saved setting appears ineffective, verify target identity and scope, then documented propagation. If access changes unexpectedly, restore the recorded baseline for the disposable lab scope and stop expanding the change.

Cleanup evidence: record objects removed, settings restored and any retained artifact or unresolved dependency. Never delete a protected record just to obtain a clean lab.

LAB-07 - Information protection and governance

Prerequisites and access: Purview read roles for inspection; appropriate DLP, information-protection and lifecycle roles for isolated policy work; feature-specific licensing.

1. Use synthetic documents only.
2. Classify three sample data shapes and justify the chosen detector.
3. Draft a sensitivity label and publishing scope, distinguishing optional protection settings.
4. In a licensed sandbox, publish a harmless test label to pilot users and verify availability; otherwise complete the configuration worksheet.
5. Specify DLP requirements for all six covered locations, noting unsupported combinations.
6. If creating a policy, use a supported simulation or test mode and a bounded pilot; do not enable tenant-wide enforcement.
7. Analyze the supplied blocked-sharing alert.
8. Build a retention timeline without creating irreversible records settings.
9. Review a posture finding and assign an owner.
10. Remove reversible lab policies only after checking whether retention or records restrictions affect cleanup..

Expected deliverable and acceptance: Classification rationale, six-location DLP matrix, alert triage and retention timeline. Pass requires separate allowed and denied test cases.

Troubleshooting: if the feature is absent, check role, subscription and rollout status using the linked source. If a saved setting appears ineffective, verify target identity and scope, then documented propagation. If access changes unexpectedly, restore the recorded baseline for the disposable lab scope and stop expanding the change.

Cleanup evidence: record objects removed, settings restored and any retained artifact or unresolved dependency. Never delete a protected record just to obtain a clean lab.

LAB-08 - Copilot readiness and rollout

Prerequisites and access: Readiness and usage-report access plus scoped Copilot administration; licensed pilot users for live experiences.

1. Define a fictional approved task: summarize a public project-status document for internal review.
2. Record a manual baseline time and quality checklist.
3. Complete the readiness worksheet for two pilot users.
4. Identify one unsupported client and one overshared source in the fixtures.
5. Propose and verify the intended corrections offline or in a licensed test tenant.
6. Run the task only with synthetic content.
7. Compare the output against its source and record omissions.
8. Draft a short user guide and escalation path.

9. Hold a pilot exit review using the predefined criteria.
10. Document hold or expand with evidence and a review date..

Expected deliverable and acceptance: Readiness register, task-quality evaluation and launch decision. A useful answer must show the source checking performed by the learner.

Troubleshooting: if the feature is absent, check role, subscription and rollout status using the linked source. If a saved setting appears ineffective, verify target identity and scope, then documented propagation. If access changes unexpectedly, restore the recorded baseline for the disposable lab scope and stop expanding the change.

Cleanup evidence: record objects removed, settings restored and any retained artifact or unresolved dependency. Never delete a protected record just to obtain a clean lab.

LAB-09 - Copilot controls and integrations

Prerequisites and access: AI administration and relevant Microsoft 365 Apps Cloud Policy role; external-source permission for connector work; approved billing only if exercising metered services.

1. Create a control register covering web search, Search, purchase settings, admin experiences, release preferences, disclaimer and media generation.
2. Read the current Cowork product page linked in this guide and record its date.
3. Inspect spending-policy scope without enabling chargeable usage.
4. Draft a Cloud Policy web-search pilot with actual user scope.
5. Review a hypothetical third-party provider and record processing questions requiring resolution.
6. Use the connector fixture to detect a mismatched external identity.
7. Draw the corrected mapping.
8. In an authorized sandbox only, configure a limited synthetic source and test allowed and denied users.
9. Record propagation and removal behavior.
10. Remove disposable connections only after confirming whether indexed content remains..

Expected deliverable and acceptance: Control register, provider review and connector access-test record. Do not claim a live integration was tested when only the fixture was used.

Troubleshooting: if the feature is absent, check role, subscription and rollout status using the linked source. If a saved setting appears ineffective, verify target identity and scope, then documented propagation. If access changes unexpectedly, restore the recorded baseline for the disposable lab scope and stop expanding the change.

Cleanup evidence: record objects removed, settings restored and any retained artifact or unresolved dependency. Never delete a protected record just to obtain a clean lab.

LAB-10 - Agent identity and access

Prerequisites and access: Appropriate Entra Agent ID governance and Conditional Access roles; feature-specific Agent 365 and Entra licensing.

1. Use the supplied agent inventory to distinguish blueprint, child identity, optional user account and legacy service principal.
2. Assign business and technical ownership in the worksheet.
3. Draft a sponsor-transfer workflow for the departing owner.
4. Review the current supported Conditional Access controls for agent identities.
5. Define a report-only pilot with an included child, a future child and an excluded identity.
6. In a licensed authorized tenant, inspect the supported configuration and fresh sign-in evidence; otherwise annotate synthetic events.
7. Check the separate user-account policy requirement.
8. Document any API-key path as outside Entra token enforcement.
9. Remove disposable pilot configurations after review..

Expected deliverable and acceptance: Identity diagram, ownership matrix, lifecycle workflow and scoped policy evidence. Pass requires explicitly separating agent identities from agent user accounts.

Troubleshooting: if the feature is absent, check role, subscription and rollout status using the linked source. If a saved setting appears ineffective, verify target identity and scope, then documented propagation. If access changes unexpectedly, restore the recorded baseline for the disposable lab scope and stop expanding the change.

Cleanup evidence: record objects removed, settings restored and any retained artifact or unresolved dependency. Never delete a protected record just to obtain a clean lab.

LAB-11 - Agent registry and tool governance

Prerequisites and access: Appropriate agent registry curation and tool-management permissions; do not assume view access includes install or consent rights.

1. Load the five-agent fixture.
2. Reconcile owner, platform, registry state and observed use.
3. Record the missing inventory coverage.
4. Review the proposed mailbox-writing agent and choose approve, return or reject with reasons.
5. Prepare an intake record for a custom agent upload without publishing it.
6. Inspect a permitted registry entry and distinguish available actions from your current role's rights.
7. Review the unwanted-tool scenario and document tool versus whole-agent containment.
8. Where available, use a disposable tool or agent configuration to test an allowed and denied path.
9. Do not install an external executable or grant broad consent to complete the exercise.
10. Record rollback and remaining dependencies..

Expected deliverable and acceptance: Reconciled registry, publication decision and tool-control test. Pass requires matching containment to the actual capability.

Troubleshooting: if the feature is absent, check role, subscription and rollout status using the linked source. If a saved setting appears ineffective, verify target identity and scope, then documented propagation. If access changes unexpectedly, restore the recorded baseline for the disposable lab scope and stop expanding the change.

Cleanup evidence: record objects removed, settings restored and any retained artifact or unresolved dependency. Never delete a protected record just to obtain a clean lab.

LAB-12 - Agent security and compliance operations

Prerequisites and access: Agent 365 observation and appropriate security or compliance roles; connected data sources and audit availability for live evidence.

1. Plot or tabulate the supplied hourly agent activity.
2. Identify the repeated-error interval and distinguish it from the planned batch.
3. Draft a proportionate containment decision and expected recovery evidence.
4. Map the synthetic sensitive-data send scenario from source to destination.
5. Specify a supported control at each relevant boundary, marking unsupported or unknown coverage.
6. Review three compliance requirements and assign met, unmet or untested status.
7. Prepare a remediation plan with owners.
8. Run a tabletop incident handoff between service, identity and data owners.
9. Capture a closure note with residual limitations..

Expected deliverable and acceptance: Activity investigation, data-flow control map and compliance evidence matrix. Pass requires an evidence-backed reason for each status.

Troubleshooting: if the feature is absent, check role, subscription and rollout status using the linked source. If a saved setting appears ineffective, verify target identity and scope, then documented propagation. If access changes unexpectedly, restore the recorded baseline for the disposable lab scope and stop expanding the change.

Cleanup evidence: record objects removed, settings restored and any retained artifact or unresolved dependency. Never delete a protected record just to obtain a clean lab.

LAB-13 - AI economics, adoption and service health

Prerequisites and access: Billing or cost read access, usage-report access and Service health read access appropriate to the task.

1. Read the cost and adoption fixture.
2. Reconcile duplicated metered charges.
3. Calculate active-license rate and cost per accepted outcome.
4. State the reporting period and population.
5. Separate workload-level activity from whole-tenant availability.

6. Review the supplied service advisory and decide whether its timing supports the suspected effect.
7. Write a one-page operations report with a confirmed observation, a bounded inference and a next action.
8. Set a review date and owner.
9. Have a peer challenge any unsupported ROI claim..

Expected deliverable and acceptance: Reconciled cost ledger and operational report. Pass requires the correct denominator, overlap treatment and a distinction between adoption and value.

Troubleshooting: if the feature is absent, check role, subscription and rollout status using the linked source. If a saved setting appears ineffective, verify target identity and scope, then documented propagation. If access changes unexpectedly, restore the recorded baseline for the disposable lab scope and stop expanding the change.

Cleanup evidence: record objects removed, settings restored and any retained artifact or unresolved dependency. Never delete a protected record just to obtain a clean lab.

Read-only Graph PowerShell lab

Prerequisites: supported PowerShell, approved Microsoft Graph modules and consent for the read scope. Use a lab tenant, replace the tenant placeholder and review the scope before authenticating. User.Read.All is appropriate here for requested user properties; it is not claimed to be the least scope for every possible user query. This script reads and exports; it does not create, update or delete directory objects. The code is documentation-based and has not been executed against your tenant.

```
# If required, install approved modules using your organization's normal process.
# Install-Module Microsoft.Graph.Authentication -Scope CurrentUser
# Install-Module Microsoft.Graph.Users -Scope CurrentUser
$LabTenantId = "REPLACE-WITH-AUTHORIZED-LAB-TENANT-ID"
Connect-MgGraph -TenantId $LabTenantId -Scopes "User.Read.All" -ContextScope Process
Get-MgContext | Select-Object TenantId, Account, Scopes
# STOP if the tenant or account is not the approved lab identity.
Get-MgUser -All -Property Id,DisplayName,UserPrincipalName,UserType,AccountEnabled |
Select-Object Id,DisplayName,UserPrincipalName,UserType,AccountEnabled |
Export-Csv -Path ".\AB650-Lab-Users.csv" -NoTypeInformation -Encoding UTF8
Disconnect-MgGraph
```

Explain: Connect-MgGraph establishes the scoped session; Get-MgContext checks identity and tenant; Get-MgUser reads directory users; -All handles all returned pages; -Property requests selected properties; Select-Object controls output columns; Export-Csv creates local evidence; Disconnect-MgGraph ends the session. Inspect the CSV and compare record counts with a trusted reference. Store exports appropriately because real directory identifiers are personal or organizational data.

Source: <https://learn.microsoft.com/en-us/powershell/microsoftgraph/authentication-commands>
Source: <https://learn.microsoft.com/en-us/powershell/module/microsoft.graph.users/get-mguser>

Objective-to-lesson traceability

The identifiers below are LOCAL course pointers, not Microsoft-assigned objective codes. They locate the numbered bullet under each skills subsection in the current official study guide. For example, D1.T1.4 means domain 1, tenant subsection, bullet 4. W1 denotes workloads; D2 I1 identities, A1 authentication, S1 threat security, P1 Purview; D3 C1 Copilot, I1 agent implementation, G1 agent governance, M1 monitoring. Compound pointers mean one lesson covers more than one bullet. This preserves precise traceability without reproducing the blueprint. Foundation lessons are explicitly prerequisites; capstones and revision inherit the mapped lessons.

Blueprint pointer	Lessons / lab	Question set
D1.T1.1	L01 / LAB-01	Q001-Q005
D1.T1.2	L02 / LAB-01	Q006-Q010
D1.T1.3	L03 / LAB-01	Q011-Q015
D1.T1.4	L04 / LAB-01	Q016-Q020
D1.T1.5	L05 / LAB-02	Q021-Q025
D1.T1.6	L06 / LAB-02	Q026-Q030
D1.T1.7	L07 / LAB-02	Q031-Q035
D1.T1.8	L08 / LAB-02	Q036-Q040
D1.W1.1	L09 / LAB-03	Q041-Q045
D1.W1.2	L10 / LAB-03	Q046-Q050
D1.W1.3	L11 / LAB-03	Q051-Q055
D1.W1.4	L12 / LAB-03	Q056-Q060
D1.W1.5	L13 / LAB-03	Q061-Q065
D2.I1.1	L14 / LAB-04	Q066-Q070
D2.I1.2	L15 / LAB-04	Q071-Q075
D2.I1.3	L16 / LAB-04	Q076-Q080
D2.I1.4	L17 / LAB-04	Q081-Q085
D2.I1.5	L18 / LAB-04	Q086-Q090
D2.I1.6	L19 / LAB-04	Q091-Q095
D2.I1.7	L20 / LAB-04	Q096-Q100
D2.A1.1	L21 / LAB-05	Q101-Q105
D2.A1.2	L22 / LAB-05	Q106-Q110
D2.A1.3	L23 / LAB-05	Q111-Q115

Blueprint pointer	Lessons / lab	Question set
D2.A1.4	L24 / LAB-05	Q116-Q120
D2.A1.5	L25 / LAB-05, L26 / LAB-05	Q121-Q125, Q126-Q130
D2.S1.1	L27 / LAB-06	Q131-Q135
D2.S1.2	L28 / LAB-06, L29 / LAB-06	Q136-Q140, Q141-Q145
D2.S1.3	L30 / LAB-06	Q146-Q150
D2.S1.4	L31 / LAB-06	Q151-Q155
D2.P1.1	L34 / LAB-07, L35 / LAB-07	Q166-Q170, Q171-Q175
D2.P1.2	L32 / LAB-07, L33 / LAB-07	Q156-Q160, Q161-Q165
D2.P1.3	L36 / LAB-07	Q176-Q180
D2.P1.4	L37 / LAB-07	Q181-Q185
D2.P1.5	L38 / LAB-07	Q186-Q190
D3.C1.1	L39 / LAB-08, L42 / LAB-08	Q191-Q195, Q206-Q210
D3.C1.2	L40 / LAB-08	Q196-Q200
D3.C1.3	L43 / LAB-09	Q211-Q215
D3.C1.4	L44 / LAB-09	Q216-Q220
D3.C1.5	L45 / LAB-09	Q221-Q225
D3.C1.6	L41 / LAB-08	Q201-Q205
D3.C1.7	L46 / LAB-09	Q226-Q230
D3.C1.8	L47 / LAB-09	Q231-Q235
D3.C1.9	L47 / LAB-09	Q231-Q235
D3.I1.1	L48 / LAB-10	Q236-Q240
D3.I1.2	L49 / LAB-10, L51 / LAB-10	Q241-Q245, Q251-Q255
D3.I1.3	L50 / LAB-10	Q246-Q250
D3.I1.4	L52 / LAB-11	Q256-Q260
D3.I1.5	L53 / LAB-11	Q261-Q265
D3.I1.6	L54 / LAB-11	Q266-Q270
D3.I1.7	L55 / LAB-11	Q271-Q275
D3.I1.8	L55 / LAB-11	Q271-Q275
D3.G1.1	L56 / LAB-12	Q276-Q280
D3.G1.2	L57 / LAB-12	Q281-Q285
D3.G1.3	L58 / LAB-12	Q286-Q290

Blueprint pointer	Lessons / lab	Question set
D3.M1.1	L59 / LAB-13	Q291-Q295
D3.M1.2	L60 / LAB-13	Q296-Q300
D3.M1.3	L60 / LAB-13	Q296-Q300
D3.M1.4	L60 / LAB-13	Q296-Q300

Integrated enterprise case studies

Case A - Acquisition launch

Scenario: Northstar adds a 30-person division. Twenty users are direct members of a licensed group; ten are in a nested group. A verified domain must keep its old MX route until Friday. Finance has a broadly shared library and a shared support mailbox whose delegate cannot send as support.

Mapping: L01-L13; D1 and D3.C1.2

Task: Produce a launch checklist, identify three independent control failures, choose targeted corrections and prove continuity.

Model answer and rationale: Expected: preserve mail routing until cutover; repair direct license coverage; trace and remove unnecessary finance access; configure the specific mailbox send permission. Validate users, DNS and direct access separately. Reject changing MX early, assuming nested inheritance or using search exclusion as the complete access repair.

Case B - Identity and phishing incident

Scenario: A phishing message reaches six people. One clicks. A support administrator has a permanently active broad role. A risky sign-in coincides with the click, but compromise is unconfirmed. A second applicable policy blocks the resource even after MFA.

Mapping: L14-L31; D2.I1, D2.A1, D2.S1

Task: Write a time-ordered evidence table, targeted response, access prediction and privileged-role improvement.

Model answer and rationale: Expected: distinguish delivery, click and compromise; correlate sign-in evidence; predict the applicable block; record remediation outcomes; use suitable least privilege and time-bound activation. Do not infer six compromised devices or that MFA cancels blocking.

Case C - Sensitive records and AI pilot

Scenario: The pilot retrieves HR content via an unintended group grant. A regulated record requires seven years of preservation from creation, while the site requires three. A blocked DLP attempt is reported as a confirmed leak.

Mapping: L32-L42; D2.P1 and D3.C1

Task: Create a data remediation plan, retention timeline and corrected incident statement; decide whether to expand the pilot.

Model answer and rationale: Expected: fix source access; apply appropriate item retention and distribution; preserve the longer active retain requirement; distinguish prevented attempt from successful exposure; hold expansion until required controls pass. Do not trade a mandatory data control for high satisfaction.

Case D - Agent estate and economics

Scenario: A sponsor leaves. An agent blueprint has two current children and another planned. A separate user account exists. A registry tool is unexpectedly used. Costs show an overlapping 300 charge; 40 of 100 licensed users are active.

Mapping: L43-L60; D3

Task: Produce an identity diagram, lifecycle decision, tool containment, reconciled cost model and one-page operational review.

Model answer and rationale: Expected: transfer or retire sponsorship; distinguish identity, user account and registry ownership; use supported scope for future children and separately assess user-account policies; block the relevant capability and verify; remove duplicate charges; state 40% activity for the defined period without claiming proven ROI.

Case scoring: 20 points each: 4 for identifying the real requirement and scope, 4 for the appropriate control, 4 for rejecting tempting alternatives, 4 for positive and negative evidence, and 4 for rollback, ownership and limitations. Pass gate: at least 16/20 and no critical error such as knowingly broadening sensitive access, confusing report-only with enforcement or claiming untested behavior as verified.

Instructor and facilitator guide

Delivery options: independent reading with self-marking; 6-10 minute narrated micro-lessons recorded from the outline; or four optional 60-minute workshops. These are delivery designs, not claims of affiliation with Microsoft Learn, CBT Nuggets or Udemy. Use the original examples and your own recording assets.

Before delivery: confirm the blueprint and high-change feature pages; verify the training tenant and license matrix; preload synthetic fixtures; test demonstration identities; keep a fallback offline worksheet; separate learner and administrator accounts; prepare captions and readable screenshots. Do not promise attendees a free tenant or access to every feature.

Workshop	Preparation	60-minute agenda	Expected evidence
1: tenant and collaboration	M01-M03	10 retrieval; 15 demo; 20 paired diagnosis; 10 review; 5 next steps	Case A launch record
2: identity and threat response	M04-M06	10 retrieval; 15 log reading; 20 incident exercise; 10 challenge; 5 debrief	Case B evidence timeline
3: data and Copilot rollout	M07-M09	10 retrieval; 15 control comparison; 20 gate review; 10 feedback; 5 actions	Case C go/hold decision
4: agent operations	M10-M13	10 retrieval; 15 identity diagram; 20 operations review; 10 challenge; 5 exam plan	Case D operating model

Repeatable facilitation method: ask learners to predict before demonstrating; show one successful and one denied operation; ask which observation could falsify their explanation; reveal the model answer after discussion. If learners choose a broad role, ask for the exact operation and object scope. If they quote a dashboard, ask for the underlying event. Avoid turning the exercise into memorizing portal clicks.

Assess each module using its knowledge check, associated five-question lesson sets and lab rubric. Give feedback in requirement → choice → evidence → limitation order. Advanced extension: introduce a conflicting control, stale token, unavailable entitlement or disconnected source and ask what additional evidence is needed. Do not give an advanced learner permission to skip the negative test.

Quality assurance for a training provider: pilot the bank with independent learners, inspect item-level errors and ambiguity, recheck product facts and revise before making any validated-readiness claim. Track first-attempt exam results separately from course completion and repeated practice. No cohort outcome data exists for this package; its mastery thresholds are course design choices.

Slide deck outline and recording plan

Each module has eight proposed slides. This is the requested outline, not a completed branded slide deck. Record the theory segment and demonstration separately so volatile UI footage can be replaced. Add alt text to diagrams and captions to recordings.

Deck segment M01 - Tenant foundations and licensing

- Slide 01.1: Business problem - use the first lesson scenario; ask for an initial prediction.
- Slide 01.2: Concepts - explain Branding, Domain verification, Organization settings, Group licensing.
- Slide 01.3: Diagram - Create a two-layer tenant map: identity and administration above workload resources. Mark domain verification separately from mail-routing arrows.
- Slide 01.4: Decision method - compare the correct control with the lesson distractors.
- Slide 01.5: Demonstration - use LAB-01; show baseline, scope and result.
- Slide 01.6: Failure investigation - Confusing appearance with trust, a group assignment with every user entitlement, and a saved value with effective behavior.
- Slide 01.7: Retrieval check - Why can a verified domain still deliver mail elsewhere? Answer: verification proves ownership; routing follows the configured mail records.
- Slide 01.8: Practice and transfer - submit Tenant baseline, DNS change ticket and entitlement reconciliation. Pass when the trainee separates verification from routing and detects the nested-membership gap.

Deck segment M02 - Operational resilience and funding

- Slide 02.1: Business problem - use the first lesson scenario; ask for an initial prediction.
- Slide 02.2: Concepts - explain AI entitlements and funding, Backup and recovery, Network diagnosis, Service health.
- Slide 02.3: Diagram - Draw a fault tree with client, identity, network and cloud-service branches. Add a separate recovery timeline with incident time and chosen restore point.
- Slide 02.4: Decision method - compare the correct control with the lesson distractors.
- Slide 02.5: Demonstration - use LAB-02; show baseline, scope and result.
- Slide 02.6: Failure investigation - Confusing funding with permission, advisory scope with every failure, and retention with recovery.
- Slide 02.7: Retrieval check - Why can a successful restore job still fail business acceptance? Answer: it may restore the wrong point, wrong scope or incomplete expected content.

- Slide 02.8: Practice and transfer - submit Incident triage worksheet, restore acceptance record and funding matrix. Offline completion is a design exercise, not evidence of a live restore.

Deck segment M03 - Collaboration and safe discovery

- Slide 03.1: Business problem - use the first lesson scenario; ask for an initial prediction.
- Slide 03.2: Concepts - explain Shared mailbox permissions, Teams membership, Meeting Copilot and transcripts, SharePoint permissions, Discovery versus access.
- Slide 03.3: Diagram - Draw user → permission → resource, with a separate discovery branch. Make a table comparing mailbox read and send permissions.
- Slide 03.4: Decision method - compare the correct control with the lesson distractors.
- Slide 03.5: Demonstration - use LAB-03; show baseline, scope and result.
- Slide 03.6: Failure investigation - Assuming channel membership is universal, Full Access includes Send As, or transcription allowed means a saved transcript exists.
- Slide 03.7: Retrieval check - Can hiding a site in broad search remove a saved link permission? Answer: no; test and repair the underlying access separately.
- Slide 03.8: Practice and transfer - submit Permission matrix, two-identity access tests and meeting-state worksheet. Pass requires distinguishing discovery, authorization and transcript state.

Deck segment M04 - Identity administration at scale

- Slide 04.1: Business problem - use the first lesson scenario; ask for an initial prediction.
- Slide 04.2: Concepts - explain User lifecycle, Guests and external collaboration, Group purpose, Privileged roles and PIM, Administrative units, Contacts, Graph PowerShell bulk administration.
- Slide 04.3: Diagram - Draw joiner → mover → leaver with access and ownership checks at each transition; add capability, scope and time as three role dimensions.
- Slide 04.4: Decision method - compare the correct control with the lesson distractors.
- Slide 04.5: Demonstration - use LAB-04; show baseline, scope and result.
- Slide 04.6: Failure investigation - Equating contacts with sign-in users, confusing role eligibility with activation, and treating administrative units as separate tenants.
- Slide 04.7: Retrieval check - Why is a successful Graph connection not proof that 200 updates succeeded? Answer: it validates a session, not each target operation.
- Slide 04.8: Practice and transfer - submit Lifecycle checklist, delegated-scope test and read-only inventory. Pass requires explaining all exported fields and identifying the duplicate before any write.

Deck segment M05 - Authentication and access decisions

- Slide 05.1: Business problem - use the first lesson scenario; ask for an initial prediction.

- Slide 05.2: Concepts - explain Authentication methods, Password Protection, Self-service password reset, Sign-in investigation, Conditional Access interaction, Risk-based access.
- Slide 05.3: Diagram - Draw a policy decision tree: in scope? conditions match? report-only or enforced? applicable block? required controls satisfied?
- Slide 05.4: Decision method - compare the correct control with the lesson distractors.
- Slide 05.5: Demonstration - use LAB-05; show baseline, scope and result.
- Slide 05.6: Failure investigation - Treating enablement as registration, no prompt as no MFA, or the newest Conditional Access policy as the winner.
- Slide 05.7: Retrieval check - What does report-only failure establish? Answer: the policy evaluation indicates a failure under its logic, but report-only did not enforce blocking.
- Slide 05.8: Practice and transfer - submit Policy truth table and event-based explanation. Fail if the learner claims that report-only blocked a user or that one grant overrides a block.

Deck segment M06 - Threat protection and response

- Slide 06.1: Business problem - use the first lesson scenario; ask for an initial prediction.
- Slide 06.2: Concepts - explain Alert triage, Threat policy scope, Threat policy precedence, Threat investigation and response, Attack simulation training.
- Slide 06.3: Diagram - Draw an incident evidence chain from alert to message to affected identity to remediation result. Add a separate per-feature policy ladder.
- Slide 06.4: Decision method - compare the correct control with the lesson distractors.
- Slide 06.5: Demonstration - use LAB-06; show baseline, scope and result.
- Slide 06.6: Failure investigation - Assuming custom priority zero beats a preset, every recipient was compromised, or simulation delivery proves learning.
- Slide 06.7: Retrieval check - Why does closing an alert not prove containment? Answer: status may change without removing the message or addressing compromised access.
- Slide 06.8: Practice and transfer - submit Investigation timeline, policy-source analysis and simulation design. Live remediation or campaign delivery is not required for completion.

Deck segment M07 - Information protection and governance

- Slide 07.1: Business problem - use the first lesson scenario; ask for an initial prediction.
- Slide 07.2: Concepts - explain Classification techniques, Sensitivity labels and publishing, DLP requirements, Copilot DLP boundaries, Retention design, DLP alert response, DSPM for AI.
- Slide 07.3: Diagram - Draw five labeled boxes: classify, protect, restrict action, preserve, recover. Connect each to a different business requirement.
- Slide 07.4: Decision method - compare the correct control with the lesson distractors.
- Slide 07.5: Demonstration - use LAB-07; show baseline, scope and result.

- Slide 07.6: Failure investigation - Confusing sensitivity with retention, label creation with publication, a citation with processing, and posture visibility with remediation.
- Slide 07.7: Retrieval check - Why is a DLP alert not automatically proof of a leak? Answer: it may describe a prevented attempt; inspect the action outcome.
- Slide 07.8: Practice and transfer - submit Classification rationale, six-location DLP matrix, alert triage and retention timeline. Pass requires separate allowed and denied test cases.

Deck segment M08 - Copilot readiness and rollout

- Slide 08.1: Business problem - use the first lesson scenario; ask for an initial prediction.
- Slide 08.2: Concepts - explain Copilot readiness, Data readiness, User experiences, Phased rollout decisions.
- Slide 08.3: Diagram - Draw a launch gate with technical readiness, data readiness, support readiness and task-quality evidence all feeding the decision.
- Slide 08.4: Decision method - compare the correct control with the lesson distractors.
- Slide 08.5: Demonstration - use LAB-08; show baseline, scope and result.
- Slide 08.6: Failure investigation - Calling interest adoption, assuming AI creates new permission grants, and averaging mandatory control failures into a good score.
- Slide 08.7: Retrieval check - Why is license assignment insufficient for launch? Answer: it proves neither supported experiences nor safe data access or useful outcomes.
- Slide 08.8: Practice and transfer - submit Readiness register, task-quality evaluation and launch decision. A useful answer must show the source checking performed by the learner.

Deck segment M09 - Copilot controls and integrations

- Slide 09.1: Business problem - use the first lesson scenario; ask for an initial prediction.
- Slide 09.2: Concepts - explain Copilot web search, Copilot Search, Tenant Copilot settings, Cowork administration, Providers and connectors.
- Slide 09.3: Diagram - Draw separate lanes for experience visibility, access grant, model selection, data retrieval and billing. Add indexed versus runtime connector flows.
- Slide 09.4: Decision method - compare the correct control with the lesson distractors.
- Slide 09.5: Demonstration - use LAB-09; show baseline, scope and result.
- Slide 09.6: Failure investigation - Treating disclaimers as enforcement, hidden features as revoked access, tiny budgets as no access, and Ready status as verified isolation.
- Slide 09.7: Retrieval check - Why does no indexing not mean no data transfer? Answer: a runtime connector still retrieves information for the requested operation.
- Slide 09.8: Practice and transfer - submit Control register, provider review and connector access-test record. Do not claim a live integration was tested when only the fixture was used.

Deck segment M10 - Agent identity and access

- Slide 10.1: Business problem - use the first lesson scenario; ask for an initial prediction.
- Slide 10.2: Concepts - explain Agent identity lifecycle, Agent access governance, Agent ownership layers, Agent Conditional Access.
- Slide 10.3: Diagram - Draw sponsor and technical owner connected to a blueprint, two child identities and one distinct agent user account. Mark the policy subject for each request.
- Slide 10.4: Decision method - compare the correct control with the lesson distractors.
- Slide 10.5: Demonstration - use LAB-10; show baseline, scope and result.
- Slide 10.6: Failure investigation - Confusing a child with its blueprint, registry owner with sponsor, and identity-targeted policy with user-account coverage.
- Slide 10.7: Retrieval check - Why cannot an agent identity complete human MFA? Answer: the noninteractive subject requires controls supported for that identity type.
- Slide 10.8: Practice and transfer - submit Identity diagram, ownership matrix, lifecycle workflow and scoped policy evidence. Pass requires explicitly separating agent identities from agent user accounts.

Deck segment M11 - Agent registry and tool governance

- Slide 11.1: Business problem - use the first lesson scenario; ask for an initial prediction.
- Slide 11.2: Concepts - explain Agent settings, Agent inventory, Agent publication review, Registry actions and tools.
- Slide 11.3: Diagram - Draw request → review → approve, return or reject, then show agent access and tool access as separate gates.
- Slide 11.4: Decision method - compare the correct control with the lesson distractors.
- Slide 11.5: Demonstration - use LAB-11; show baseline, scope and result.
- Slide 11.6: Failure investigation - Equating installed with approved for everyone, registry inventory with all possible agents, or view permission with management permission.
- Slide 11.7: Retrieval check - Why is a registry owner not proof that the agent has safe tools? Answer: accountability and permitted capabilities are different properties.
- Slide 11.8: Practice and transfer - submit Reconciled registry, publication decision and tool-control test. Pass requires matching containment to the actual capability.

Deck segment M12 - Agent security and compliance operations

- Slide 12.1: Business problem - use the first lesson scenario; ask for an initial prediction.
- Slide 12.2: Concepts - explain Agent activity monitoring, Agent data protection, Agent compliance gaps.

- Slide 12.3: Diagram - Draw source → agent → tool → destination, with identity and data-control checkpoints; add an evidence-to-requirement compliance table.
- Slide 12.4: Decision method - compare the correct control with the lesson distractors.
- Slide 12.5: Demonstration - use LAB-12; show baseline, scope and result.
- Slide 12.6: Failure investigation - Treating approval as permanent safety, prompt wording as enforcement and product availability as organizational compliance.
- Slide 12.7: Retrieval check - Why is a tenfold activity spike not proof of compromise? Answer: legitimate demand, retries or a loop can also explain it.
- Slide 12.8: Practice and transfer - submit Activity investigation, data-flow control map and compliance evidence matrix. Pass requires an evidence-backed reason for each status.

Deck segment M13 - AI economics, adoption and service health

- Slide 13.1: Business problem - use the first lesson scenario; ask for an initial prediction.
- Slide 13.2: Concepts - explain AI cost reconciliation, Adoption and AI service health.
- Slide 13.3: Diagram - Draw three columns: entitlement, activity and accepted outcomes. Add a cost-reconciliation table with overlap removal.
- Slide 13.4: Decision method - compare the correct control with the lesson distractors.
- Slide 13.5: Demonstration - use LAB-13; show baseline, scope and result.
- Slide 13.6: Failure investigation - Double-counting cost surfaces, comparing unequal report windows and equating prompts with productivity.
- Slide 13.7: Retrieval check - What is adoption for 40 active users among 100 licensed users? Answer: 40 percent for that defined population and period; it is not a value claim.
- Slide 13.8: Practice and transfer - submit Reconciled cost ledger and operational report. Pass requires the correct denominator, overlap treatment and a distinction between adoption and value.

Revision guide and exam priorities

Microsoft does not publish the frequency of individual live questions. The priorities below are inferred from domain weight and conceptual risk, not leaked or measured item frequency. Security/governance and AI deserve most study time. Do not omit the smaller tenant domain because its concepts underpin later scenarios.

Priority	Why candidates struggle	Revision action
Control boundaries	Access, discovery, licensing, billing and labels can look like interchangeable toggles.	Explain each boundary with an allowed and denied example.
Policy interaction	Candidates reason from one policy or one prompt rather than effective evaluation.	Work policy truth tables and inspect event evidence.
Purview requirements	Classification, protection, DLP, retention and recovery answer different questions.	Choose data shape, location, action and time before naming a feature.
Agent subject and ownership	Blueprints, child identities, user accounts and registry owners are distinct.	Draw objects and token subjects before scoping policies.
New product behavior	Cowork controls, connectors and agent capabilities can evolve quickly.	Recheck current product pages and note the verified date.
Operational evidence	Saved configuration or a green dashboard is mistaken for an outcome.	Collect positive, negative and post-change evidence.

Beginner challenges: vocabulary overload, understanding workload boundaries, interpreting role scope, reading logs and distinguishing authentication from authorization. Remediation: redo the foundation bridge and explain a single scenario without product jargon. Advanced concepts: multi-policy evaluation, identity-layer targeting, data-flow enforcement, propagation effects, mixed billing reconciliation and compliance evidence. Extension: analyze a failure where the first plausible explanation is wrong.

Exam cram - last review session

- A license permits a product; a role permits administration; a billing arrangement funds supported consumption.
- Verify direct group membership and effective user assignments; do not assume nested licensing.
- TXT ownership verification is separate from MX cutover.
- Mailbox read and send permissions are distinct; channel membership can differ from team membership.
- Allowed transcription is not an existing transcript; verify the meeting state.
- Missing discovery is not proof of denied access. Test a direct URL with the right identity.
- Eligible PIM role is not an active permanent assignment.
- An applicable block survives a successful MFA requirement elsewhere.

- No fresh MFA prompt can still be consistent with valid policy enforcement.
- Identify the threat-protection feature and applicable policy source before comparing priority.
- Classification identifies content; labels configure classification/protection; DLP governs supported actions; retention preserves/disposes; backup recovers.
- A blocked DLP attempt is not automatically a successful leak.
- Correct source oversharing before expanding AI retrieval.
- Agent blueprint, child identity and agent user account are different policy subjects.
- Inventory, approval, access and tool invocation each need their own evidence.
- Cowork access and visibility are separate; verify current spending-policy behavior.
- No indexing does not mean no runtime transfer; test connector identities and permissions.
- Count overlapping charges once and define adoption denominators.
- Read the question for scope, required action, constraints and evidence; reject an option that solves a different problem.
- Review actual appointment instructions. Do not infer exam timing, lab presence or allowed resources from this workbook.

30-day study plan - 60 hours

Day	Two-hour focus	Output
1	Foundation	F01-F04; glossary and object map
2	M01	L01-L02; domain and branding exercises
3	M01 / M02	L03-L05; licensing and funding
4	M02	L06-L08; health and recovery
5	M03	L09-L10; mailbox and channels
6	M03	L11-L13; transcript and discovery
7	M04	L14-L16; lifecycle, guests and groups
8	M04	L17-L19; privilege, units and contacts
9	M04 / M05	L20-L21; Graph and methods
10	M05	L22-L23; password and reset
11	M05	L24-L26; sign-ins and access
12	M06	L27-L29; alerts and policies
13	M06	L30-L31; response and training
14	M07	L32-L33; classification and labels
15	M07	L34-L35; DLP boundaries
16	M07	L36-L38; retention, alerts and posture
17	M08	L39-L40; readiness and data
18	M08	L41-L42; experiences and pilot gate
19	M09	L43-L44; search controls
20	M09	L45-L47; settings, Cowork and integration
21	M10	L48-L49; lifecycle and access
22	M10	L50-L51; ownership and policy subjects
23	M11	L52-L53; settings and inventory
24	M11	L54-L55; publication and tools
25	M12	L56-L58; monitoring and compliance
26	M13	L59-L60; economics and health
27	Remediation	Complete unfinished labs; retest weakest boundary distinctions

Day	Two-hour focus	Output
28	Integrated assessment	Cases A and B; use closed-book decisions and justify alternatives
29	Integrated assessment	Cases C and D; independent peer or instructor scoring
30	Readiness review	Cram guide; timed unseen Learn exercises where available; booking-rule check

The first 55 scheduled hours cover foundation and modules, including embedded revision. The final five hours across days 28-30 are the formal case and readiness assessment; the remaining hour revisits weak module material. On each ordinary day use roughly 45 minutes of explanation, 45 of practical work and 30 of retrieval or error-log review; adjust to the task. Workshops are optional additional time. If a gate fails, extend the calendar rather than cram through it.

Readiness decision

Local mastery gates: at least 95% overall and at least 90% in each domain on both mixed checks, completed on separate days without notes or prior answer-key review; at least 16/20 on each integrated case; satisfactory lab evidence; and correct explanation of critical control distinctions. A 30-item check requires 29 correct answers to exceed 95%. Repeated or previously exposed questions are retrieval practice, not independent evidence. These gates are course design choices, not a validated exam prediction. Report live lab competence separately from fixture reasoning.

Flashcards - 60 prompts

FC01 - Branding

Front: What distinction matters most, and how would you prove the result?

Back: Branding gives users recognizable organization and sign-in experiences; it does not establish trust in a message or authorize access. Evidence: Capture the organization experience and a fresh sign-in view using the pilot identity, checking the intended URLs and imagery.

FC02 - Domain verification

Front: What distinction matters most, and how would you prove the result?

Back: Domain ownership verification and mail routing are separate decisions. Proving control of a namespace does not itself move email. Evidence: Check verified domain status and independently confirm the production MX destination still matches the approved pre-cutover configuration.

FC03 - Organization settings

Front: What distinction matters most, and how would you prove the result?

Back: A tenant contains organization-wide service choices and profile information; each setting needs a business owner and an understood scope. Evidence: Compare the saved values and effective experience against the approved requirements, including any propagation delay.

FC04 - Group licensing

Front: What distinction matters most, and how would you prove the result?

Back: A license is an entitlement assembled from service plans. Group assignment must be verified for the intended users rather than inferred from membership somewhere in a hierarchy. Evidence: Inspect direct membership, user-level entitlement sources, service-plan state and assignment errors for the ten affected accounts.

FC05 - AI entitlements and funding

Front: What distinction matters most, and how would you prove the result?

Back: Product access, permission to create agents and the mechanism that pays for consumption answer different questions. Evidence: Produce a matrix of persona, product, role, billing source and available capacity, verified against current subscription details.

FC06 - Backup and recovery

Front: What distinction matters most, and how would you prove the result?

Back: Backup prepares a recoverable copy or restore point for supported workloads; retention governs preservation and deletion obligations. Evidence: Compare restored test files with known expected contents, review job status and document missing or changed items.

FC07 - Network diagnosis

Front: What distinction matters most, and how would you prove the result?

Back: Connectivity is a path from client through local network and security intermediaries to a cloud service. A healthy service can still be unreachable from one office. Evidence: Correlate timestamps, connectivity results and proxy observations, then retest the same pilot operation after the targeted fix.

FC08 - Service health

Front: What distinction matters most, and how would you prove the result?

Back: Service health reports provider incidents and advisories; local symptom reports still need correlation before attribution. Evidence: Record matching or nonmatching workload, region, time window and incident details alongside pilot tests.

FC09 - Shared mailbox permissions

Front: What distinction matters most, and how would you prove the result?

Back: Opening a mailbox and sending with its identity require different delegated permissions. Evidence: Send a synthetic message through the delegate and inspect the recipient-visible sender, while confirming existing read access.

FC10 - Teams membership

Front: What distinction matters most, and how would you prove the result?

Back: Teams owners manage the workspace; channel membership and associated content boundaries depend on the channel type. Evidence: Test discussion and associated file access with an included and an excluded pilot identity.

FC11 - Meeting Copilot and transcripts

Front: What distinction matters most, and how would you prove the result?

Back: Meeting policy, organizer options and the existence of a transcript jointly determine the experience. Allowing transcription is not evidence that it ran. Evidence: Check effective meeting settings and the actual saved transcript and post-meeting experience with an eligible licensed participant.

FC12 - SharePoint permissions

Front: What distinction matters most, and how would you prove the result?

Back: Access is the result of effective permissions, group membership and sharing links, not the folder's reassuring name. Evidence: Use direct URLs with a finance identity and a nonfinance identity, recording an allow and a deny after the change.

FC13 - Discovery versus access

Front: What distinction matters most, and how would you prove the result?

Back: Restricted discovery reduces broad surfacing; access restrictions decide who can open content. These controls need different tests. Evidence: Record direct access and broad discovery separately; account for indexing delay and documented discovery exceptions.

FC14 - User lifecycle

Front: What distinction matters most, and how would you prove the result?

Back: A user identity needs controlled creation, change and retirement. Departure handling must consider sessions, data and business ownership. Evidence: Verify access revocation and successor access separately, with a record of preserved data and remaining dependencies.

FC15 - Guests and external collaboration

Front: What distinction matters most, and how would you prove the result?

Back: A guest can participate through a governed external identity; external chat capability is not equivalent to membership or file access. Evidence: Test the intended resource and a prohibited resource with the guest, and record the review date.

FC16 - Group purpose

Front: What distinction matters most, and how would you prove the result?

Back: Security groups and Microsoft 365 groups support different management and collaboration needs. Membership changes can affect multiple connected resources. Evidence: Inspect the group type, owners, membership and the expected connected resources.

FC17 - Privileged roles and PIM

Front: What distinction matters most, and how would you prove the result?

Back: Least privilege limits capability; time-bound activation reduces how long privileged capability remains available. Evidence: Inspect activation duration, approval or justification evidence and the role state after expiry.

FC18 - Administrative units

Front: What distinction matters most, and how would you prove the result?

Back: An administrative unit can scope supported delegated directory administration; it is not a separate tenant or a universal data boundary. Evidence: Test a permitted operation on an in-scope test user and its rejection on an out-of-scope test user.

FC19 - Contacts

Front: What distinction matters most, and how would you prove the result?

Back: A mail contact represents an external recipient in the directory; it is not an interactive employee account. Evidence: Verify address-book resolution and the external delivery address without creating sign-in access.

FC20 - Graph PowerShell bulk administration

Front: What distinction matters most, and how would you prove the result?

Back: Graph PowerShell exposes directory operations as commands. Safe bulk work uses explicit scope, validated input and per-record results. Evidence: Reconcile input targets with successful, failed and skipped records, and sample the resulting directory state.

FC21 - Authentication methods

Front: What distinction matters most, and how would you prove the result?

Back: Authentication establishes identity. Method availability, registration and actual use must be distinguished. Evidence: Inspect the effective method policy, registration status and authentication details for the pilot sign-in.

FC22 - Password Protection

Front: What distinction matters most, and how would you prove the result?

Back: Password Protection reduces weak password choices using prohibited-password evaluation; it does not replace multifactor authentication. Evidence: Record that prohibited test choices are rejected and acceptable synthetic choices work, without logging passwords.

FC23 - Self-service password reset

Front: What distinction matters most, and how would you prove the result?

Back: SSPR requires an eligible, in-scope user with sufficient usable registered methods. Configuration and readiness are separate. Evidence: Complete a controlled reset for the test identity and inspect the reset or audit evidence.

FC24 - Sign-in investigation

Front: What distinction matters most, and how would you prove the result?

Back: A missing MFA prompt does not establish a bypass; an existing authentication claim can satisfy the requirement. Evidence: Capture the relevant sign-in ID and explain each applicable policy's result and how authentication was satisfied.

FC25 - Conditional Access interaction

Front: What distinction matters most, and how would you prove the result?

Back: Applicable Conditional Access requirements are cumulative; an applicable block is not canceled by another policy's grant. Evidence: Inspect the final blocked result and the contribution of each applicable policy for the same token request.

FC26 - Risk-based access

Front: What distinction matters most, and how would you prove the result?

Back: Sign-in risk concerns a particular authentication event; user risk concerns potential compromise of the identity. Remediation must match the signal and supported controls. Evidence: Record the risk type, matched policy, intended action and observed test outcome without assuming every risk means confirmed compromise.

FC27 - Alert triage

Front: What distinction matters most, and how would you prove the result?

Back: An alert is a signal requiring assessment; its status is not the same as proof that the underlying threat was removed. Evidence: Link the alert to message or incident evidence and the recorded remediation result.

FC28 - Threat policy scope

Front: What distinction matters most, and how would you prove the result?

Back: Threat-protection settings must match the feature and recipients. A policy's name does not prove effective coverage. Evidence: Confirm the applicable protection feature, protected identities or domains where relevant and effective recipients.

FC29 - Threat policy precedence

Front: What distinction matters most, and how would you prove the result?

Back: Custom policy priority orders comparable custom policies; applicable preset protection can take precedence for a feature. Evidence: Identify the actual policy source for the feature and explain why it wins for that recipient.

FC30 - Threat investigation and response

Front: What distinction matters most, and how would you prove the result?

Back: Investigation connects evidence to affected assets; response must record what action was taken and whether it worked. Evidence: Correlate message identifiers, recipient evidence, click or sign-in signals where available and the result of each remediation action.

FC31 - Attack simulation training

Front: What distinction matters most, and how would you prove the result?

Back: A simulation is a controlled learning activity with approved scope and follow-up; it is not permission to send uncontrolled phishing. Evidence: Compare reporting and repeat-behavior measures across the planned exercise and follow-up, noting cohort differences.

FC32 - Classification techniques

Front: What distinction matters most, and how would you prove the result?

Back: Pattern detection, matching known records and recognizing document meaning address different classification problems. Evidence: Test known records, lookalike nonrecords and realistic benign content, recording false positives and false negatives.

FC33 - Sensitivity labels and publishing

Front: What distinction matters most, and how would you prove the result?

Back: A sensitivity label describes and can protect content; a publishing policy makes labels available to the intended users. Encryption depends on configuration. Evidence: Use an in-scope account to label test content and an excluded account to verify the intended publishing boundary.

FC34 - DLP requirements

Front: What distinction matters most, and how would you prove the result?

Back: A DLP requirement connects a location, detection condition and supported action. One workload's action should not be assumed for all others. Evidence: Test matching external mail, permitted internal mail and benign external mail; inspect policy events and user experience.

FC35 - Copilot DLP boundaries

Front: What distinction matters most, and how would you prove the result?

Back: Copilot DLP must use supported scenario-specific conditions and actions; it does not automatically govern every agent interaction. Evidence: Record the relevant policy scope, labelled test item and processing behavior; distinguish a citation from use of item content.

FC36 - Retention design

Front: What distinction matters most, and how would you prove the result?

Back: Retention rules govern preservation and disposal over time. Labels can express item-specific needs within a broader location policy. Evidence: Document the record's applicable retain actions, start event and expected earliest ordinary disposal time.

FC37 - DLP alert response

Front: What distinction matters most, and how would you prove the result?

Back: A DLP alert needs context about policy, item, actor and action before a responder can judge exposure and tune controls. Evidence: Record attempted versus completed activity, blocked versus allowed outcome and any justified follow-up.

FC38 - DSPM for AI

Front: What distinction matters most, and how would you prove the result?

Back: Data security posture management identifies and prioritizes data risks; findings need owners, remediation and later verification. Evidence: Compare targeted before-and-after access findings and direct tests, noting assessment refresh timing.

FC39 - Copilot readiness

Front: What distinction matters most, and how would you prove the result?

Back: Readiness combines identity, entitlement, supported clients, service configuration and operational support. A purchased seat is only one component. Evidence: Maintain a pass, fail or untested checklist for prerequisites and complete a representative pilot task.

FC40 - Data readiness

Front: What distinction matters most, and how would you prove the result?

Back: AI can make already-accessible information easier to find; existing oversharing therefore matters before rollout. Evidence: Prove the affected user's direct access changes as intended and then retest representative retrieval after propagation.

FC41 - User experiences

Front: What distinction matters most, and how would you prove the result?

Back: Adoption support should target actual job tasks, supported experiences and user feedback while preserving controls. Evidence: Review workload-level activity together with task completion and user-reported quality, not merely seat counts.

FC42 - Phased rollout decisions

Front: What distinction matters most, and how would you prove the result?

Back: A pilot is an experiment with acceptance criteria, not a ceremony before a predetermined launch. Evidence: Re-run the failed access and escalation scenarios and record approval against the same launch criteria.

FC43 - Copilot web search

Front: What distinction matters most, and how would you prove the result?

Back: Web grounding and work-data grounding have different information boundaries. Policy scope must be checked in the effective experience. Evidence: Inspect the effective policy assignment and test relevant modes after policy retrieval, recording uncertainty where runtime output is inconclusive.

FC44 - Copilot Search

Front: What distinction matters most, and how would you prove the result?

Back: Search configuration influences finding permitted information; permission and source quality still govern useful results. Evidence: Compare a known permitted item and a known denied item, documenting source access and discovery independently.

FC45 - Tenant Copilot settings

Front: What distinction matters most, and how would you prove the result?

Back: AI tenant settings need distinct decisions for purchasing, admin experiences, releases, disclaimers and generation features. Evidence: Capture the effective settings and test the intended user experiences; note any setting that lacks the desired granularity.

FC46 - Cowork administration

Front: What distinction matters most, and how would you prove the result?

Back: Cowork access, discovery, model availability and web-search behavior are separate controls. Current access administration must be checked against current documentation. Evidence: Review all applicable spending policies plus product prerequisites, and distinguish effective access from discoverability.

FC47 - Providers and connectors

Front: What distinction matters most, and how would you prove the result?

Back: External AI and source integrations need explicit data-flow and identity analysis; connector models differ in indexing and query-time retrieval. Evidence: Trace a synthetic record through the documented flow and test both authorized retrieval and the source or indexed access boundary.

FC48 - Agent identity lifecycle

Front: What distinction matters most, and how would you prove the result?

Back: An agent identity needs accountable ownership from creation through retirement. Its sponsor's departure is an operational event. Evidence: Inspect the resulting sponsor or retirement

state and verify required access changes and audit evidence.

FC49 - Agent access governance

Front: What distinction matters most, and how would you prove the result?

Back: An agent's access should match its task and resource boundaries; supported access packages and policy controls need deliberate targeting. Evidence: Validate a required operation and a prohibited operation under the actual agent identity, then confirm expiry or removal behavior.

FC50 - Agent ownership layers

Front: What distinction matters most, and how would you prove the result?

Back: Business sponsorship, technical identity ownership and registry ownership need not be the same person or object. Evidence: Produce an ownership matrix tied to the relevant objects and verify that each accountable person accepts the responsibility.

FC51 - Agent Conditional Access

Front: What distinction matters most, and how would you prove the result?

Back: Agent-token policy evaluation depends on the subject, resource and supported controls; agents cannot satisfy a human authentication prompt. Evidence: Use fresh token-request evidence for included and excluded identities after propagation, and assess the separate user-account policy.

FC52 - Agent settings

Front: What distinction matters most, and how would you prove the result?

Back: Tenant agent settings shape permitted creation, sharing, templates and user availability; these are governance choices with different scopes. Evidence: Record effective settings and both permitted and prohibited sharing or access attempts.

FC53 - Agent inventory

Front: What distinction matters most, and how would you prove the result?

Back: A registry enables discovery and curation, but inventory visibility depends on integration and discovery coverage. Evidence: Maintain a reconciled inventory noting source platform, registry state, owner and uncovered agents.

FC54 - Agent publication review

Front: What distinction matters most, and how would you prove the result?

Back: Publication is a governance decision that should evaluate purpose, ownership, permissions, data flows and evidence. Evidence: Keep the review decision, required changes and verification evidence attached to the specific version or request.

FC55 - Registry actions and tools

Front: What distinction matters most, and how would you prove the result?

Back: Agent availability and the tools it can invoke are related but distinct control surfaces.

Evidence: Observe the blocked invocation and continued permitted behavior, with corresponding governance records.

FC56 - Agent activity monitoring

Front: What distinction matters most, and how would you prove the result?

Back: Monitoring links an agent's identity and actions to expected behavior, owners and anomalies. Counts alone do not establish safety. Evidence: Document the activity pattern and root-cause evidence, plus the effect of any containment or correction.

FC57 - Agent data protection

Front: What distinction matters most, and how would you prove the result?

Back: Agent data protection must match the actual workload and action path; a broad AI label is not a universal enforcement guarantee. Evidence: Test a synthetic sensitive outbound operation and a benign permitted operation under the actual agent execution path.

FC58 - Agent compliance gaps

Front: What distinction matters most, and how would you prove the result?

Back: A compliance finding is a gap between required controls and observed evidence, not merely an unfavorable dashboard color. Evidence: Link each requirement to a configured control, tested outcome, owner and unresolved limitation.

FC59 - AI cost reconciliation

Front: What distinction matters most, and how would you prove the result?

Back: Cost reporting needs consistent periods, billing boundaries and units; overlapping views must not be added as separate charges. Evidence: Produce a reconciled ledger with unique charges, overlaps, excluded taxes or credits and the reporting period.

FC60 - Adoption and AI service health

Front: What distinction matters most, and how would you prove the result?

Back: Usage, useful outcomes and service reliability are different measures. Healthy adoption reporting defines the population and time window. Evidence: Reconcile the report population and period, compare workload-level use and inspect matching service-health evidence.

Glossary - 60 working terms

AI cost reconciliation - Cost reporting needs consistent periods, billing boundaries and units; overlapping views must not be added as separate charges.

AI entitlements and funding - Product access, permission to create agents and the mechanism that pays for consumption answer different questions.

Administrative units - An administrative unit can scope supported delegated directory administration; it is not a separate tenant or a universal data boundary.

Adoption and AI service health - Usage, useful outcomes and service reliability are different measures. Healthy adoption reporting defines the population and time window.

Agent Conditional Access - Agent-token policy evaluation depends on the subject, resource and supported controls; agents cannot satisfy a human authentication prompt.

Agent access governance - An agent's access should match its task and resource boundaries; supported access packages and policy controls need deliberate targeting.

Agent activity monitoring - Monitoring links an agent's identity and actions to expected behavior, owners and anomalies. Counts alone do not establish safety.

Agent compliance gaps - A compliance finding is a gap between required controls and observed evidence, not merely an unfavorable dashboard color.

Agent data protection - Agent data protection must match the actual workload and action path; a broad AI label is not a universal enforcement guarantee.

Agent identity lifecycle - An agent identity needs accountable ownership from creation through retirement. Its sponsor's departure is an operational event.

Agent inventory - A registry enables discovery and curation, but inventory visibility depends on integration and discovery coverage.

Agent ownership layers - Business sponsorship, technical identity ownership and registry ownership need not be the same person or object.

Agent publication review - Publication is a governance decision that should evaluate purpose, ownership, permissions, data flows and evidence.

Agent settings - Tenant agent settings shape permitted creation, sharing, templates and user availability; these are governance choices with different scopes.

Alert triage - An alert is a signal requiring assessment; its status is not the same as proof that the underlying threat was removed.

Attack simulation training - A simulation is a controlled learning activity with approved scope and follow-up; it is not permission to send uncontrolled phishing.

Authentication methods - Authentication establishes identity. Method availability, registration and actual use must be distinguished.

Backup and recovery - Backup prepares a recoverable copy or restore point for supported workloads; retention governs preservation and deletion obligations.

Branding - Branding gives users recognizable organization and sign-in experiences; it does not establish trust in a message or authorize access.

Classification techniques - Pattern detection, matching known records and recognizing document meaning address different classification problems.

Conditional Access interaction - Applicable Conditional Access requirements are cumulative; an applicable block is not canceled by another policy's grant.

Contacts - A mail contact represents an external recipient in the directory; it is not an interactive employee account.

Copilot DLP boundaries - Copilot DLP must use supported scenario-specific conditions and actions; it does not automatically govern every agent interaction.

Copilot Search - Search configuration influences finding permitted information; permission and source quality still govern useful results.

Copilot readiness - Readiness combines identity, entitlement, supported clients, service configuration and operational support. A purchased seat is only one component.

Copilot web search - Web grounding and work-data grounding have different information boundaries. Policy scope must be checked in the effective experience.

Cowork administration - Cowork access, discovery, model availability and web-search behavior are separate controls. Current access administration must be checked against current documentation.

DLP alert response - A DLP alert needs context about policy, item, actor and action before a responder can judge exposure and tune controls.

DLP requirements - A DLP requirement connects a location, detection condition and supported action. One workload's action should not be assumed for all others.

DSPM for AI - Data security posture management identifies and prioritizes data risks; findings need owners, remediation and later verification.

Data readiness - AI can make already-accessible information easier to find; existing oversharing therefore matters before rollout.

Discovery versus access - Restricted discovery reduces broad surfacing; access restrictions decide who can open content. These controls need different tests.

Domain verification - Domain ownership verification and mail routing are separate decisions. Proving control of a namespace does not itself move email.

Graph PowerShell bulk administration - Graph PowerShell exposes directory operations as commands. Safe bulk work uses explicit scope, validated input and per-record results.

Group licensing - A license is an entitlement assembled from service plans. Group assignment must be verified for the intended users rather than inferred from membership somewhere in a hierarchy.

Group purpose - Security groups and Microsoft 365 groups support different management and collaboration needs. Membership changes can affect multiple connected resources.

Guests and external collaboration - A guest can participate through a governed external identity; external chat capability is not equivalent to membership or file access.

Meeting Copilot and transcripts - Meeting policy, organizer options and the existence of a transcript jointly determine the experience. Allowing transcription is not evidence that it ran.

Network diagnosis - Connectivity is a path from client through local network and security intermediaries to a cloud service. A healthy service can still be unreachable from one office.

Organization settings - A tenant contains organization-wide service choices and profile information; each setting needs a business owner and an understood scope.

Password Protection - Password Protection reduces weak password choices using prohibited-password evaluation; it does not replace multifactor authentication.

Phased rollout decisions - A pilot is an experiment with acceptance criteria, not a ceremony before a predetermined launch.

Privileged roles and PIM - Least privilege limits capability; time-bound activation reduces how long privileged capability remains available.

Providers and connectors - External AI and source integrations need explicit data-flow and identity analysis; connector models differ in indexing and query-time retrieval.

Registry actions and tools - Agent availability and the tools it can invoke are related but distinct control surfaces.

Retention design - Retention rules govern preservation and disposal over time. Labels can express item-specific needs within a broader location policy.

Risk-based access - Sign-in risk concerns a particular authentication event; user risk concerns potential compromise of the identity. Remediation must match the signal and supported controls.

Self-service password reset - SSPR requires an eligible, in-scope user with sufficient usable registered methods. Configuration and readiness are separate.

Sensitivity labels and publishing - A sensitivity label describes and can protect content; a publishing policy makes labels available to the intended users. Encryption depends on configuration.

Service health - Service health reports provider incidents and advisories; local symptom reports still need correlation before attribution.

SharePoint permissions - Access is the result of effective permissions, group membership and sharing links, not the folder's reassuring name.

Shared mailbox permissions - Opening a mailbox and sending with its identity require different delegated permissions.

Sign-in investigation - A missing MFA prompt does not establish a bypass; an existing authentication claim can satisfy the requirement.

Teams membership - Teams owners manage the workspace; channel membership and associated content boundaries depend on the channel type.

Tenant Copilot settings - AI tenant settings need distinct decisions for purchasing, admin experiences, releases, disclaimers and generation features.

Threat investigation and response - Investigation connects evidence to affected assets; response must record what action was taken and whether it worked.

Threat policy precedence - Custom policy priority orders comparable custom policies; applicable preset protection can take precedence for a feature.

Threat policy scope - Threat-protection settings must match the feature and recipients. A policy's name does not prove effective coverage.

User experiences - Adoption support should target actual job tasks, supported experiences and user feedback while preserving controls.

User lifecycle - A user identity needs controlled creation, change and retirement. Departure handling must consider sessions, data and business ownership.

Source register and update checklist

Source history: the original lesson source register was assembled on 6 October 2026. On 7 October 2026, the exam page, study guide, all three learning paths, official lab collection and selected high-change product units were rechecked. This is not a claim that every linked page was re-read. Company cases, numerical examples and assessments are fictional teaching material. No live tenant was accessed or changed.

<https://learn.microsoft.com/en-us/training/modules/provision-brand-microsoft-365-tenant/>

<https://learn.microsoft.com/en-us/training/modules/license-microsoft-365-ai-services/>

<https://learn.microsoft.com/en-us/training/modules/maintain-microsoft-365-tenant-health/>

<https://learn.microsoft.com/en-us/training/modules/configure-exchange-teams/>

<https://learn.microsoft.com/en-us/training/modules/govern-sharepoint-onedrive-copilot/>

<https://learn.microsoft.com/en-us/training/modules/provision-govern-identities-entra/>

<https://learn.microsoft.com/en-us/training/modules/implement-authentication-access-entra/>

<https://learn.microsoft.com/en-us/training/modules/secure-email-collaboration-defender-office365/>

<https://learn.microsoft.com/en-us/training/modules/protect-information-purview/>

<https://learn.microsoft.com/en-us/training/modules/prevent-data-loss-lifecycle-purview/>

<https://learn.microsoft.com/en-us/training/modules/enable-manage-microsoft-365-copilot/>

<https://learn.microsoft.com/en-us/training/modules/configure-copilot-settings-extensibility/>

<https://learn.microsoft.com/en-us/training/modules/establish-agent-identities-entra-agent-id/>

<https://learn.microsoft.com/en-us/training/modules/curate-control-agents-registry-tools/>

<https://learn.microsoft.com/en-us/training/modules/agent-365-monitor-manage/>

<https://learn.microsoft.com/en-us/training/modules/agent-365-govern/>

<https://learn.microsoft.com/en-us/training/modules/monitor-optimize-m365-ai-services/>

- The official Microsoft AB-650 lab collection was successfully retrieved on 7 October 2026. Follow its current instructions for live portal work. This course does not supply the hosted lab tenant; our fixture exercises are simulations, not evidence of completed live labs. Open official labs

- Current Cowork controls: <https://learn.microsoft.com/en-us/microsoft-365/copilot/cowork/cowork-access>

- Agent tools: <https://learn.microsoft.com/en-us/microsoft-365/admin/manage/manage-tools-for-agent>

- Agent registry: <https://learn.microsoft.com/en-us/microsoft-365/admin/manage/agent-registry>

- SharePoint discovery and access units:

<https://learn.microsoft.com/en-us/training/modules/govern-sharepoint-onedrive-copilot/>

- Agent Conditional Access unit: <https://learn.microsoft.com/en-us/training/modules/establish-agent-identities-entra-agent-id/7-apply-conditional-access-agent-licensing>

Before teaching or booking: compare all 58 blueprint pointers; inspect exam status and appointment rules; recheck Cowork access controls, connector models, agent policy subjects, Purview workload limitations and licensing; revise affected questions and increment the version. Any future date-specific assertion needs fresh verification.